

XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

The XiangShan Team

HPCA'25@Las Vegas, USA

March 02, 2025



What we will cover in this tutorial

- Introduce of XiangShan (8:30 – 9:00, 30 minutes)
- CPU Microarchitecture (9:00 – 10:00, 60 minutes)
 - Design and implementation – How to implement novel ideas on XiangShan
 - Frontend: branch prediction and instruction fetch
 - Backend: out-of-order scheduler, execution units
 - Load/Store Unit: LSQ, pipelines, TLBs, data caches
 - L2/L3 caches and prefetchers
- Development workflows (10:30 – 12:00, 90 minutes)
 - Introduction of their usages – How to develop on XiangShan with MinJie
 - Simulation and Debugging
 - Research Demo



XiangShan: Open-Source High-Performance Processor

- **1st generation: YQH (Yanqihu)**
 - 2020/6: first commit of design RTL
 - 2021/7: **28nm tape-out, 1.3GHz**
 - Performance: **SPEC CPU2006 7.01@1GHz, DDR4-1600**
- **2nd generation: NH (Nanhu)**
 - 2021/5: starting design exploration and RTL design
 - 2023/4: GDSII delivery
 - **14-nm tape-out, estimated SPEC CPU2006 20@2GHz**
- **3rd generation: KMH (Kunminghu)**
 - ISA feature: **Vector (V) extension, Hypervisor (H) extension**
 - Support RVA23 profile and L2 Cache with CHI protocol
 - **estimated SPEC CPU2006 45@3GHz**
- **Open-sourced at <https://github.com/OpenXiangShan/XiangShan>**



Fragrant Hills in Beijing

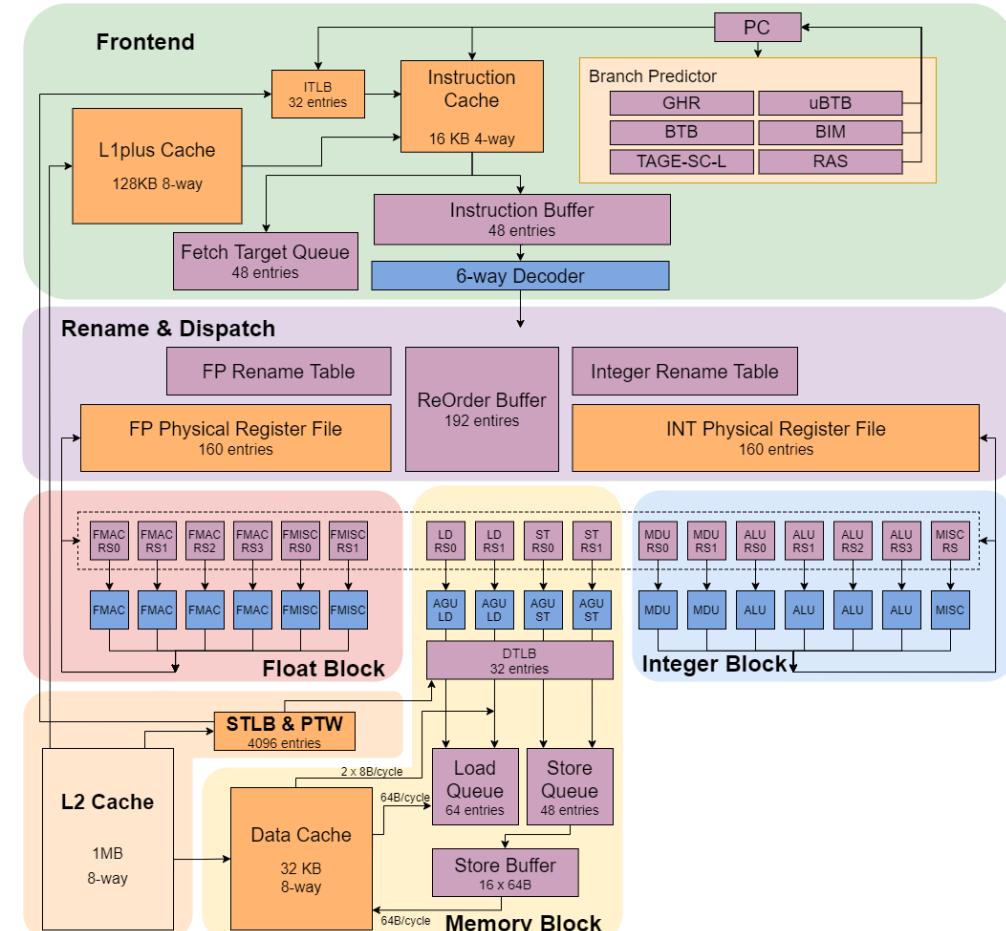
The screenshot shows the GitHub repository page for "XiangShan". It displays the repository's activity, including 153 branches and 5 tags. A list of recent commits is shown, such as "cz4e timing(MainPipe): remove set_conflict for tag/met..." and "coupledL2 @ bad8d17". The repository has 6.1k stars and 735 forks. The sidebar includes links for Readme, View license, Activity, Custom properties, 6.1k stars, 96 watching, 735 forks, and Report repository.

>6.1K stars, >735 forks on GitHub

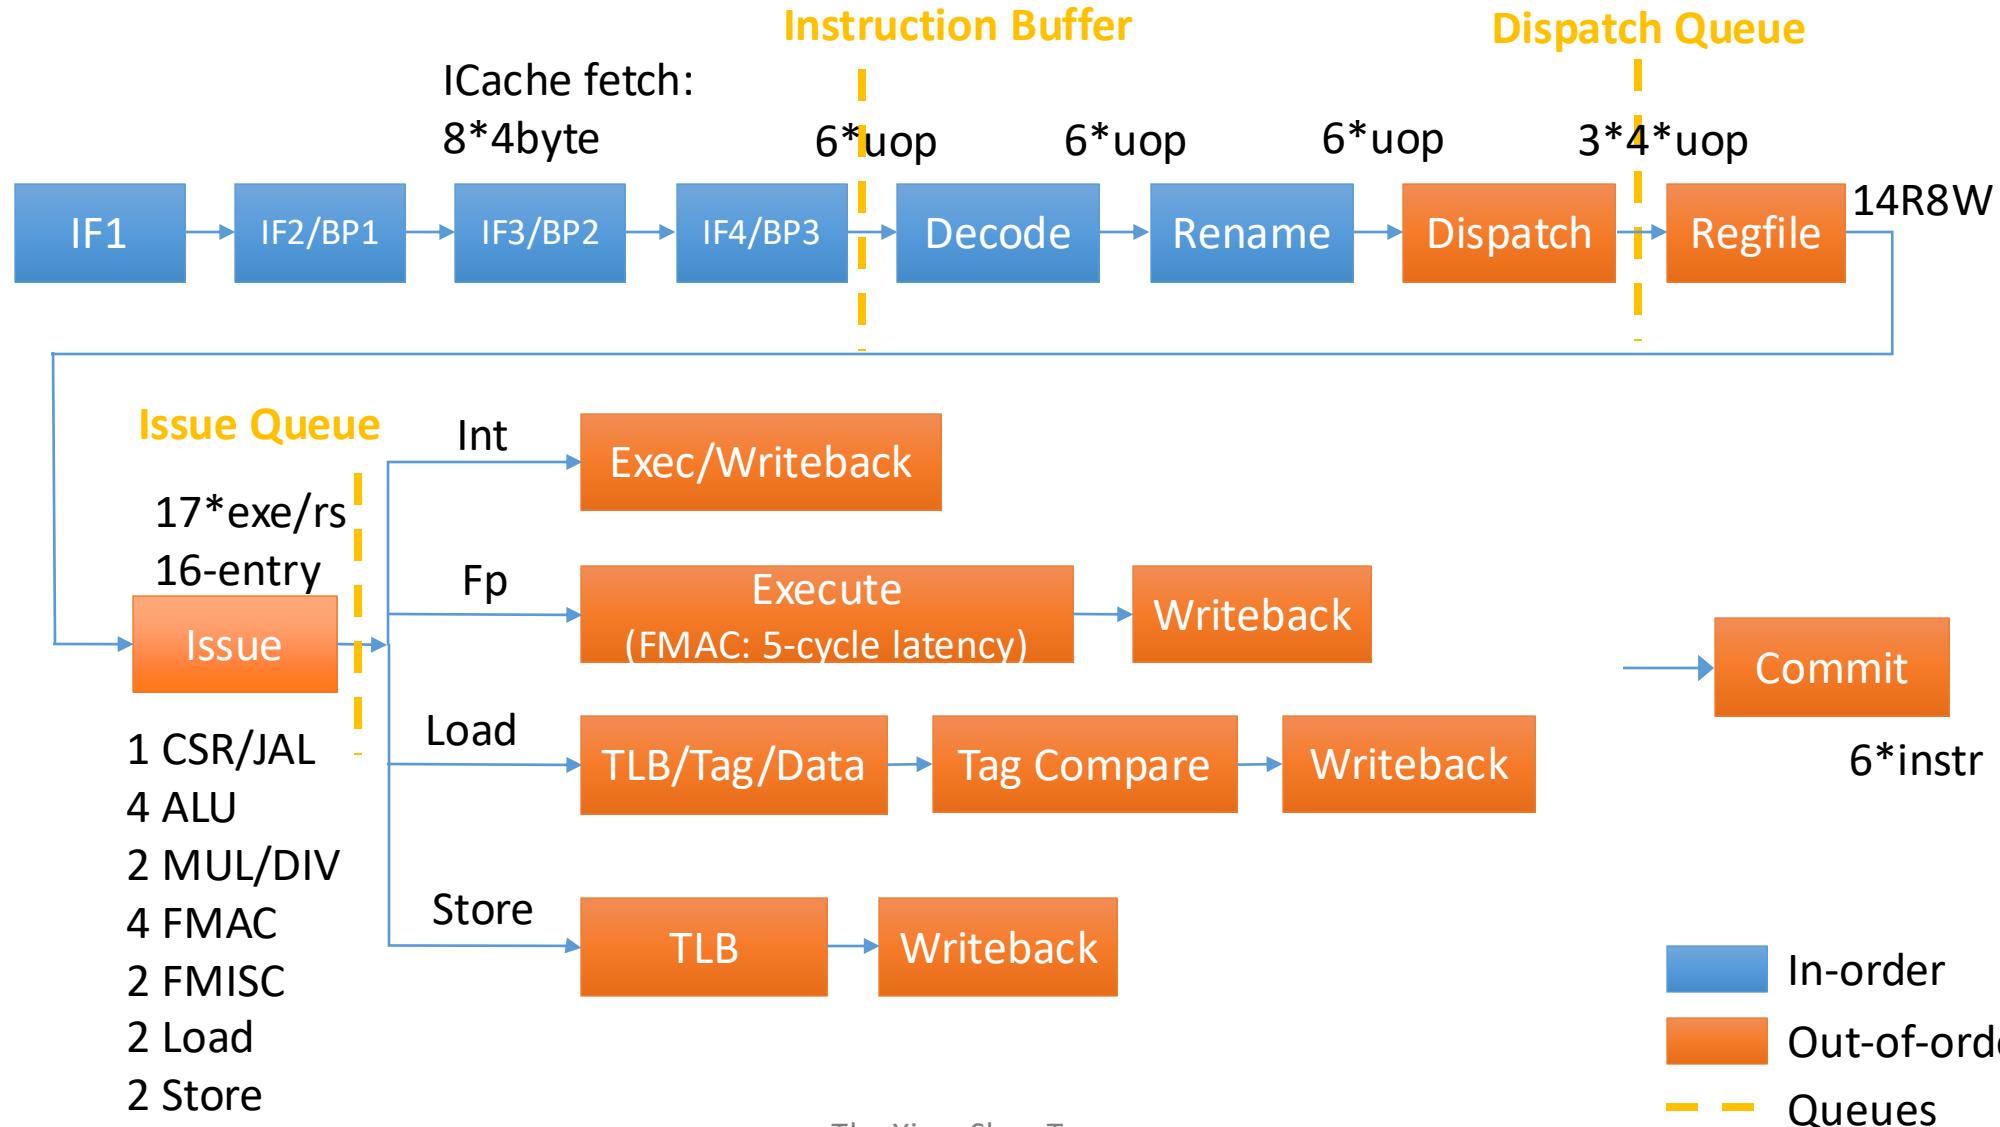


Yanqihu : 1st generation of XiangShan

- 11-stage, 6-wide decode/rename
- TAGE-SC-L branch prediction
- 160 Int PRF + 160 FP PRF
- 192-entry ROB, 64-entry LQ, 48-entry SQ
- 16-entry RS for each FU
- 16KB L1 Cache, 128KB L1plus Cache for instruction
- 32KB L1 Data Cache
- 32-entry ITLB/DTLB, 4K-entry STLB
- 1MB inclusive L2 Cache



Yanqihu µArch Overview





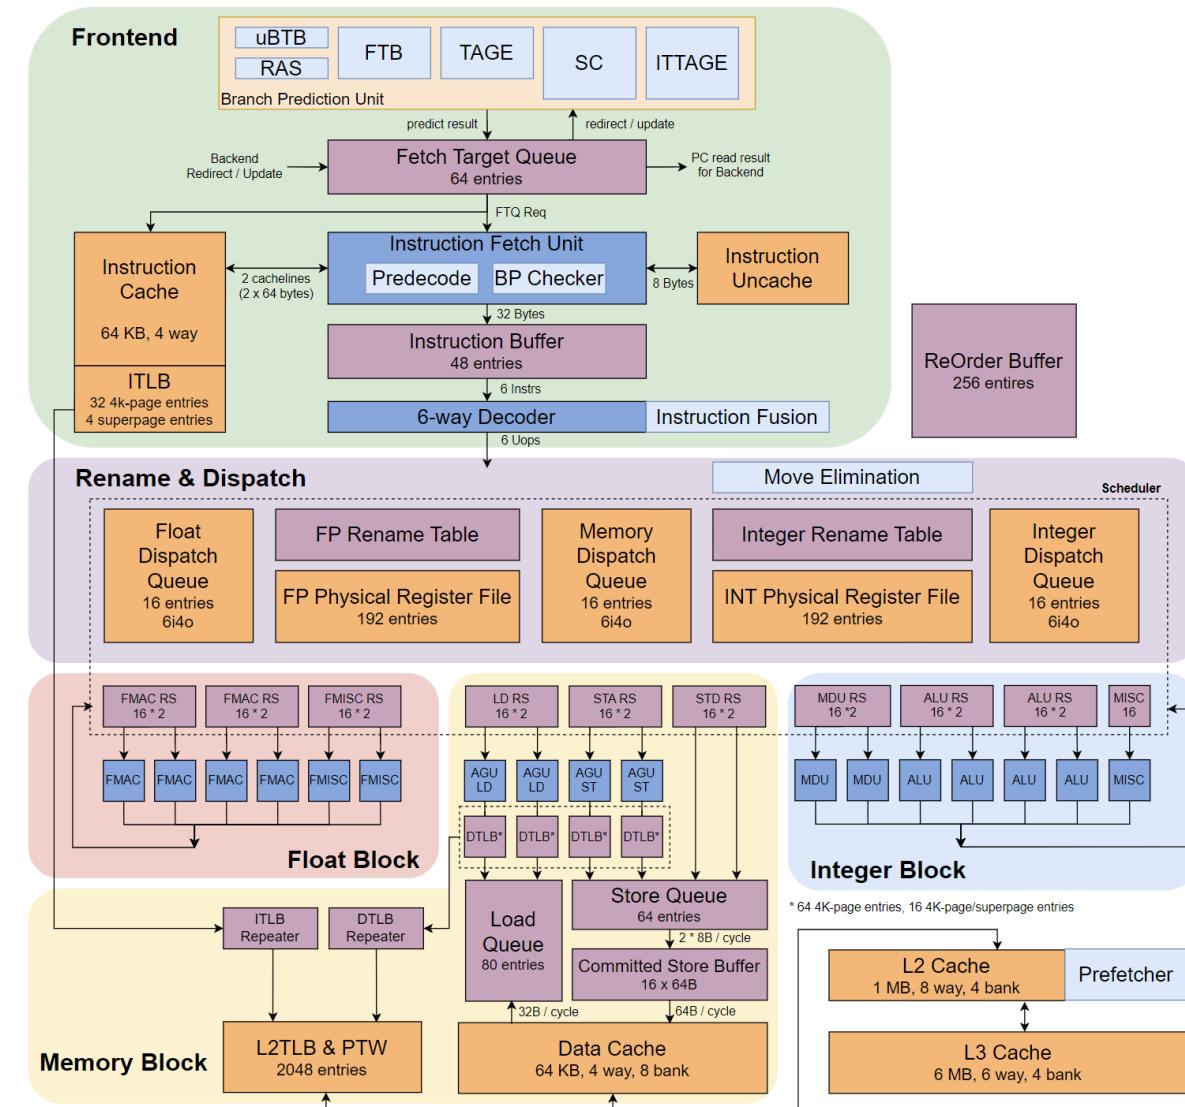
Overview: XiangShan Gen 2 Nanhu

- Named after a lake in Jiaxing, Zhejiang, China
- **Fetch & Branch Prediction**
 - Decoupled fetch and branch prediction unit in Frontend architecture
 - Higher throughput and prediction accuracy
- **New feature and Functional Support**
 - Support RISC-V Bit-Manipulation (B) and Cryptography (K) extensions
 - Support Physical Memory Protection (PMP) and configurable PMA
- **Load Store Unit**
 - Increase the capacity of TLB, Dcache and LSQ entries
 - Redesign DCache pipeline to reduce bank conflicts and improve efficiency
- **L2/L3 Cache**
 - Develop an open-source high performance L2 Cache / LLC named **HuanCun**



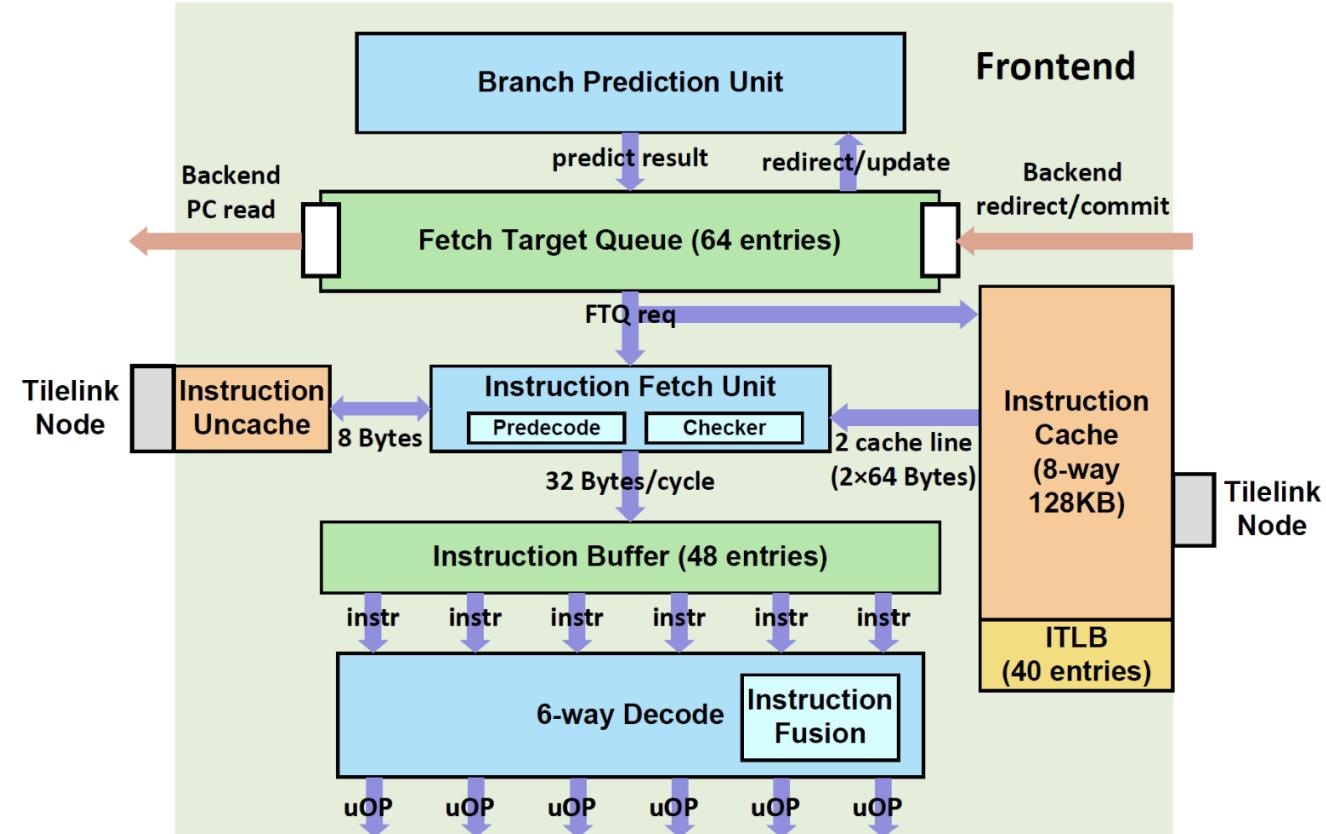
Nanhу μArch Overview

- **192** Int PRF + **192** FP PRF
- **256**-entry ROB, **80**-entry LQ, **64**-entry SQ
- **32**-entry RS for two FUs
- **64KB** L1 Instr. Cache, **64KB** L1 Data Cache
- **64**-entry DM + **16**-entry FA DTLB
- **32**-entry ITLB, **2K**-entry STLB
- **1MB** non-inclusive L2 Cache
- **6MB** non-inclusive L3 Cache (LLC)



Frontend Overview

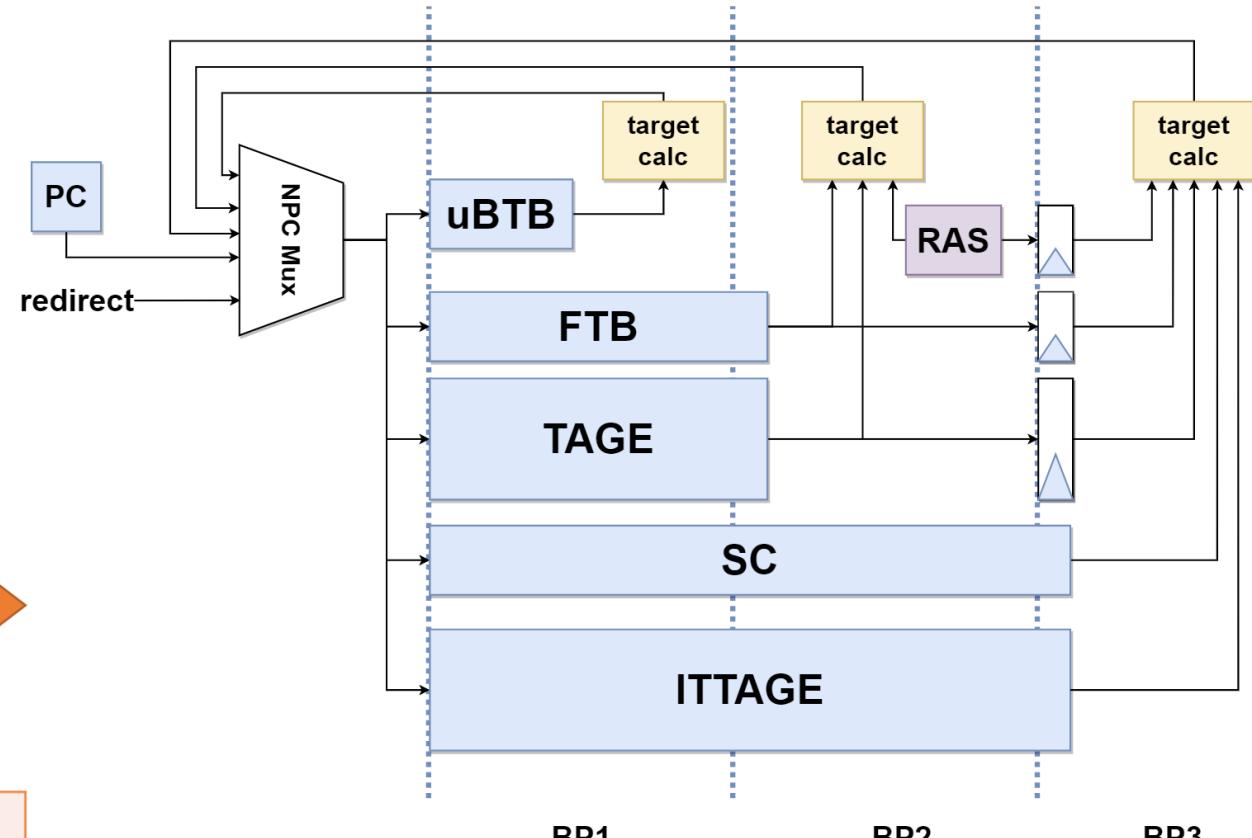
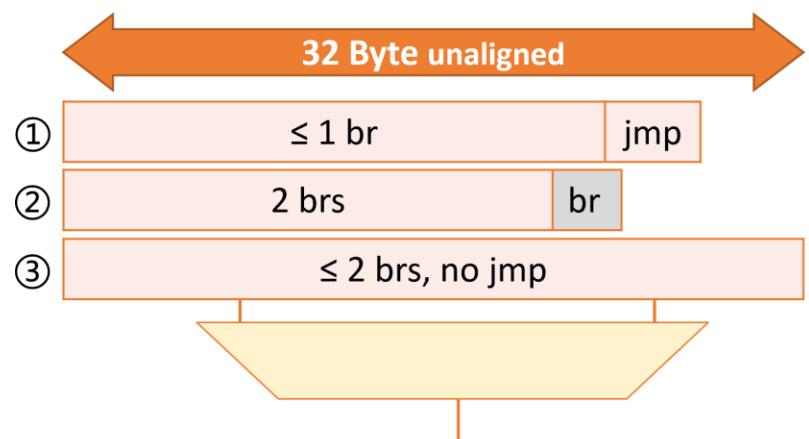
- Decoupling of Fetch and Branch Prediction Units
 - **Producer**: branch predictor unit
 - **Consumer**: instruction fetch unit
- Merit
 - Performance
 - Hide branch prediction bubbles
 - Guide instruction prefetching
 - Timing
 - Avoid interact of BPU and ICache
 - Other Detailed Optimization



Frontend — Branch Predictor

Three-stage overriding prediction

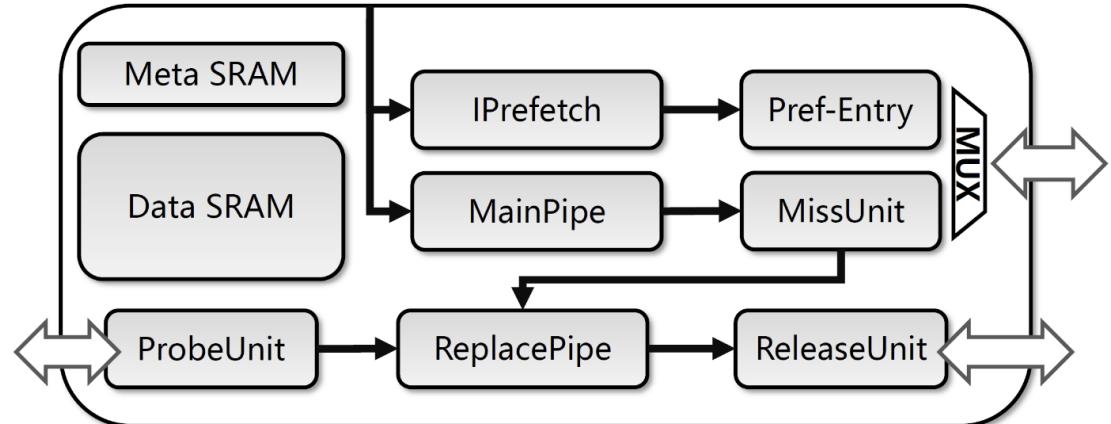
- Stage 1: uBTB
 - Stage 2: FTB + TAGE + RAS
 - Stage 3: ITTAGE + SC
-
- FTB: Fetch Target Buffer



Frontend — Instruction Cache

- **Main feature**

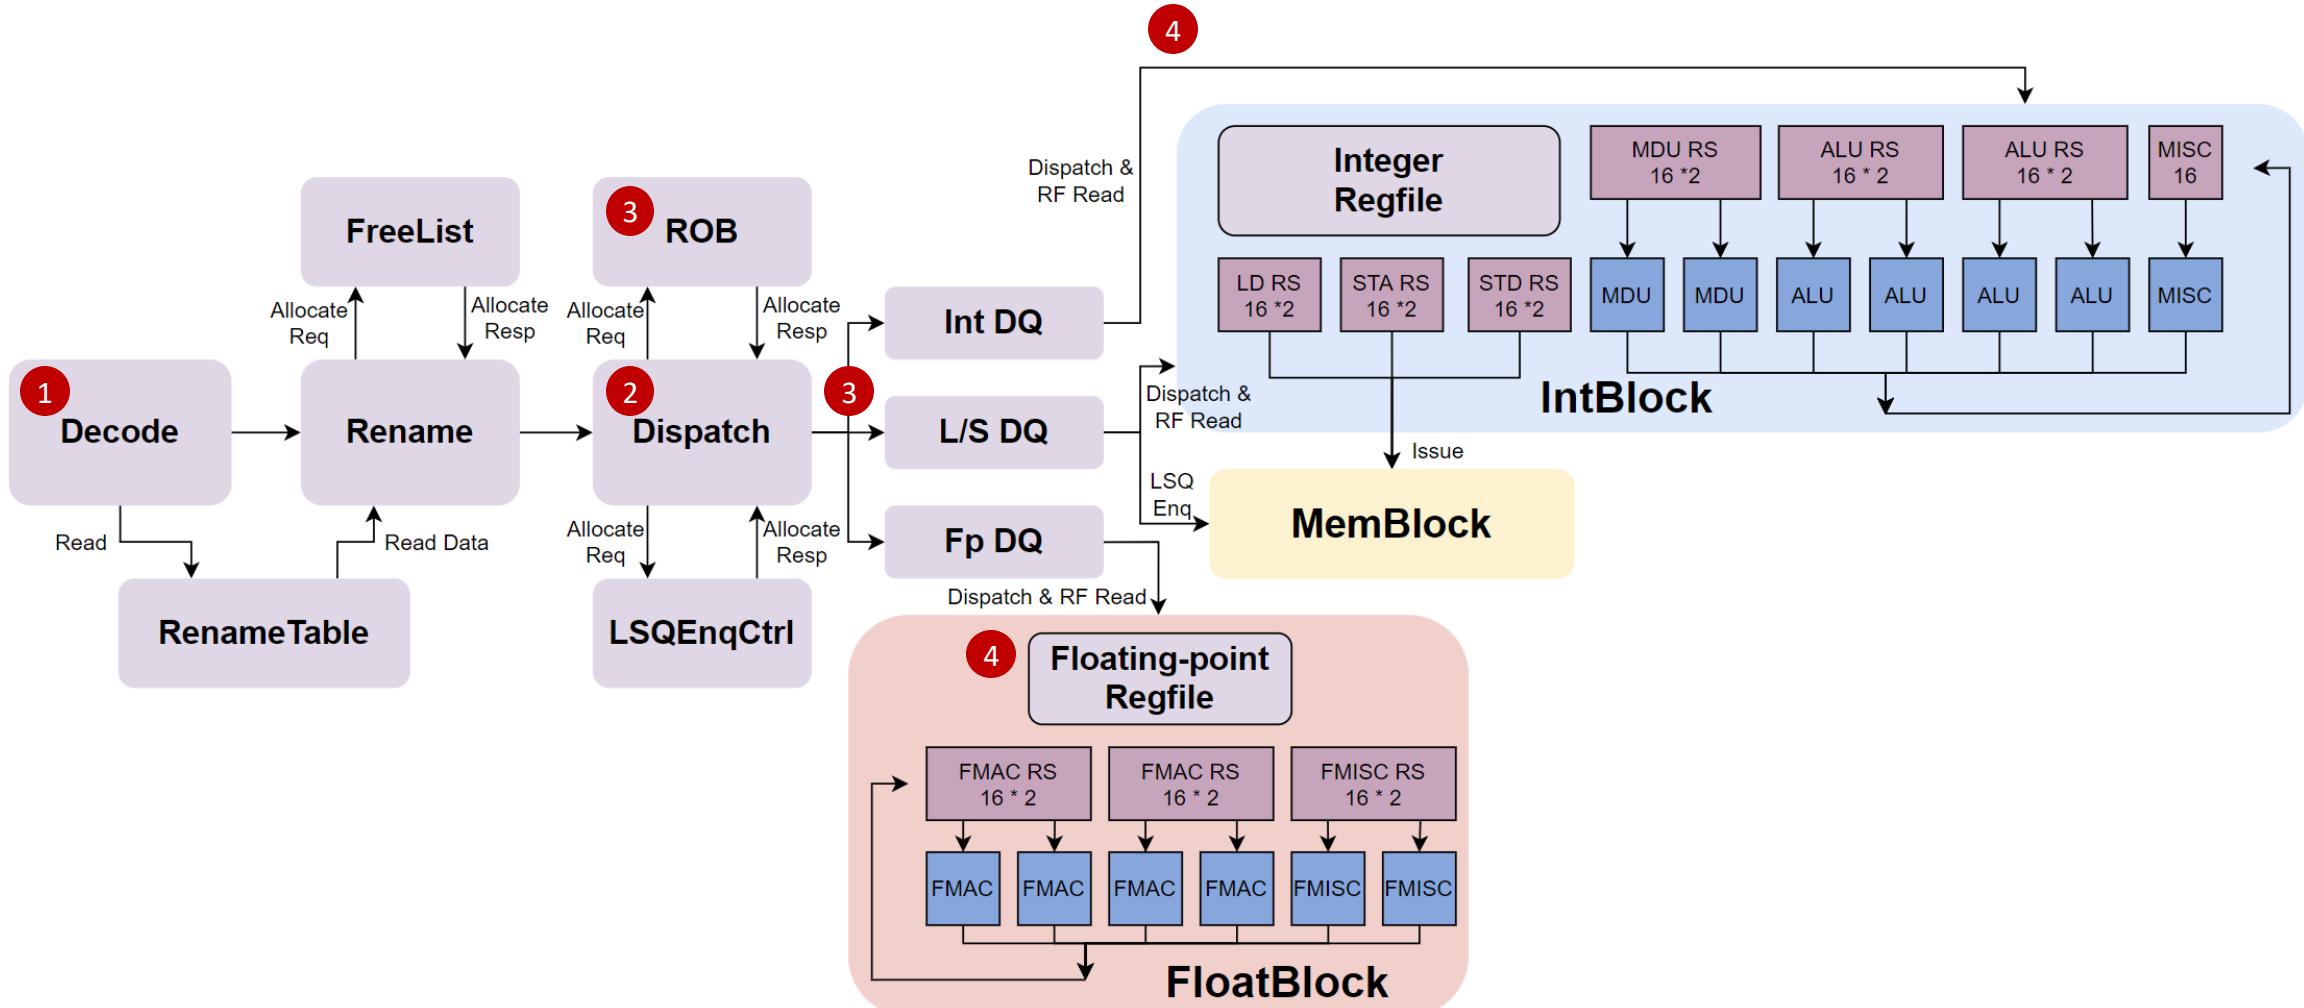
- 64KB 4-w VIPT blocking cache
- Even/odd bank interleaving read
- **Support TileLink protocol**
- FDIP instruction prefetching



	Yanqihu	Nanhu	
L1 ICache	16KB, 4-w	64KB, 4-w	4×
L1plus Cache	128KB, 8-w	-	Reduce
Read bandwidth	64 B	128 B	2×
TileLink Support	No	Yes	New
Instruction prefetch	Stream	L2 FDIP	New



Backend Overview





Backend — Instruction Fusion

- Fusion of adjacent UOPs
 - Improve backend efficiency
 - Reduce data bypass delay
- Fusion Types
 1. Reuse of the original calculation
 2. Add new type of calculation

slli r1, r0, 32
srl r1, r1, 32

add.uw r1, r0, zero

Type1: Reusing the original calculation

srl r1, r0, 8
andi r1, r1, 255

byte2 r1, r0

Type2: Creating new type of calculation

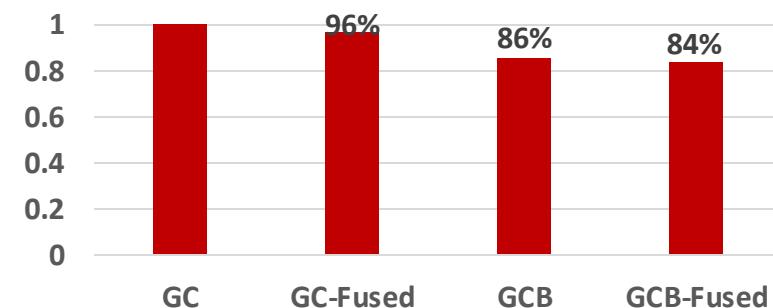


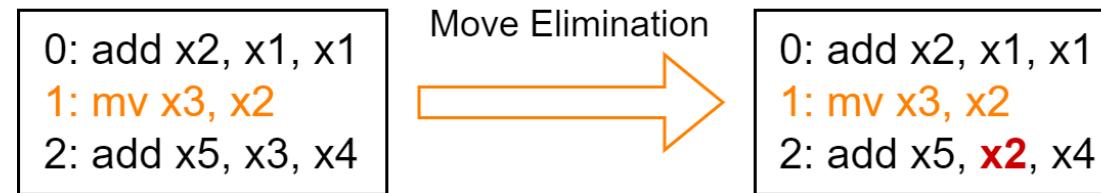
Figure. Normalized Dynamic Instruction Count for CoreMark (-O2)



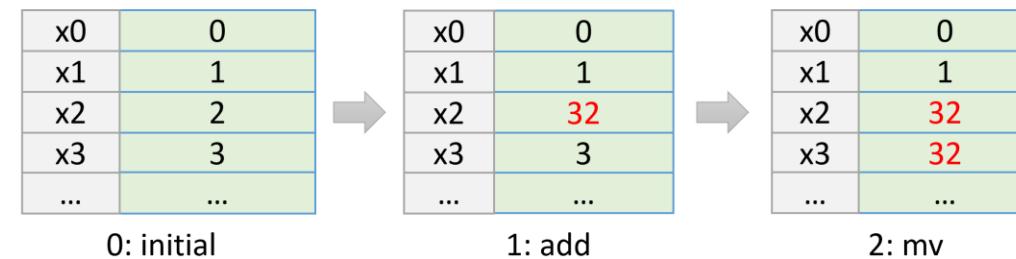
Backend — Rename & Move Elimination

- RenameMap: the mapping between **arch registers** and **physical registers**
- Move Elimination
 - No need to allocate new physical register by just changing the RenameMap
 - Mark as **Complete** when Move instruction enters ROB
 - Use **Reference-Counter** to record mapping counts of physical registers

program execution flow:



rename table mapping:

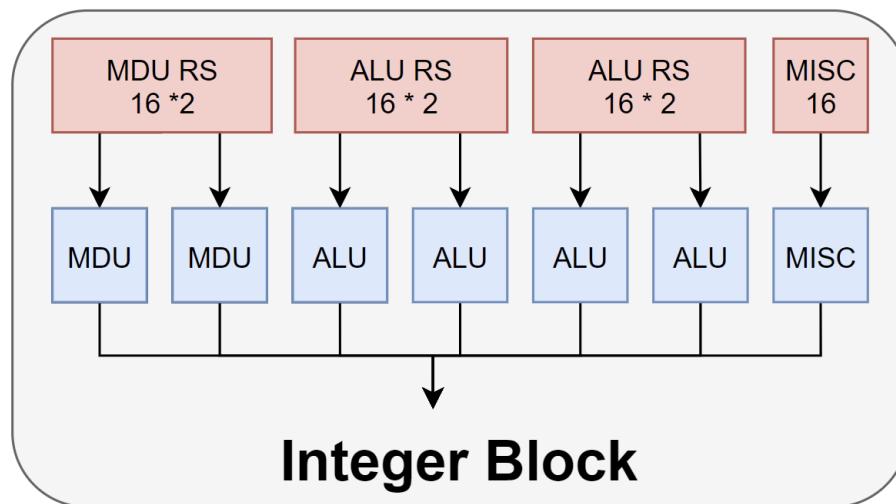




Backend — Reservation Station

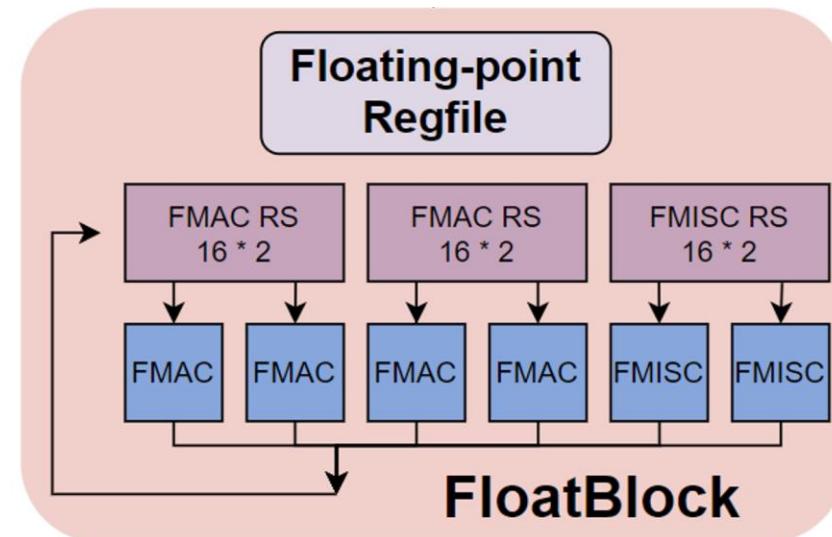
- Main Changes:

- 1i1o -> 2i2o
- based selection



0	true			
1	true	true		
2	true	false	true	
3	true	false	false	true
	0	1	2	3

Age Matrix





Backend — Floating Point Units (FPUs)

- Industry competitive Floating Point Unit, IEEE 754 compatible

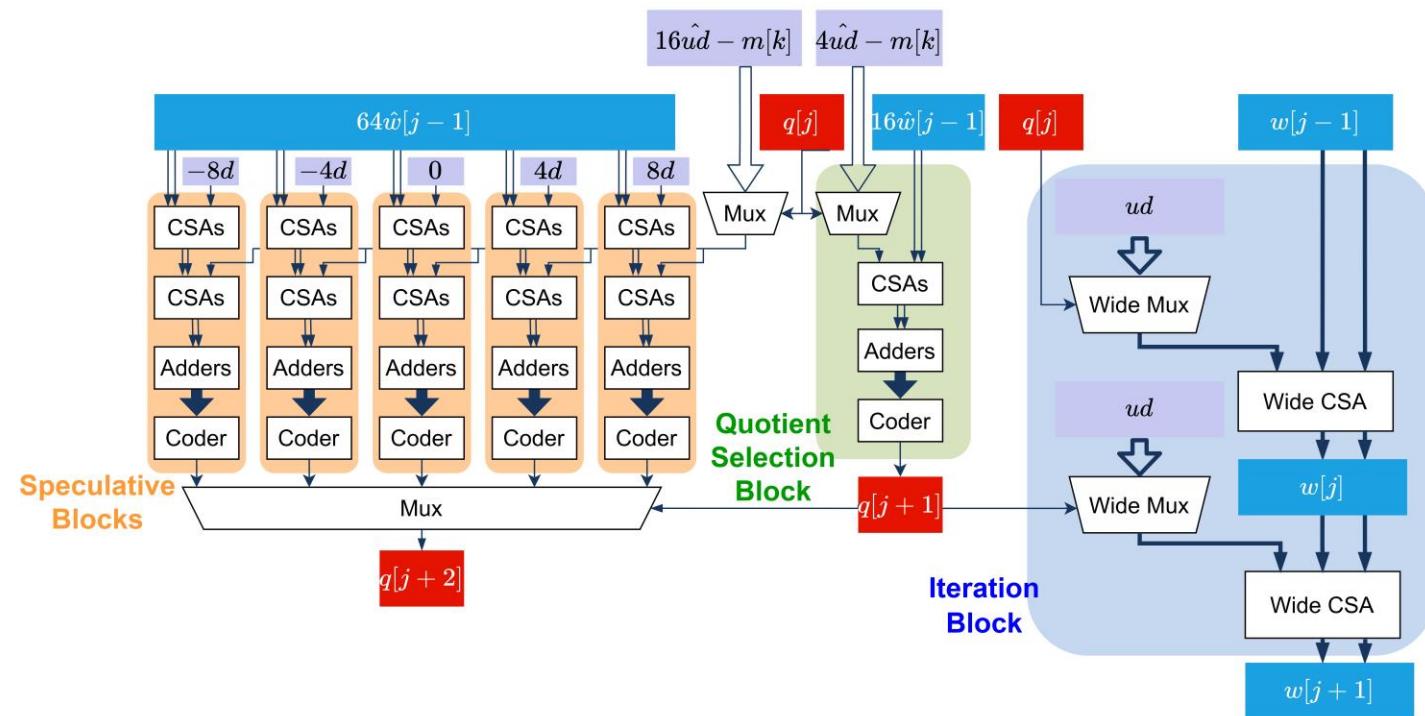
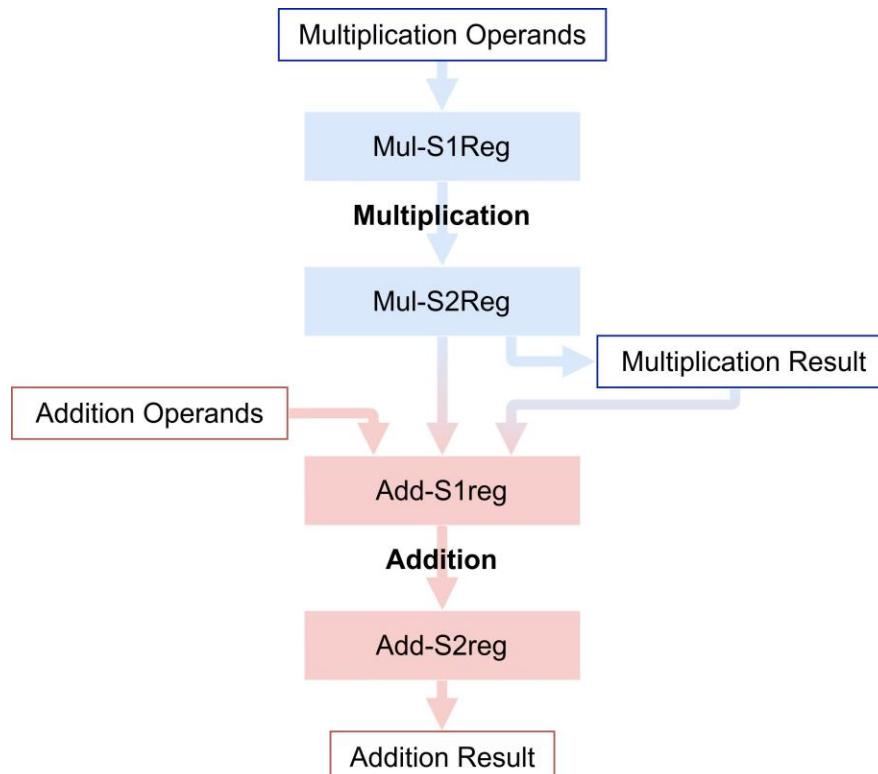
Operations	Latency
FADD	3
FMUL	3
FMA	5
FDIV	≤ 11 (float) ≤ 18 (double)
FSQRT	≤ 17 (float) ≤ 31 (double)
FCVT, FCMP	3

The screenshot shows the GitHub repository page for `OpenXiangShan/fudian`. The repository is public and has 15 branches, 0 tags, and 33 commits. The commit history includes initial commits for FADD, .gitignore, .gitmodules, .mill-version, and .scalafmt.conf. The repository is described as an open-source high-performance IEEE-754 floating unit with 31 stars, 8 watching, and 12 forks.

<https://github.com/OpenXiangShan/fudian>

Backend — FPU

- Cascade FMA
 - Reduce FADD delay 5->3
- SRT16 Division units
 - 1/2 delay compared with YQH (SRT4)





Backend — B/K Extension Implemented

- **B:** Bit-Manipulation extension
- **K:** Cryptographic extension

1. Extensions
1.1. Zba extension
1.2. Zbb: Basic bit-manipulation
1.2.1. Logical with negate
1.2.2. Count leading/trailing zero bits
1.2.3. Count population
1.2.4. Integer minimum/maximum
1.2.5. Sign- and zero-extension
1.2.6. Bitwise rotation
1.2.7. OR Combine
1.2.8. Byte-reverse
1.3. Zbc: Carry-less multiplication
1.4. Zbs: Single-bit instructions

B Extension

2. Extensions Overview
2.1. Zbkb - Bitmanip instructions for Cryptography
2.2. Zbkc - Carry-less multiply instructions
2.3. Zbx - Crossbar permutation instructions
2.4. Zknd - NIST Suite: AES Decryption
2.5. Zkne - NIST Suite: AES Encryption
2.6. Zknh - NIST Suite: Hash Function Instructions
2.7. Zksed - ShangMi Suite: SM4 Block Cipher Instructions
2.8. Zksh - ShangMi Suite: SM3 Hash Function Instructions
2.9. Zkr - Entropy Source Extension
2.10. Zkn - NIST Algorithm Suite
2.11. Zks - ShangMi Algorithm Suite
2.12. Zk - Standard scalar cryptography extension
2.13. Zkt - Data Independent Execution Latency

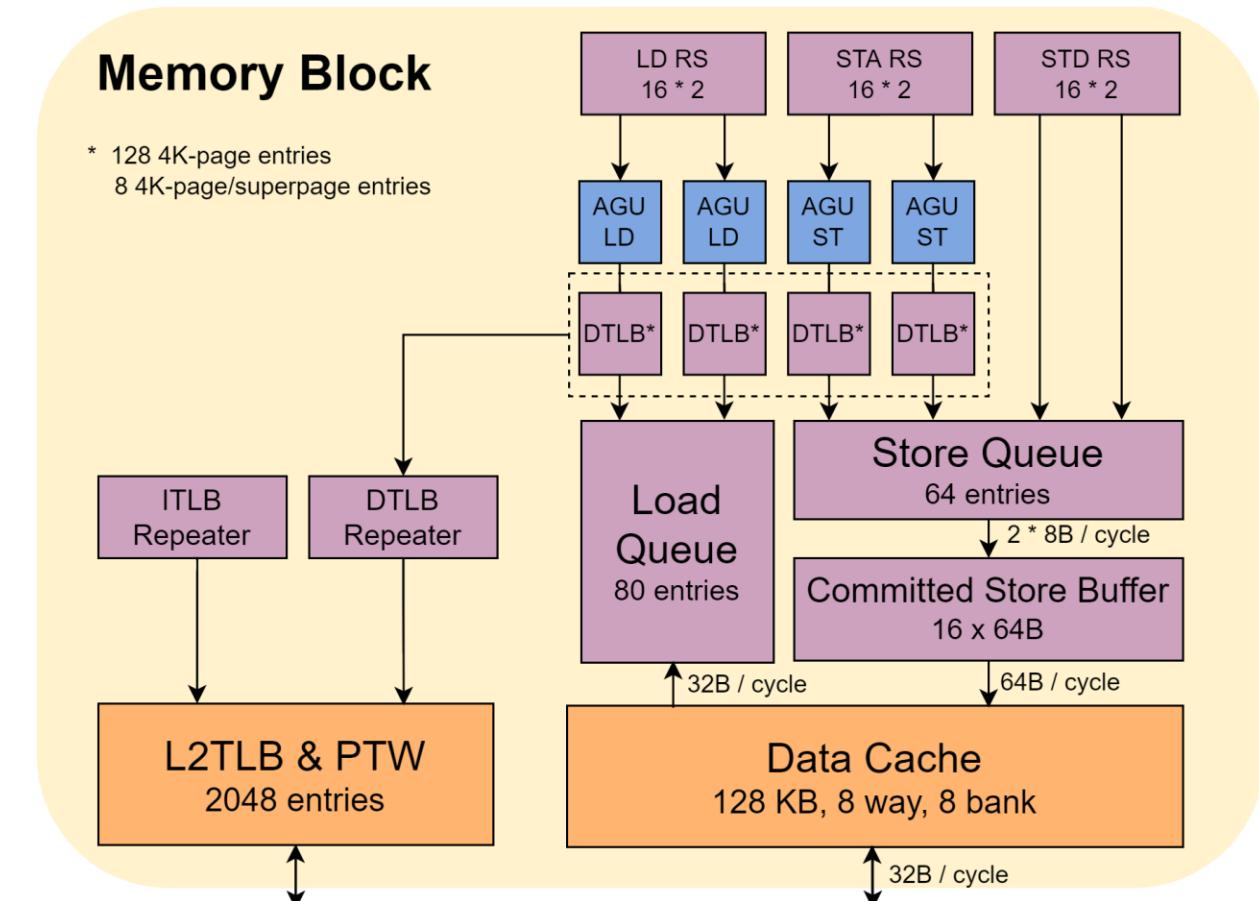
K Extension



Memory Block Overview

- Datasheets

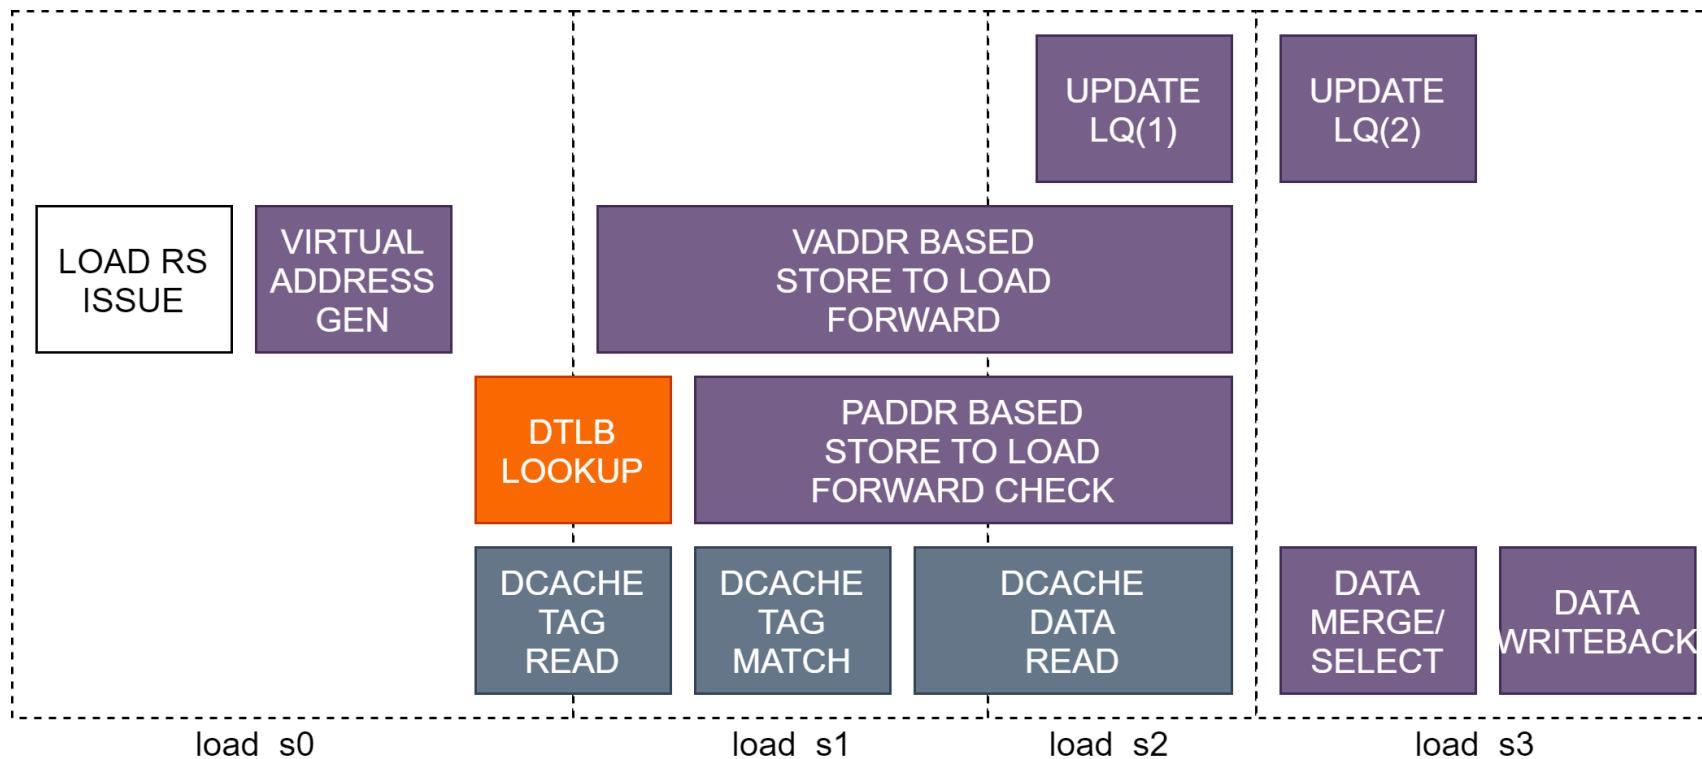
Load pipeline	2
L1 LD hit bandwidth	2x8B/cycle
L1 LD hit latency	4
Store address pipeline	2
Store data pipeline	2
Store data bandwidth	2x8B/cycle
Load Queue Entry	80
Store Queue Entry	64
Merged Store Buffer	16 cachelines
L1 Data cache size	64KB/128KB-8w
DTLB size	64 + 16
L2 TLB size	2k entries





MemBlock – Load Pipeline

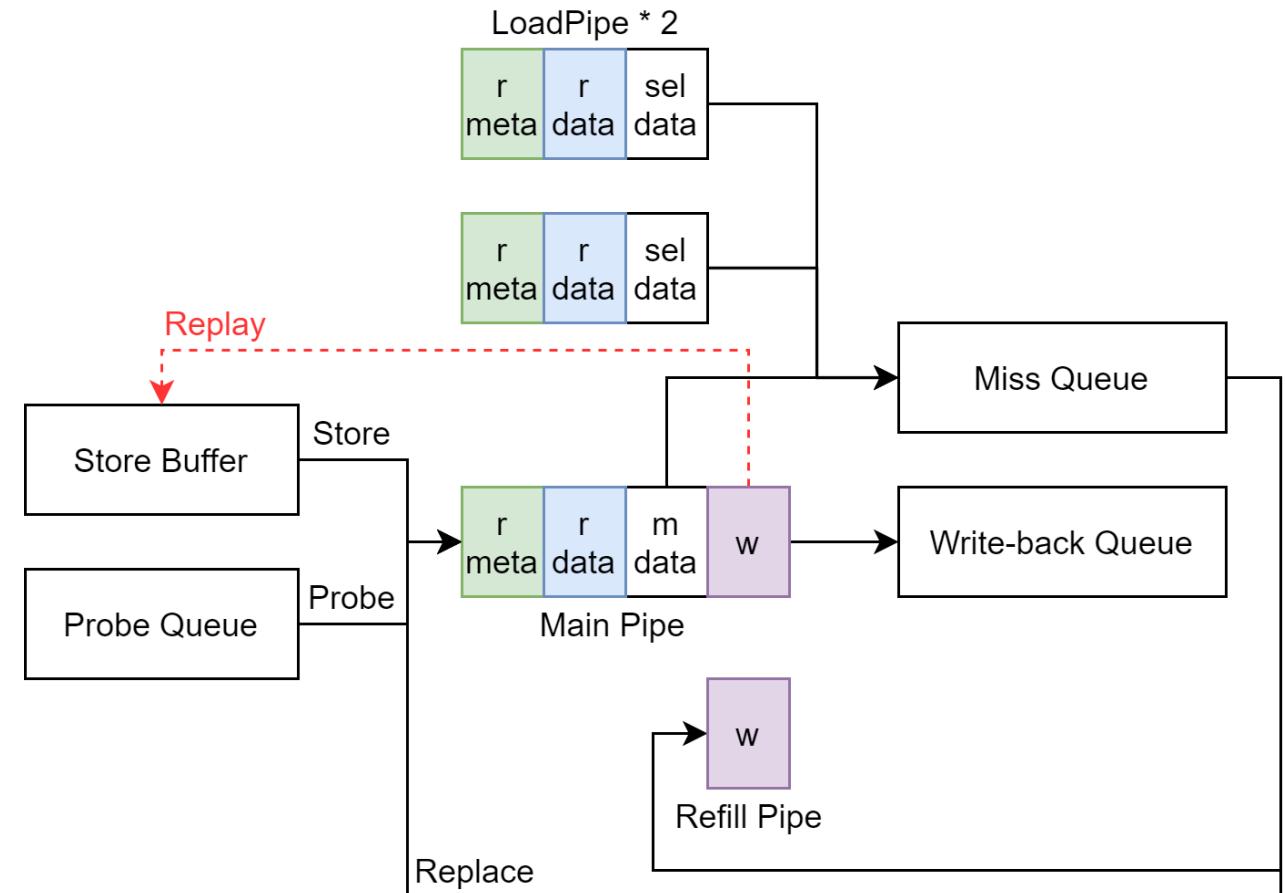
- Load Pipeline stages: three(YQH) -> **four** (NH)
- Adjust the read & write logic of tag / data / LQ
- Optimized for load-load bypass & load-store forward



MemBlock – L1 DCache

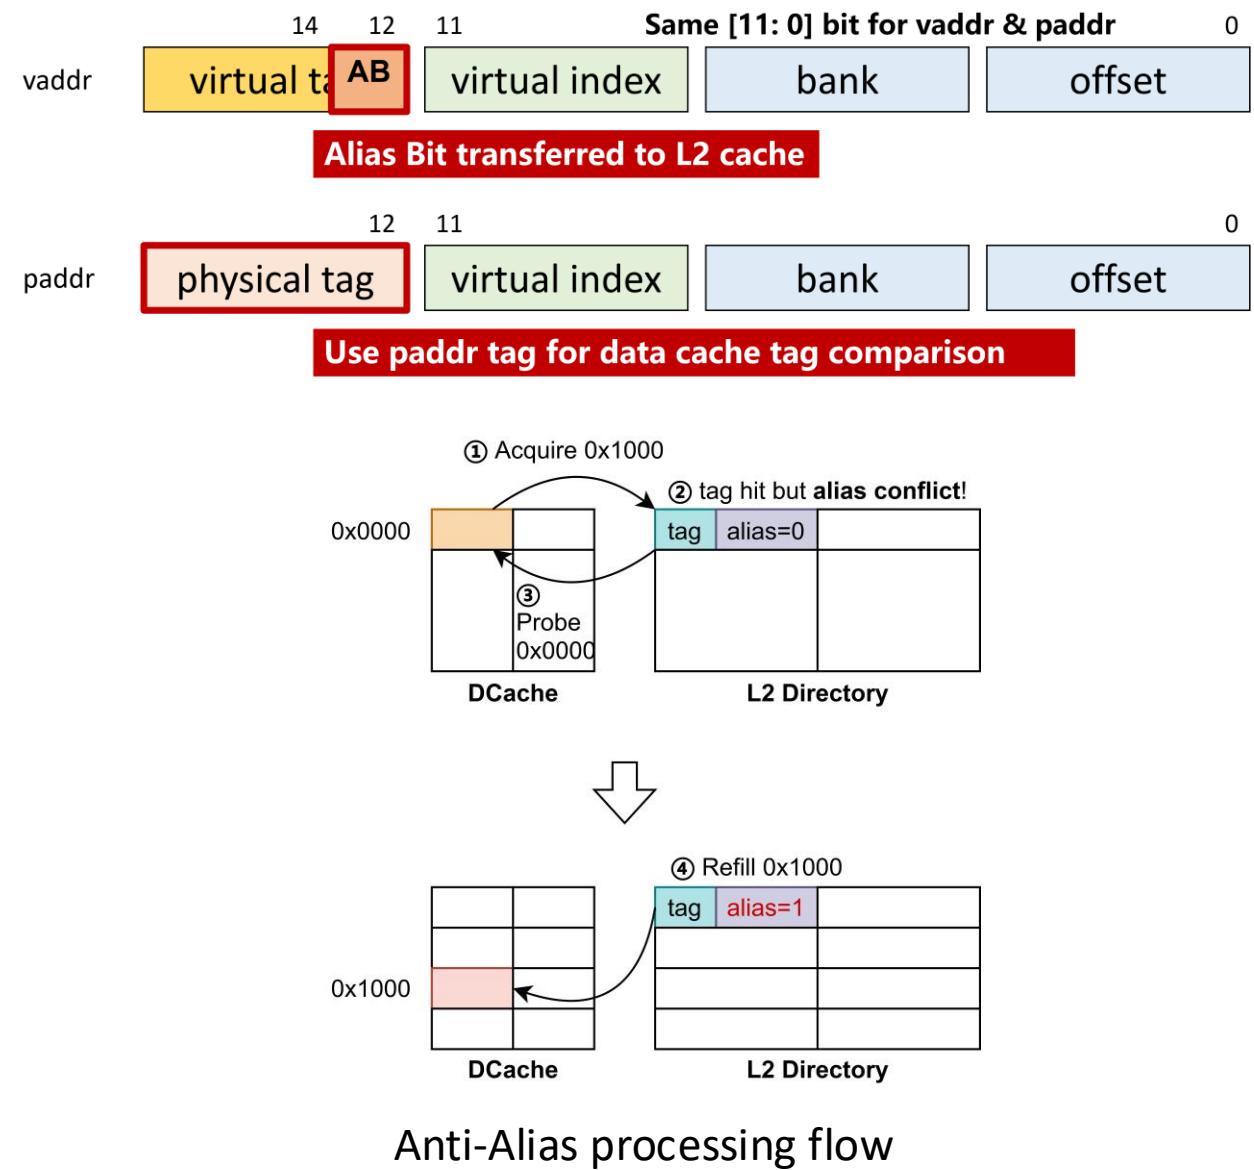
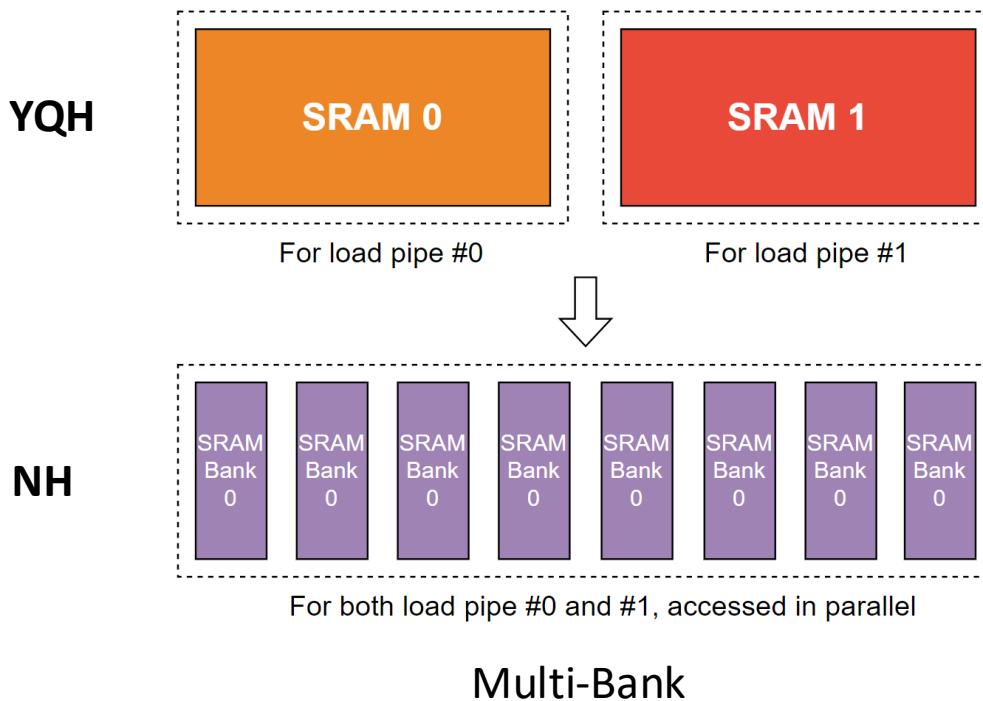
- Optimize store buffer and miss handling logic
- Support larger capacities
- I/D Cache coherence

L1 Data cache size	64KB/128KB-8w
L1/L2 Bus Width	256bit
Store Buffer	16
Probe Queue	8
Miss Queue	16
Write-back Queue	18
Software Prefetch	Support
ECC	Support



MemBlock – L1 DCache

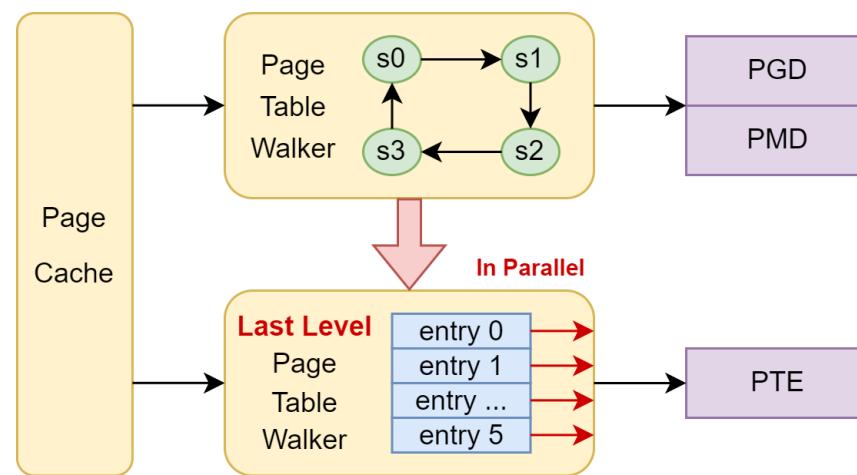
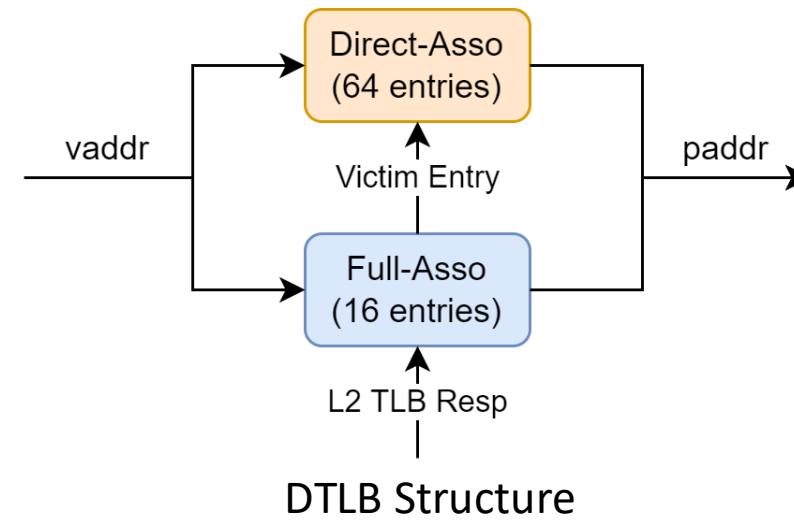
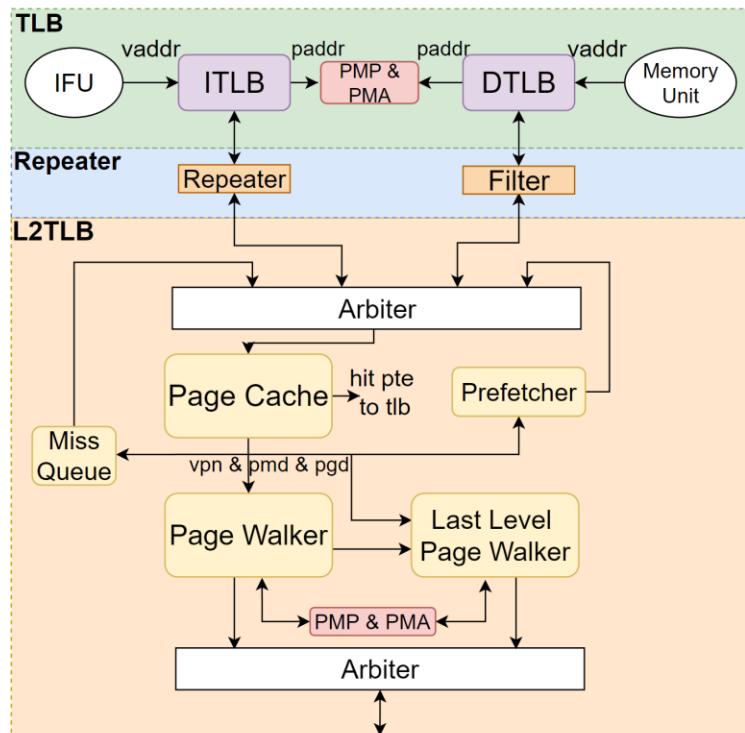
- Support larger cache capacities
 - Multi-bank for multi read ports
 - Anti-alias under larger size





MemBlock — Memory Management Unit

- Principal Improvement:
 - Larger Data TLB size
 - Parallel Page Table Walks

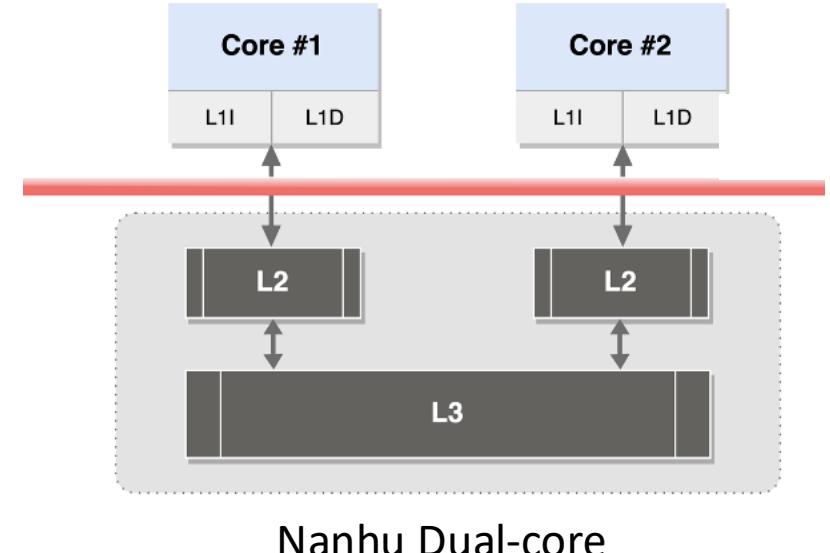


Parallel Page Table Walk



L2/L3 Cache (HuanCun)

- High performance non-blocking **L2/L3** cache
- Design features
 - **Directory based coherence**
 - **Optional Inclusive/Non-inclusive**
 - **TileLink protocol**
 - **Optimise latency and improve concurrency**
- Highly configurable parameters
 - #Size, #Banks, #Ways, #MSHRs
 - Replacement: **Random/PLRU/RRIP**
 - Prefetch: **BOP/SMS**



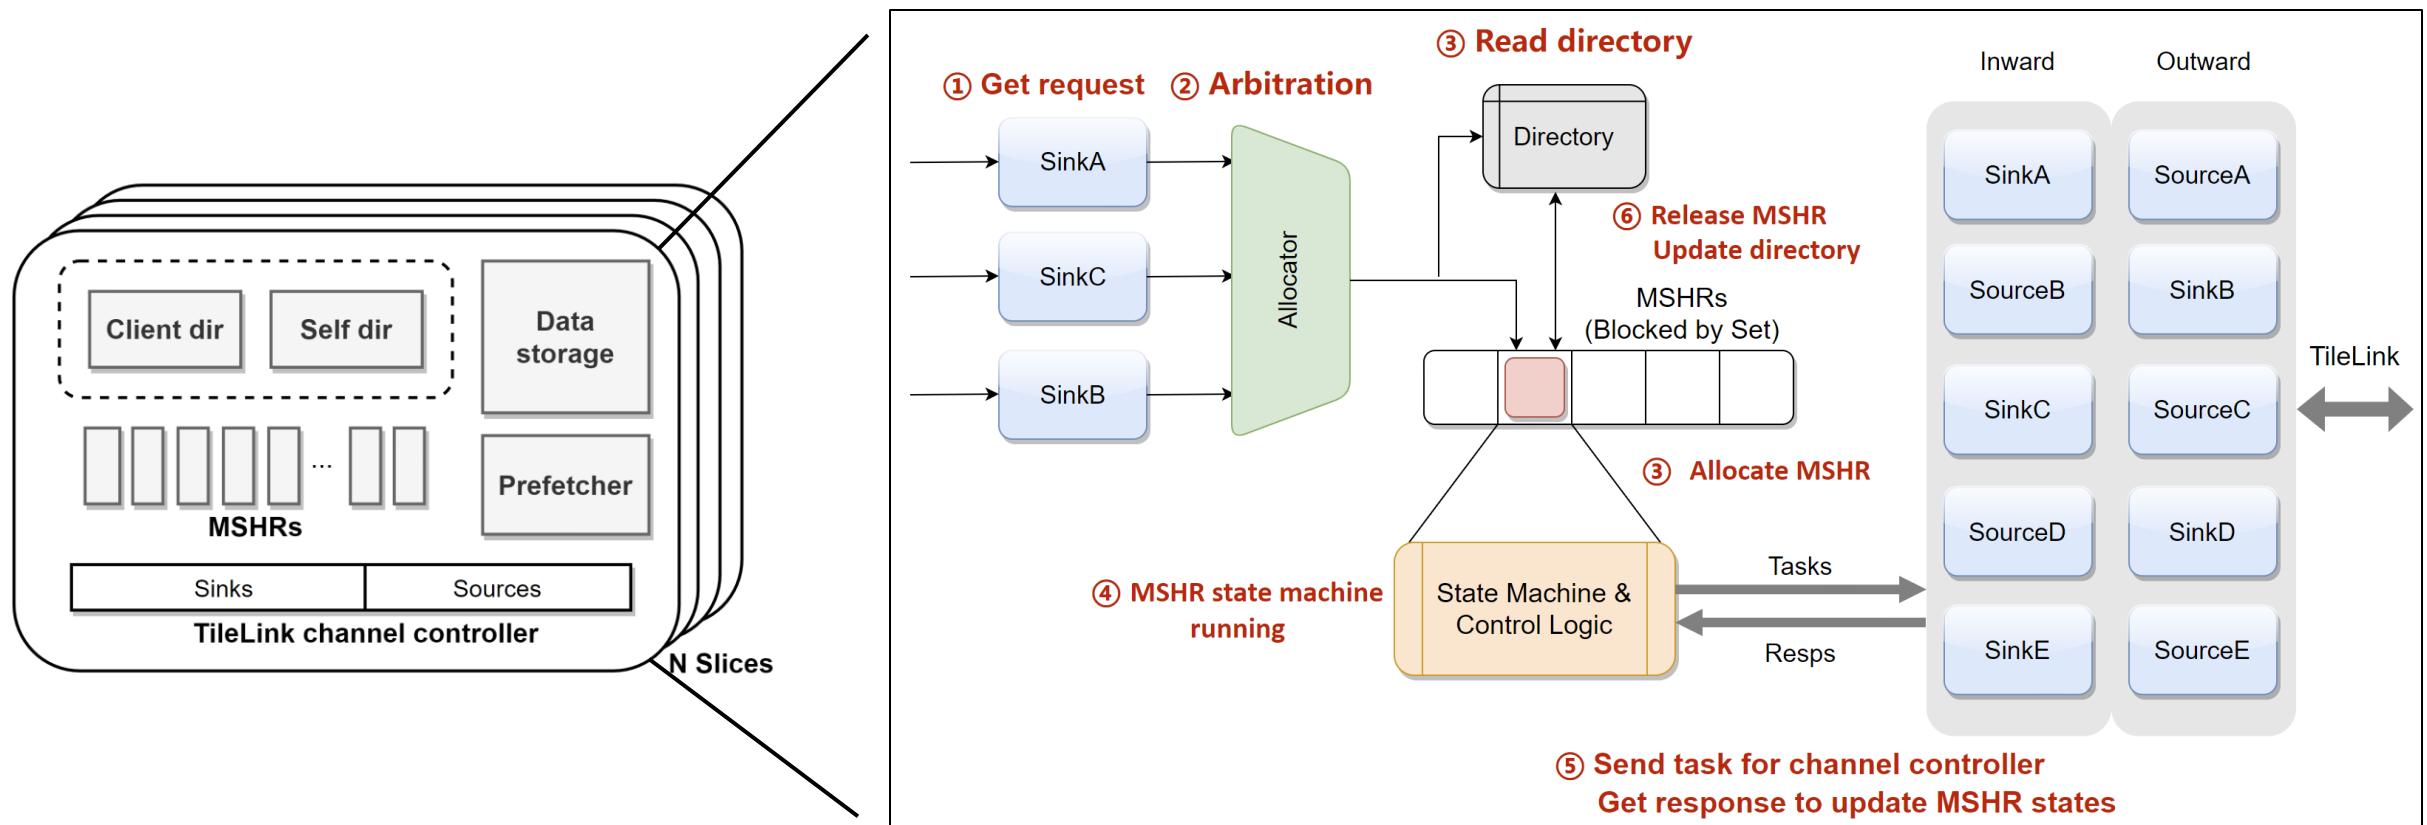
Nanhu Dual-core

A screenshot of the GitHub repository page for OpenXiangShan/HuanCun. The repository has 36 branches and 0 tags. The commit history shows several merges and updates, including changes to .github/workflows, rocket-chip, scripts, and src, along with various configuration and build-related files like .gitignore, .gitmodules, mill-version, scalafmt.conf, LICENSE, LICENSE.Sifive, Makefile, README.md, and build.sc. The commits are dated from 8 days ago to 10 months ago.

<https://github.com/OpenXiangShan/HuanCun>

L2/L3 Cache (HuanCun)

- Transaction processing flow





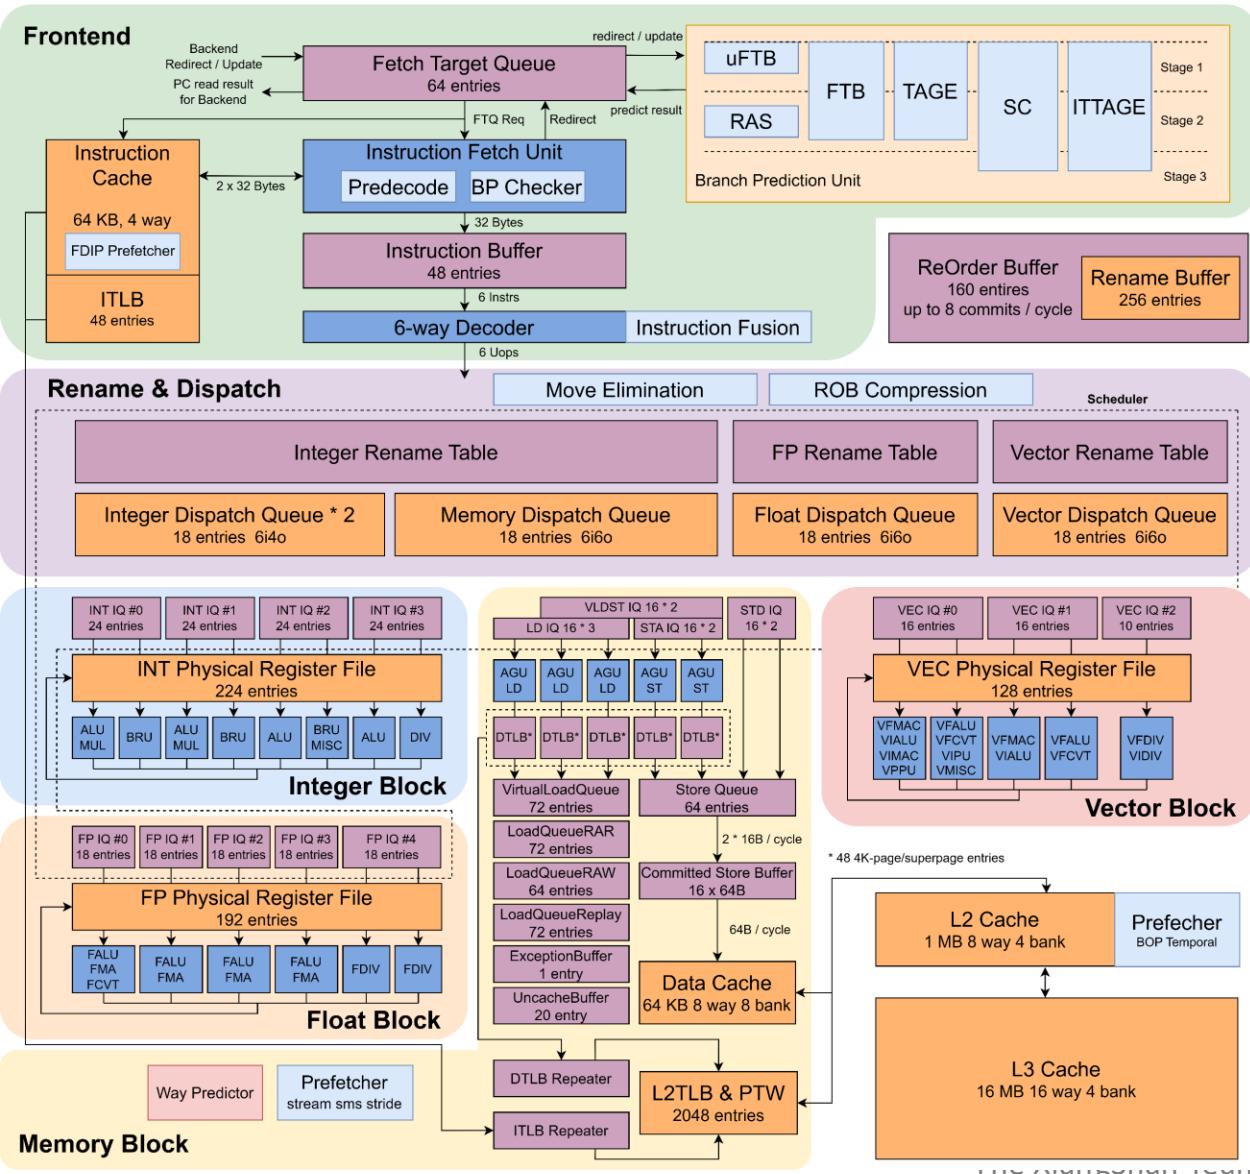
Overview: XiangShan Gen 3 Kunminghu (KMH)

- **Work with Beijing Institute of Open Source Chip (BOSC)**
 - Joint validation with partner companies based on open source
- **Functional Improvement**
 - Support RISC-V **RVA23 profile**
 - RISC-V **Vector** and **Hypervisor** extension implementation
- **Performance Exploration**
 - Performance upgrade of Frontend, Backend, load-store unit and cache
 - DSE on performance model (**XS-Gem5**) calibrated with RTL
- **Functional Verification & Physical Design**
 - Hierarchical verification flow including **UT, IT, ST + FPGA**
 - Industrial-grade verification process
 - Simultaneous iteration of RTL coding with physical design timing evaluation



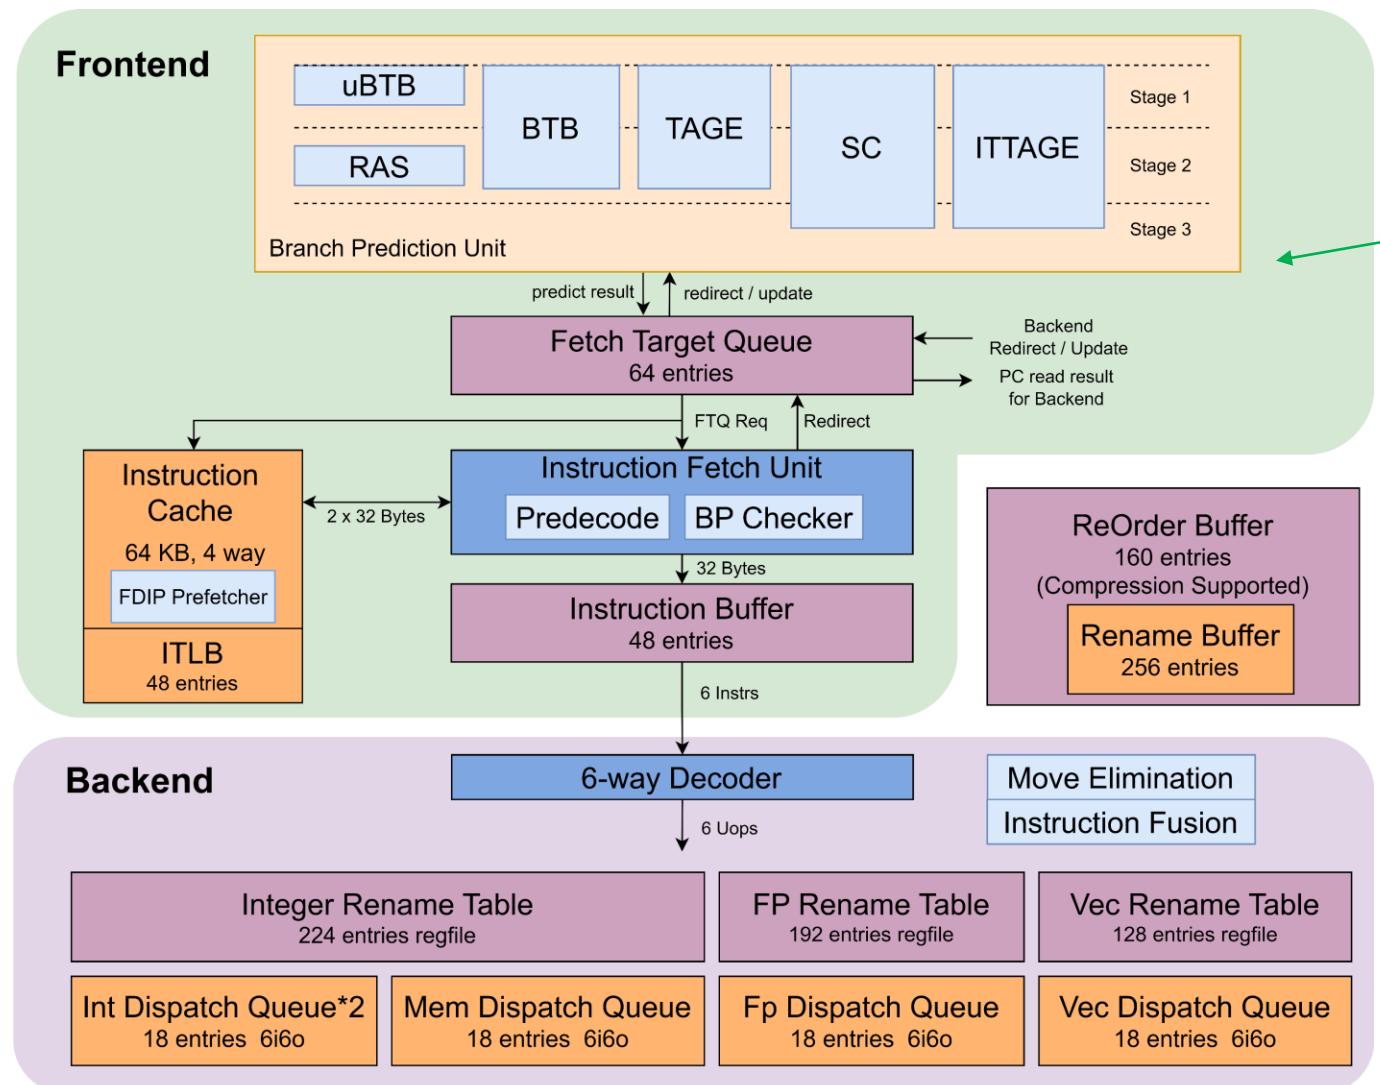


3rd Kunminghu µArch Overview



- **Decoupled frontend**
 - Reduce fetch bubbles
 - Instruction prefetching friendly
- **Aggressive outstanding instruction window**
 - Large ROB/LQ/SQ
- **Low latency & high BW \$ access**
 - Closely-coupled load/store pipe & L1/L2\$
 - Highly-banked design
 - Multi-level composite prefetching
- **ISA support**
 - Support RISC-V RVA23 profile
 - RVWMO (RISC-V Weak Memory Ordering)

3rd Kunminghu µArch Overview

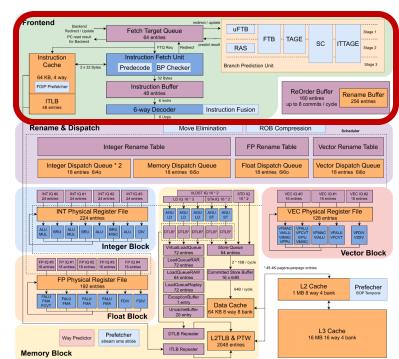


Multi-level branch predictors

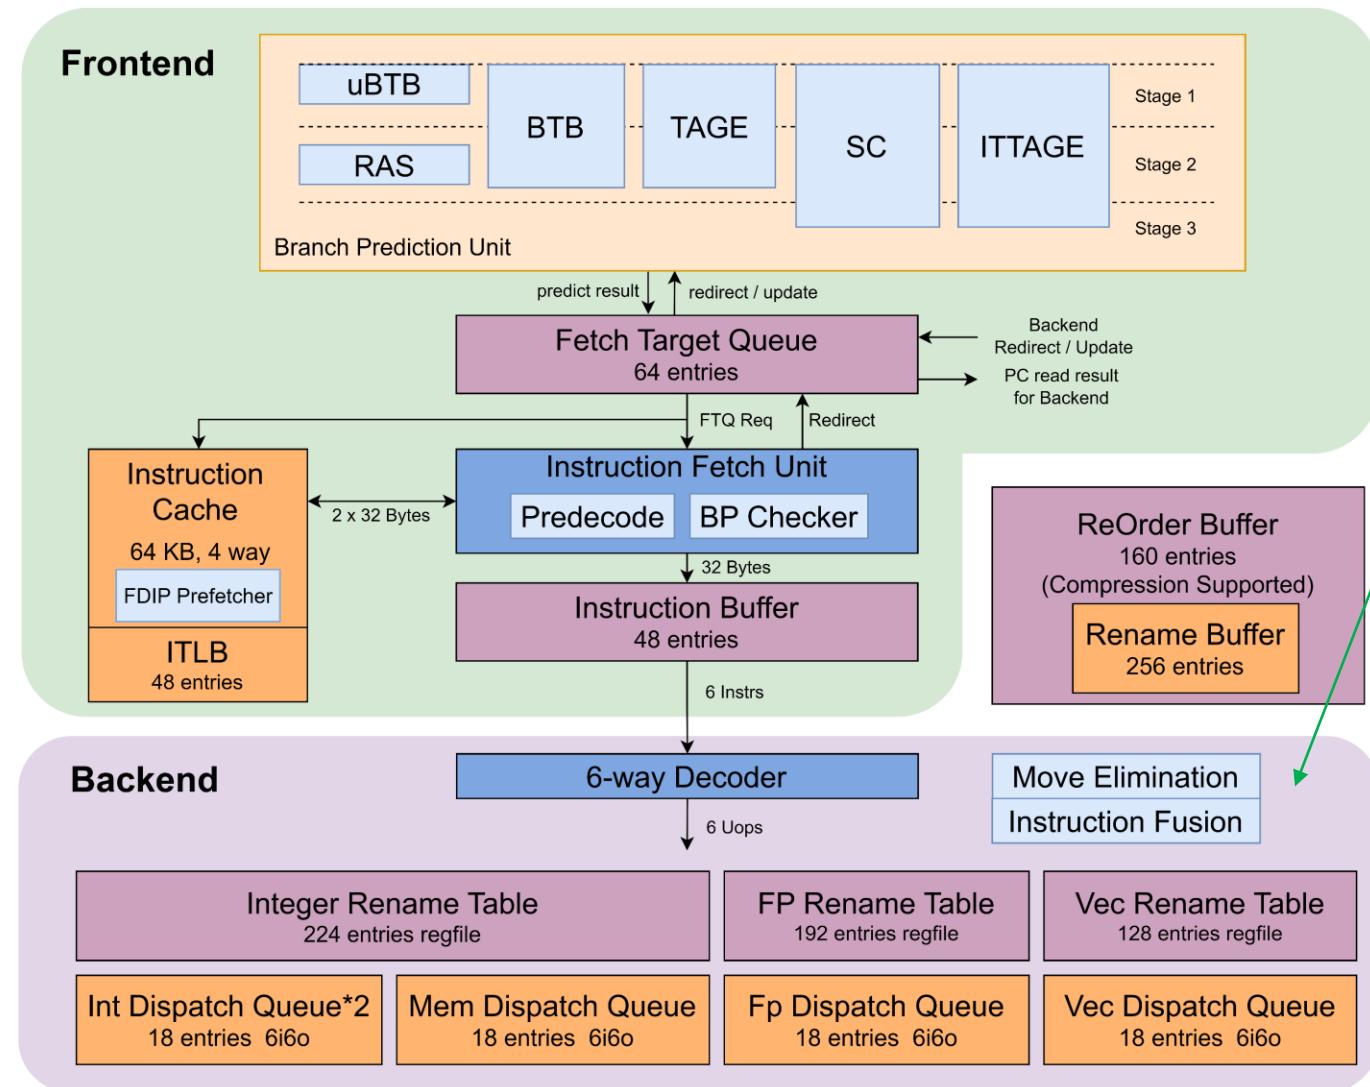
- 256-entry µBTB
- 2K-entry BTB, optional L2 BTB
- 16K TAGE-SC direction predictor
- 2K ITTAGE indirection predictor
- 48-entry return address stack

Instruction cache and TLB

- 64KB 4-way ICache
- 48-entry ITLB
- Fetch-directed instruction prefetcher



3rd Kunminghu µArch Overview



6-wide decode/ rename/ dispatch

Register rename

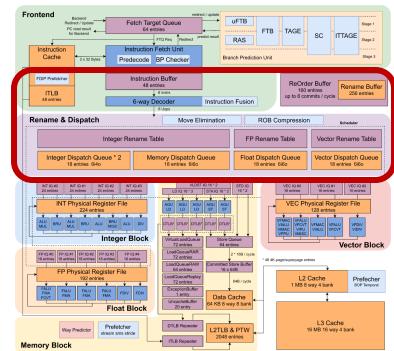
- 224-entry Int regfile
- 192-entry FP regfile
- 128-entry Vec regfile
- Move elimination
- Instruction fusion

160 entry ROB

- Support compression (up to 6-µop/entry)
- 8-entry retire/cycle
- Recovery via checkpoint + rollback

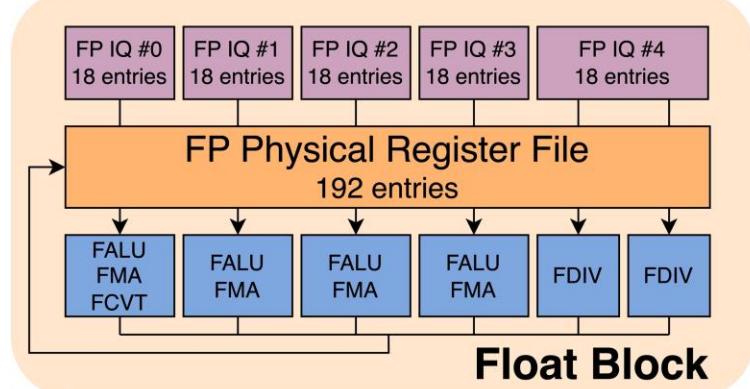
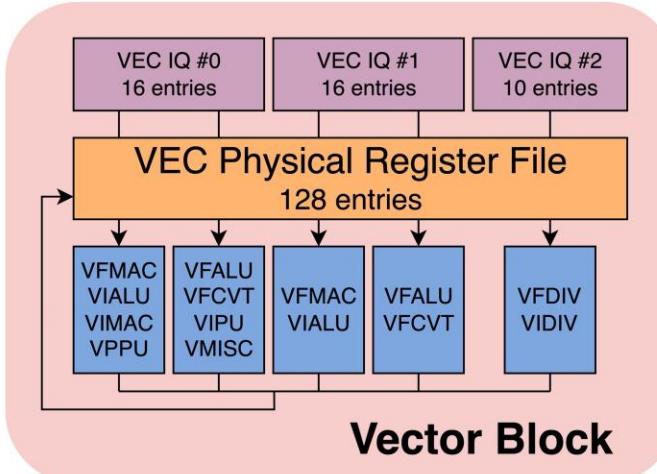
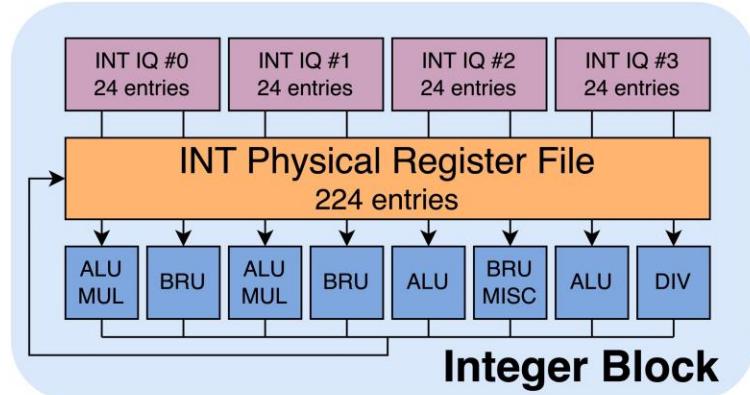
256 entry Rename Buffer

- Bridge the gap between commit & rename table update





3rd Kunminghu µArch Overview



Integer block

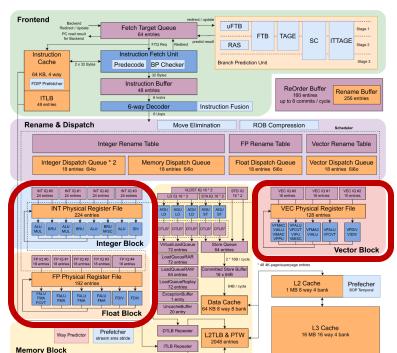
- 4 ALU (2 with 3-stage multiplier)
- 3 BRU for branch resolution
- 1 SRT radix-16 divider

Float-point block

- 4 FPU + 2 FP Div

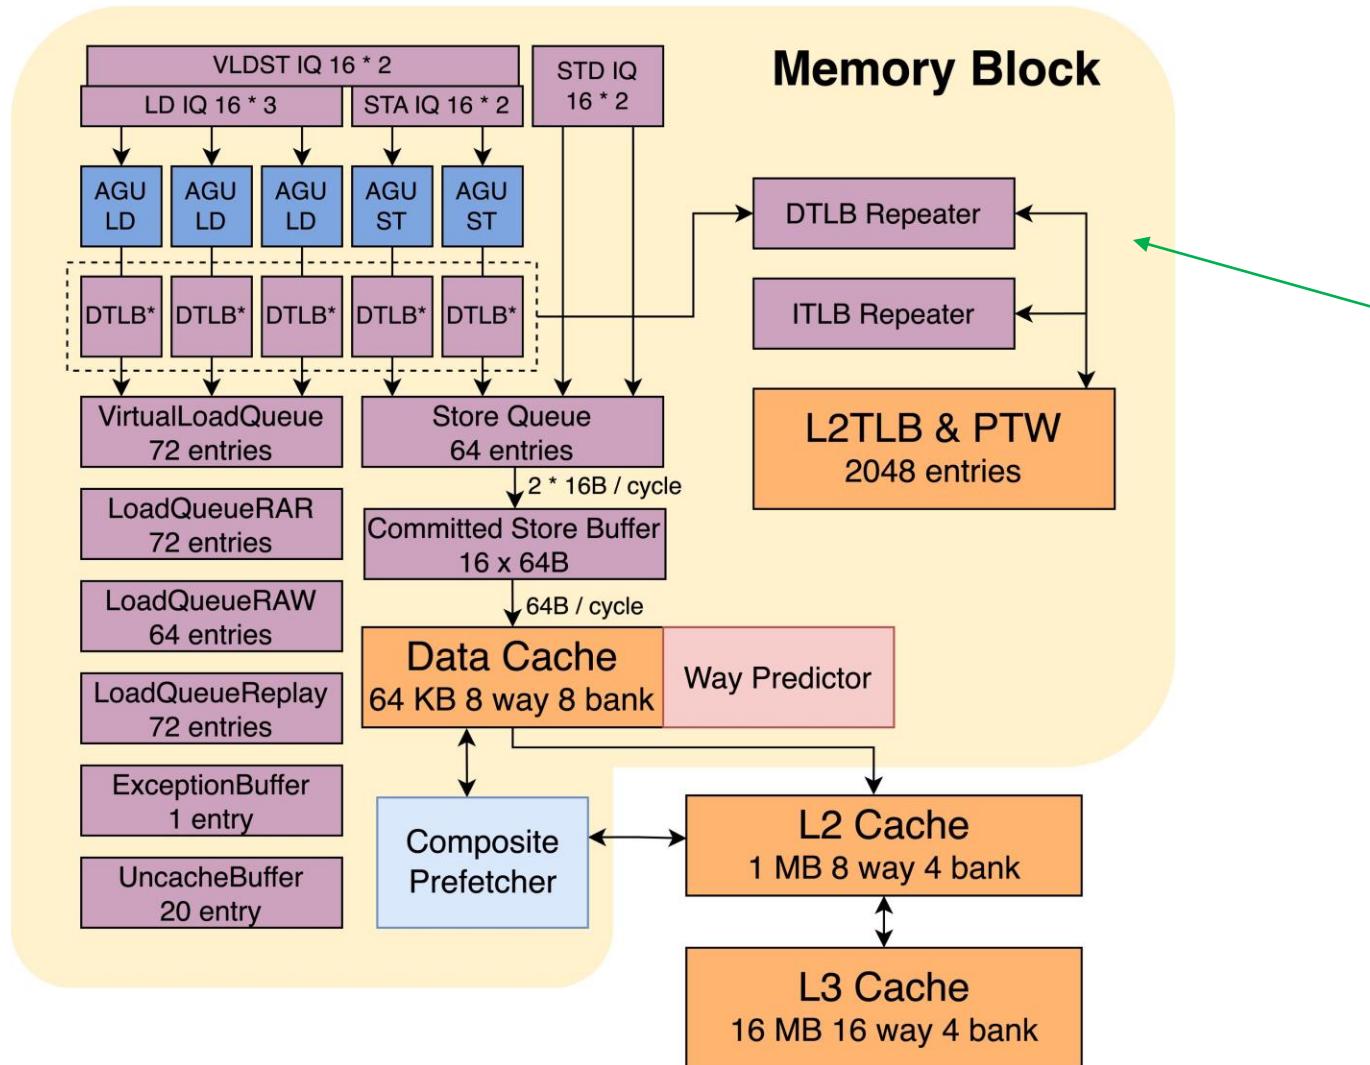
Vector block

- 4 VPU + 1 Vec Div
- **V1.0 SPEC (VLEN = 128)**





3rd Kunminghu µArch Overview



Load/Store pipeline & queue

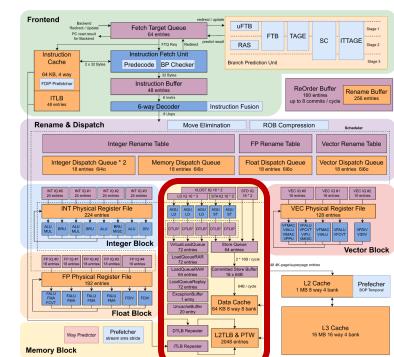
- 3 load pipes, 2 store pipes
- 72 inflight load, 64 inflight store

MMU

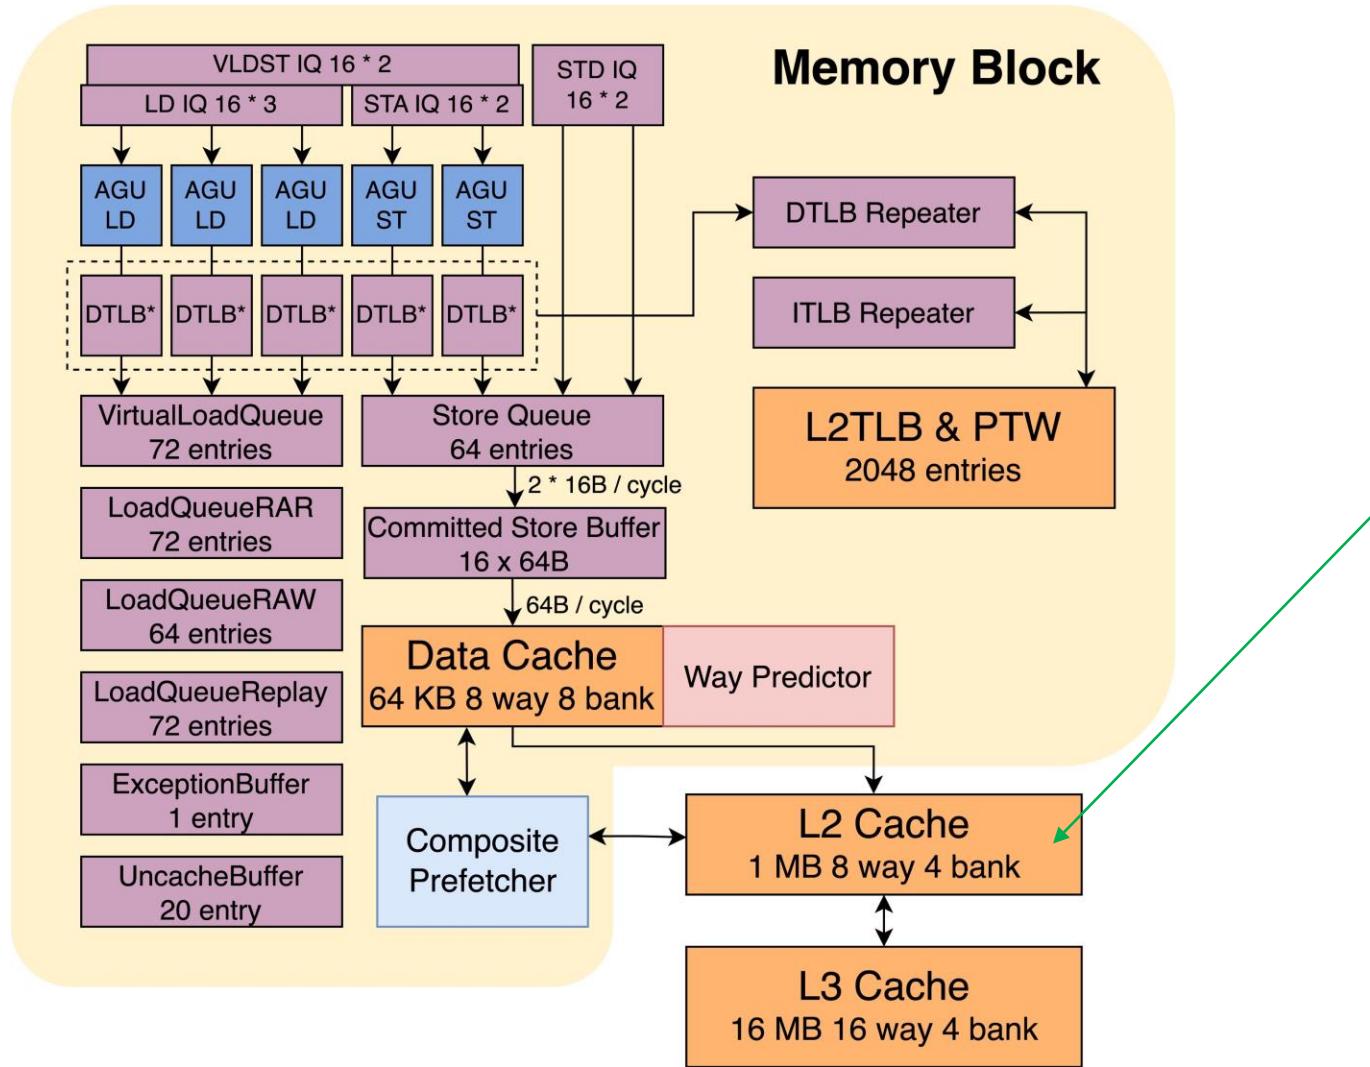
- 48b virtual/52b physical addressing
- 48-entry L1 DTLB, 2K-entry L2 TLB
- Up to 7 outstanding page table walks
- Nested walk for virtualization

Data cache

- 64KB 8-way VIPT (HW anti-aliasing)
- Way predictor for power efficiency
- Composite prefetcher
 - Stream & Stride prefetching
 - SMS & BOP prefetching
 - Temporal prefetching



3rd Kunminghu µArch Overview

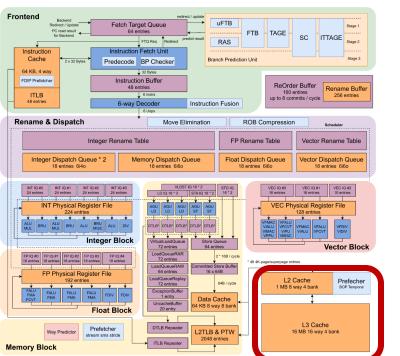


Private L2 cache

- **8-way, up to 1MB per core**
- Inclusive to D\$, non-inclusive to I\$
- **64+ OTs (configurable)**
- **10-cycle** load-use latency
- PLRU/RRIP replacement

Shared L3 cache

- **16-way, up to 16MB**
- Inclusive/Non-inclusive to L2\$
- **150+ OTs (configurable)**
- **~30-cycle** load-use latency
- PLRU/RRIP replacement





Performance Evaluation of XiangShan 3rd Kunminghu

- Method: SPEC CPU checkpoints selected by SimPoint

- Compiler: GCC 12 -O3, RV64GCB, jemalloc
- TileLink bus protocol
- Cache: 64KB I\$/D\$ + 1MB L2\$ + 16MB L3\$
- Memory modeled by DRAMsim3 DDR4@3200MHz
 - 70ns latency
 - Dual channel, 2x64

- Evaluation results

Base score	SPECint 2006	SPECfp 2006
Nanhua@2GHz	16.94	19.42
Kunminghu*@3GHz	44.17	44.52**

SPECint 2006 est.@ 3GHz	SPECfp 2006 est.@ 3GHz
400.perlbench	35.88
401.bzip2	25.55
403.gcc	46.67
429.mcf	58.13
445.gobmk	30.34
456.hammer	41.60
458.sjeng	30.50
462.libquantum	122.57
464.h264ref	56.66
471.omnetpp	39.35
473.astar	29.24
483.xalancbmk	72.01
GEOMEAN	44.17
	459.GemsFDTD
	465.tonto
	470.lbm
	481.wrf
	482.sphinx3
	GEOMEAN
	44.52

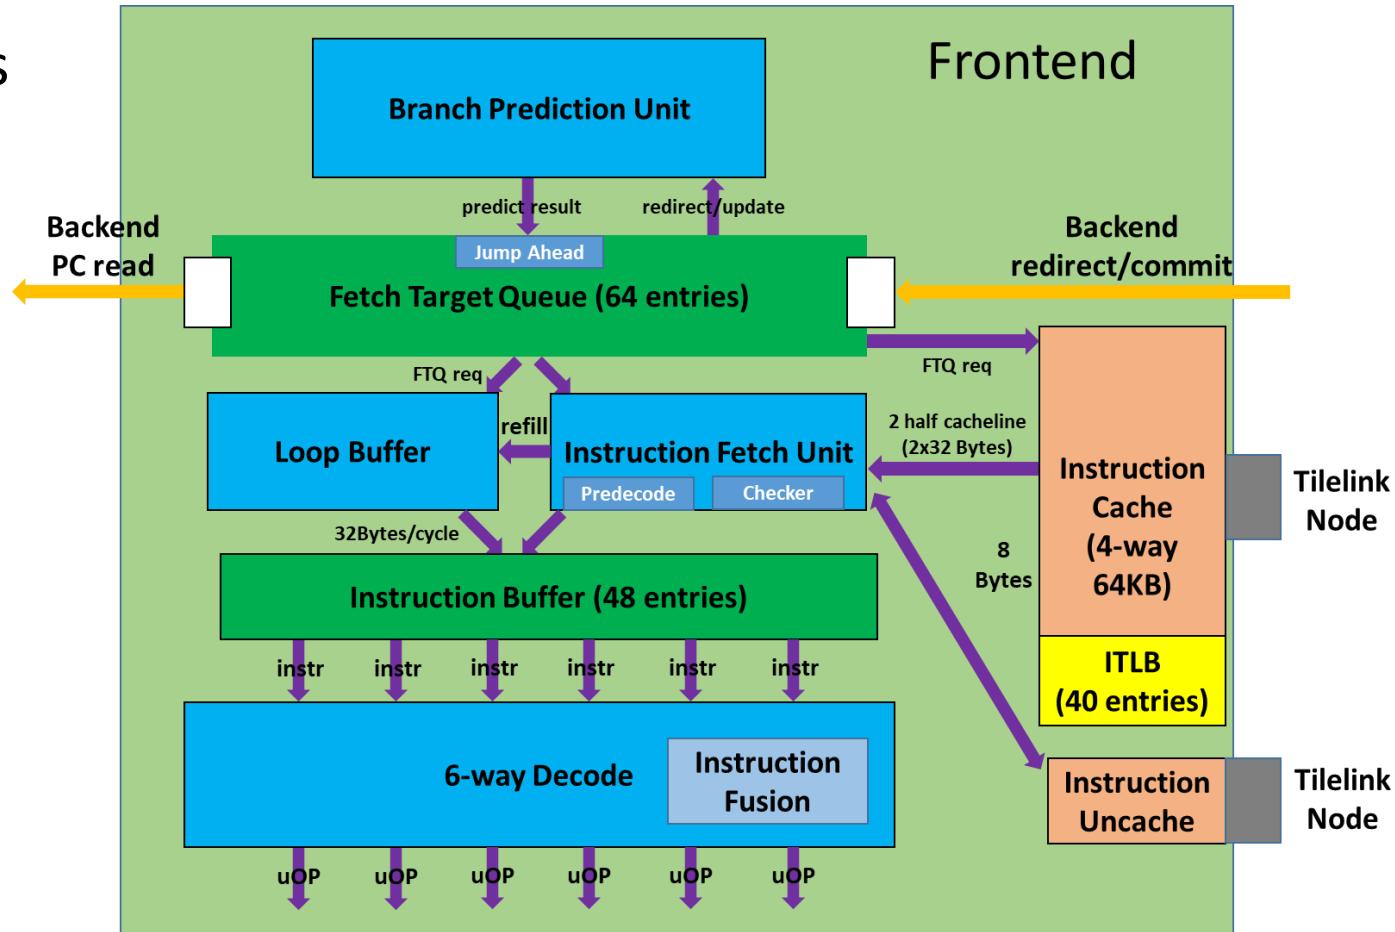
Published by XiangShan Biweekly 70

* Updated in February 2025, XiangShan commitID 7f475a2

** FP performance slightly degraded mainly due to PPA optimisation

Frontend Upgrade Overview

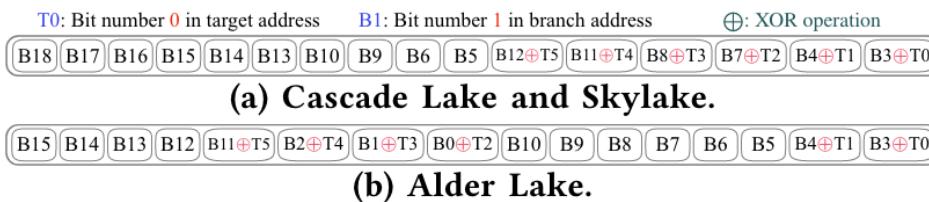
- Inherit NH decoupled architecture
 - Hiding branch prediction bubbles
 - Guiding instruction prefetch
- Iterative optimization
 - Lower prediction miss rate
 - Larger fetch/prediction width
 - More accurate branch predictor





Frontend Feature 1: TAGE Predictor Fine Tuning

- 1 Branch History Hashing



Hash algorithm affects performance

Genetic and particle swarm algorithms are employed to adjust parameters automatically.

Param of Nanhu → Param tuned: IPC +1.44%

Bonus!

- 2 #Table & History Length

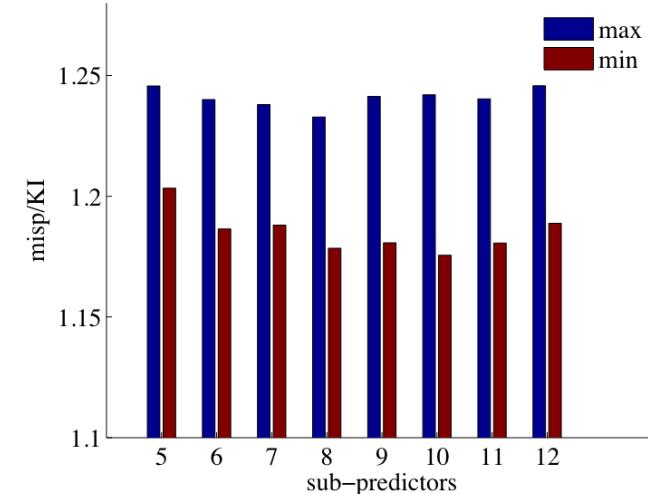


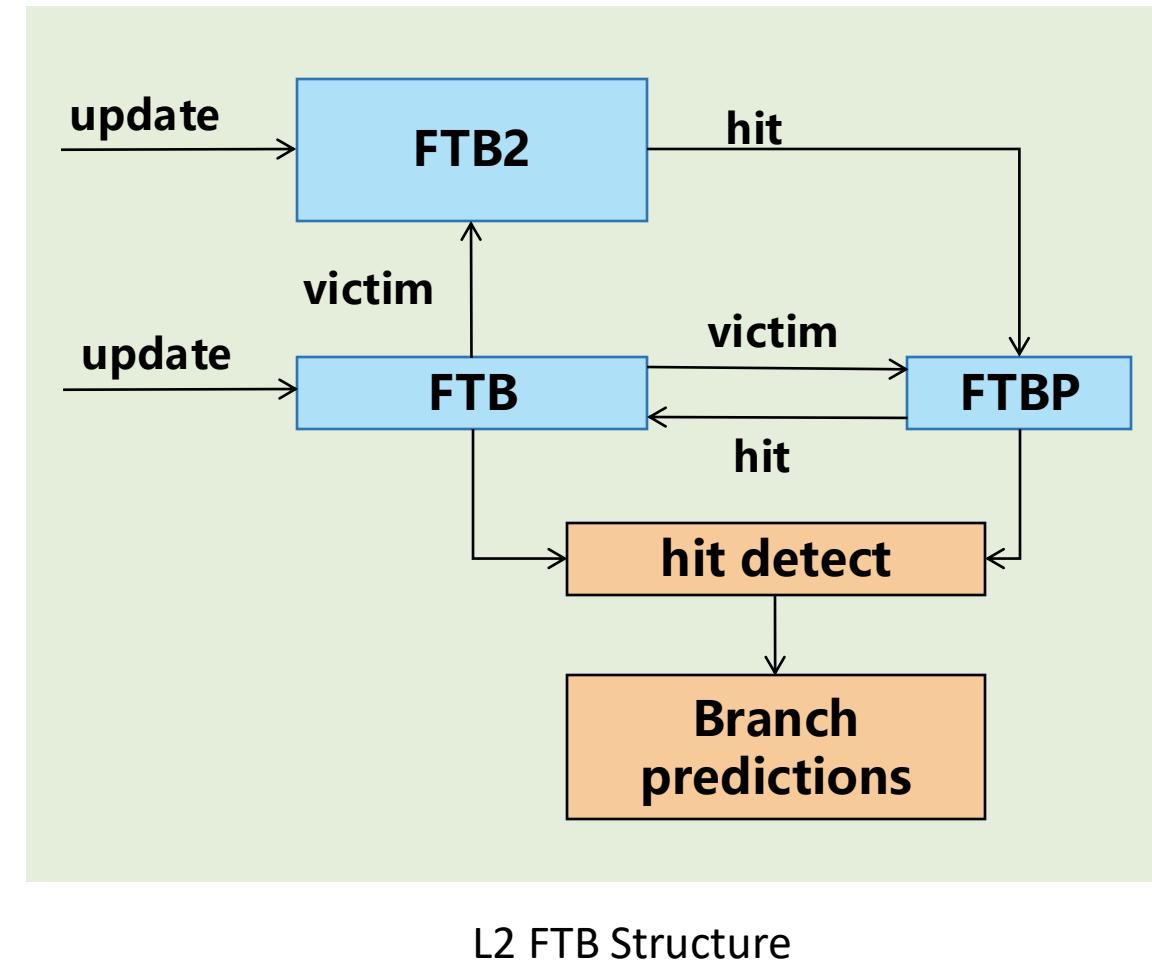
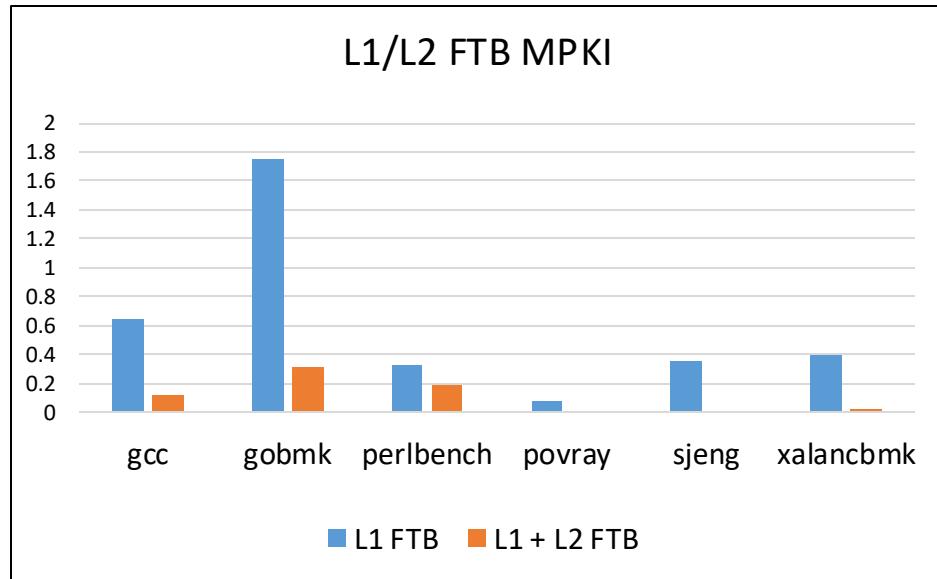
Figure 2: Performance of random parameters

Previous studies have shown that the #tables and history length have a significant effect on TAGE performance, up to 5.64% in MPKI*

* C. Zhou, L. Huang, Z. Li, T. Zhang, and Q. Dou, “Design Space Exploration of TAGE Branch Predictor with Ultra-Small RAM,” in *Proceedings of the on Great Lakes Symposium on VLSI 2017*, Banff Alberta Canada: ACM, May 2017, pp. 281–286.

Frontend Feature 2: L2 FTB (BTB)

- Hierarchical FTB
 - Support 4K L2 FTB
 - Semi-exclusive
 - FTBP acts as a buffer between L1 & L2 FTB





Frontend Feature 3: ICache Smarter Prefetcher

- Based on decoupled frontend, FDIP algorithm is implemented

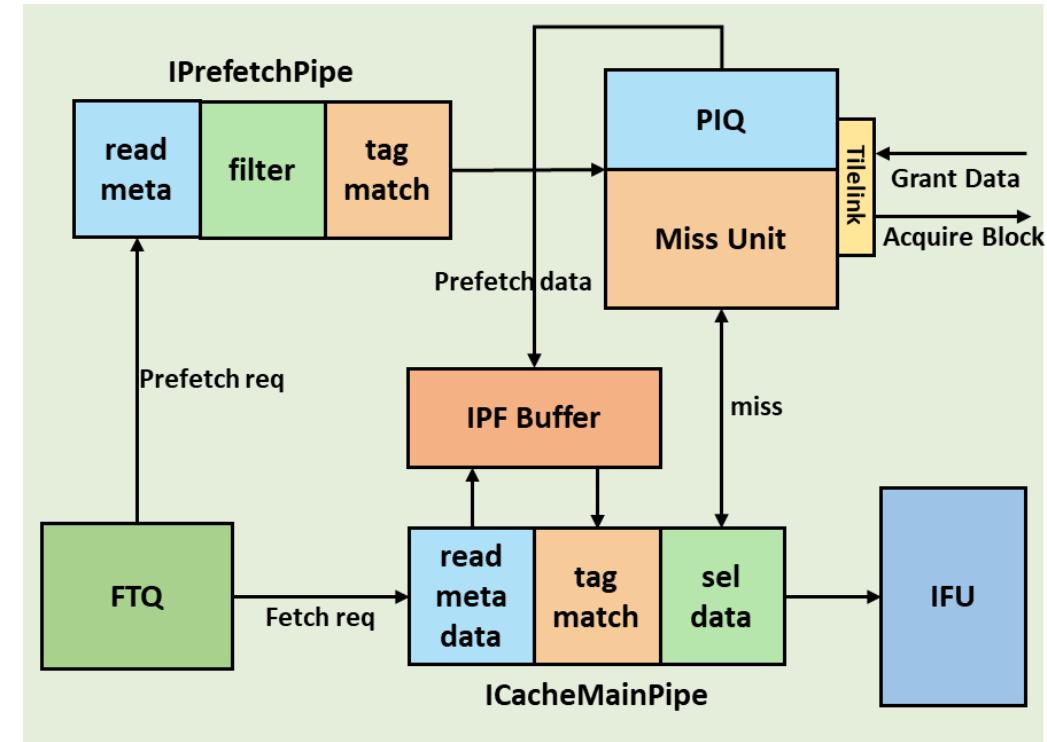
- **Prefetcher workflow**

1. Select prefetch request from PF-queue
2. Write refilled data into prefetch buffer
3. Fetch query icache & prefetch buffer parallelly

- **Optimisation compared to nanhu**

- Prefetch position: L2 Cache -> **L1 ICache**

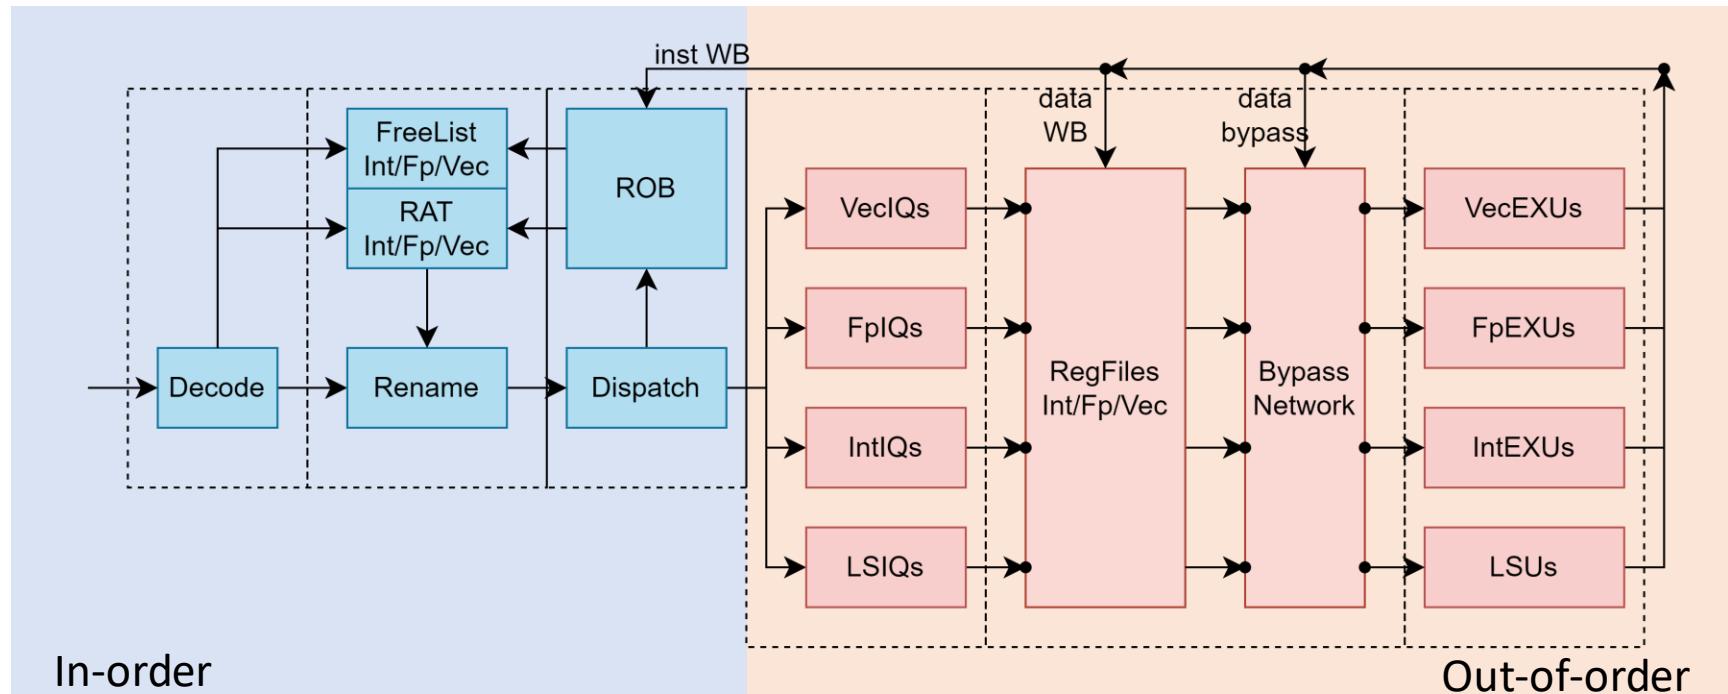
- **Dual prefetching pipelines** to enlarge bandwidth





Backend Upgrade Overview

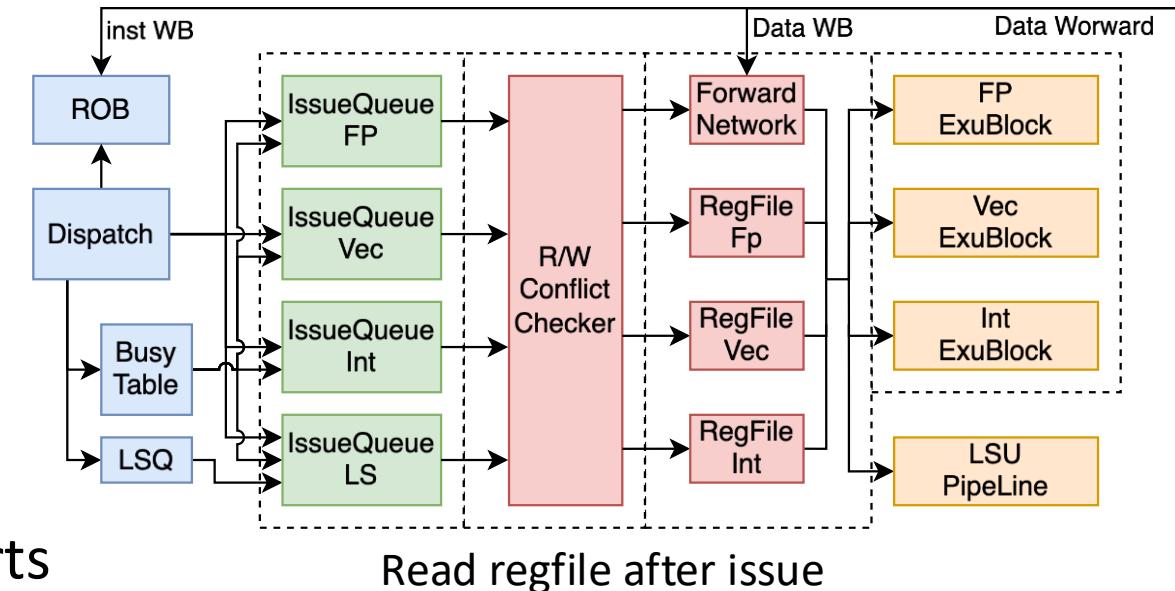
- Read regfile after issue
- Faster μop commit (ROB compression) and redirect (Rename snapshot)
- Vector extension and Hypervisor extension support





Backend Feature 1: Read Regfile after Issue

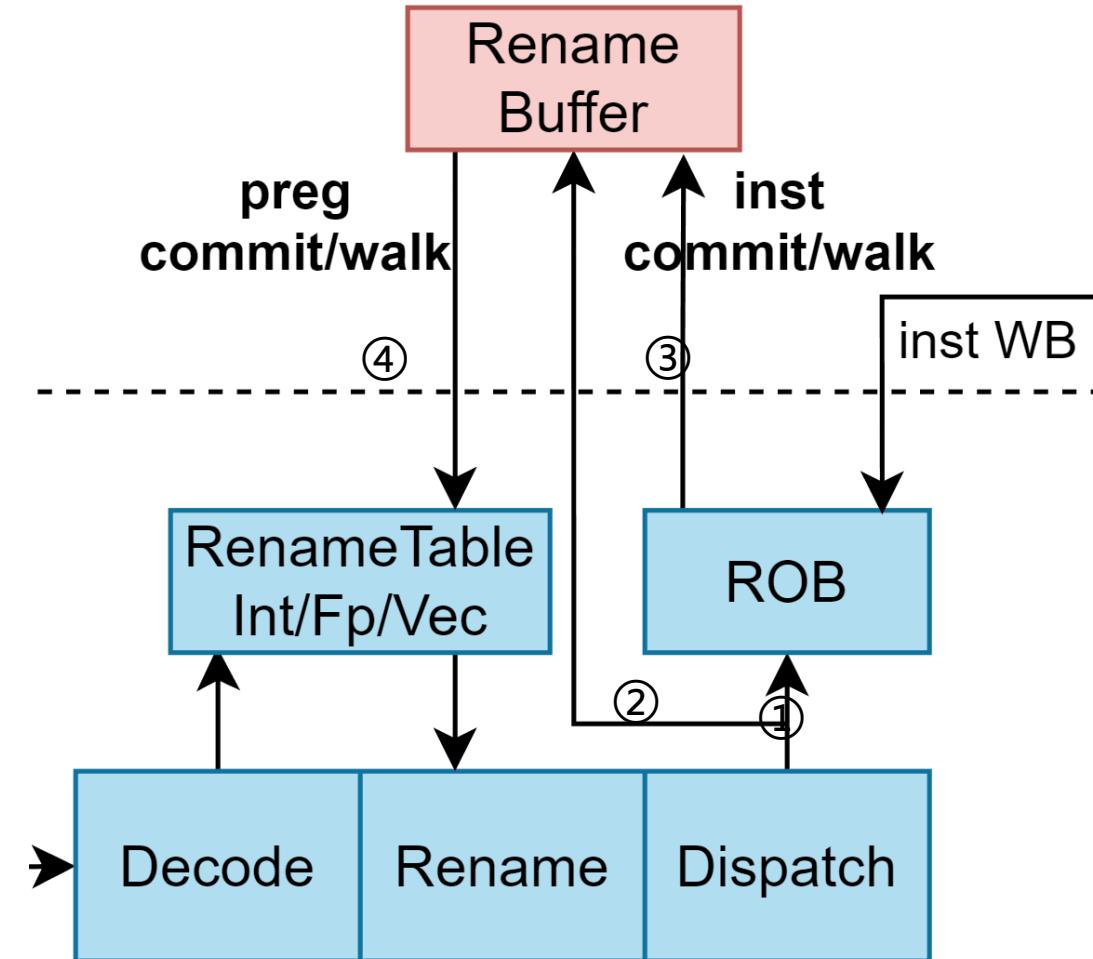
- Advantage
 - Bottleneck is shifted to after μ op issue
 - Reduce RS storage to save SRAM
 - Reduce latency and increase RS capacity
- Optimization
 - Efficient arbitration of regfile reading ports
 - Accurate speculative wakeup and cancelling





Backend Feature 2: Decouple μop & Regfile Commit

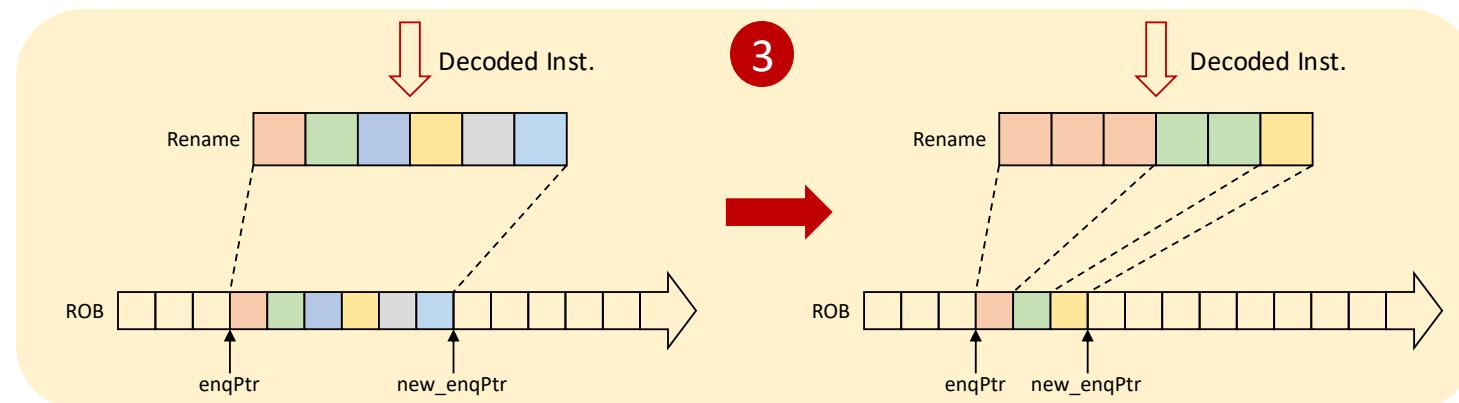
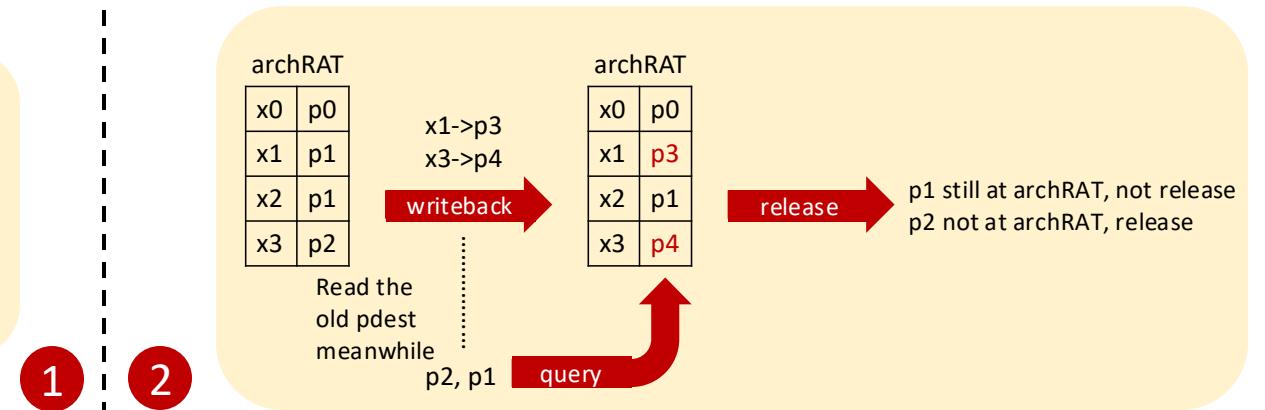
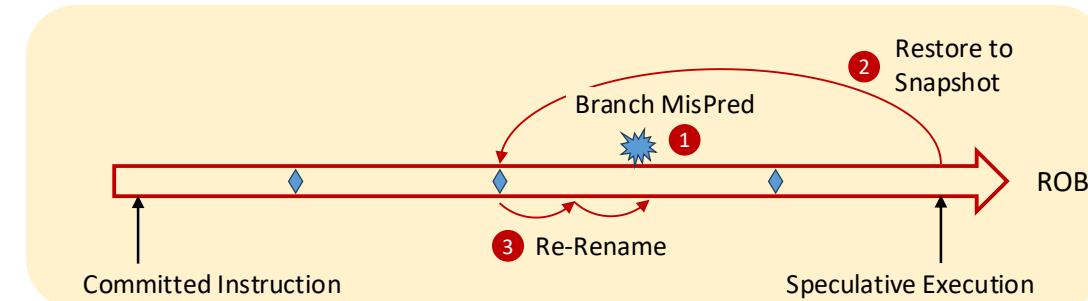
- New Rename Buffer (**RAB**) design: record regfile changing history
 - ① Save Inst in ROB after dispatch
 - ② Save regfile mapping info to Rename Buffer
 - ③ On Inst commit/redirect, wake RAB up
 - ④ RAB is responsible to update Rename Table
- Benefits
 - Faster ROB entry releasing
 - Optimize timing
 - Support μop split of vector instructions
 - Support ROB compression





Backend Feature 3: Register Renaming Optimizations

- ① Rename Snapshot: Reduce The Cost of Backend Redirect
- ② Move Elimination Based on RAT: Optimize Area and Power
- ③ ROB Compression: Reduce ROB Consuming and Enlarge OoO Windows

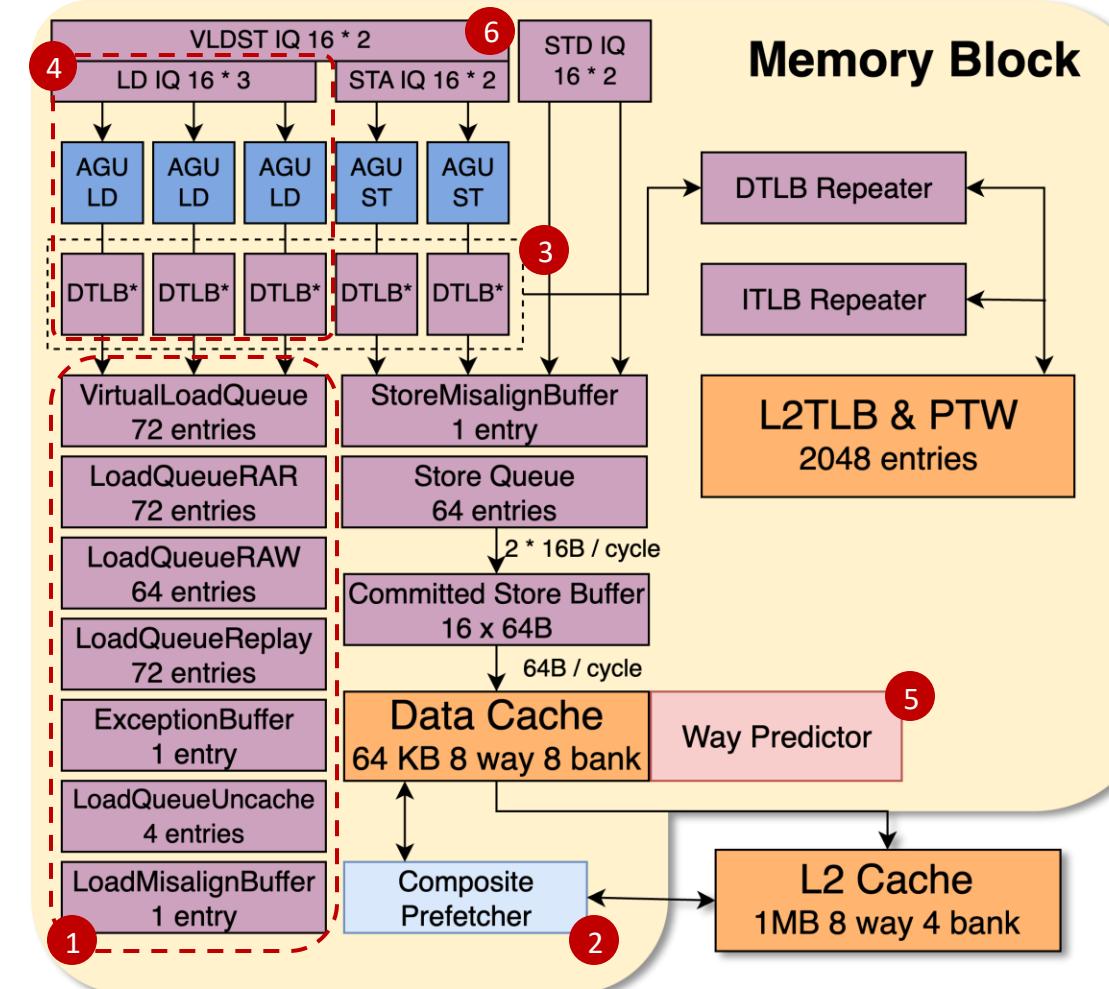




MemBlock Upgrade Overview

- Optimization compared to NH

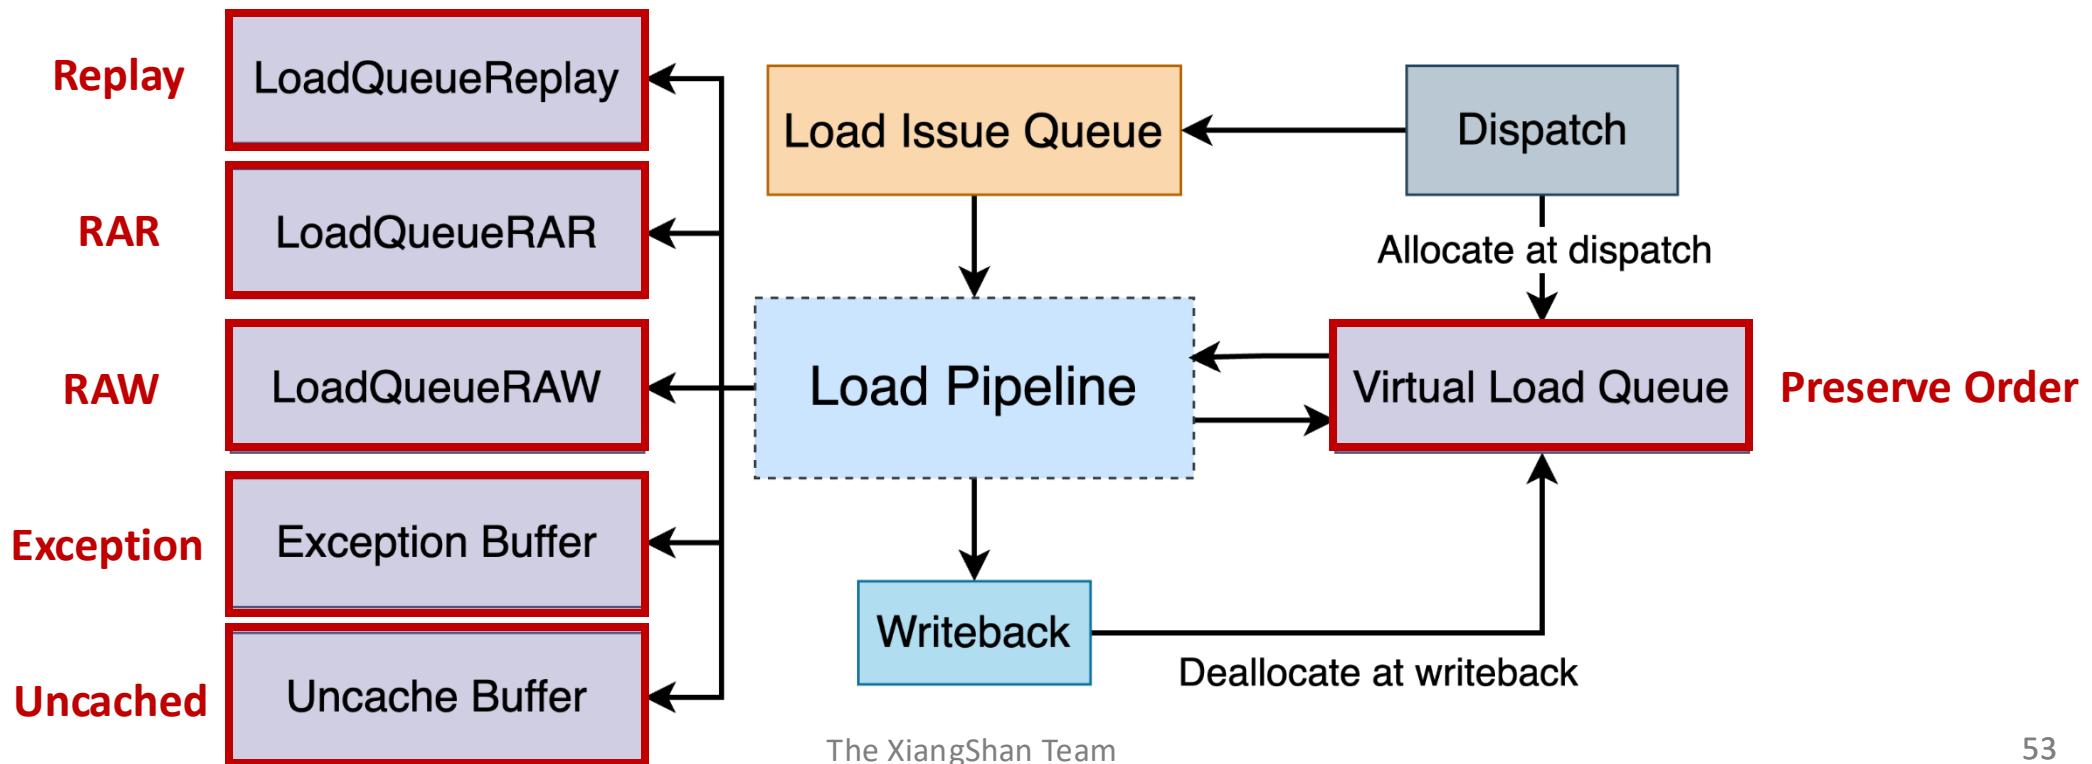
- ① Divide load queue
- ② Prefetch to L1D
- ③ Increase effective TLB capacity
- ④ Enlarge load bandwidth
- ⑤ Implement way predictor
- ⑥ Vector LSU





MemBlock Feature 1: Split Load Queue

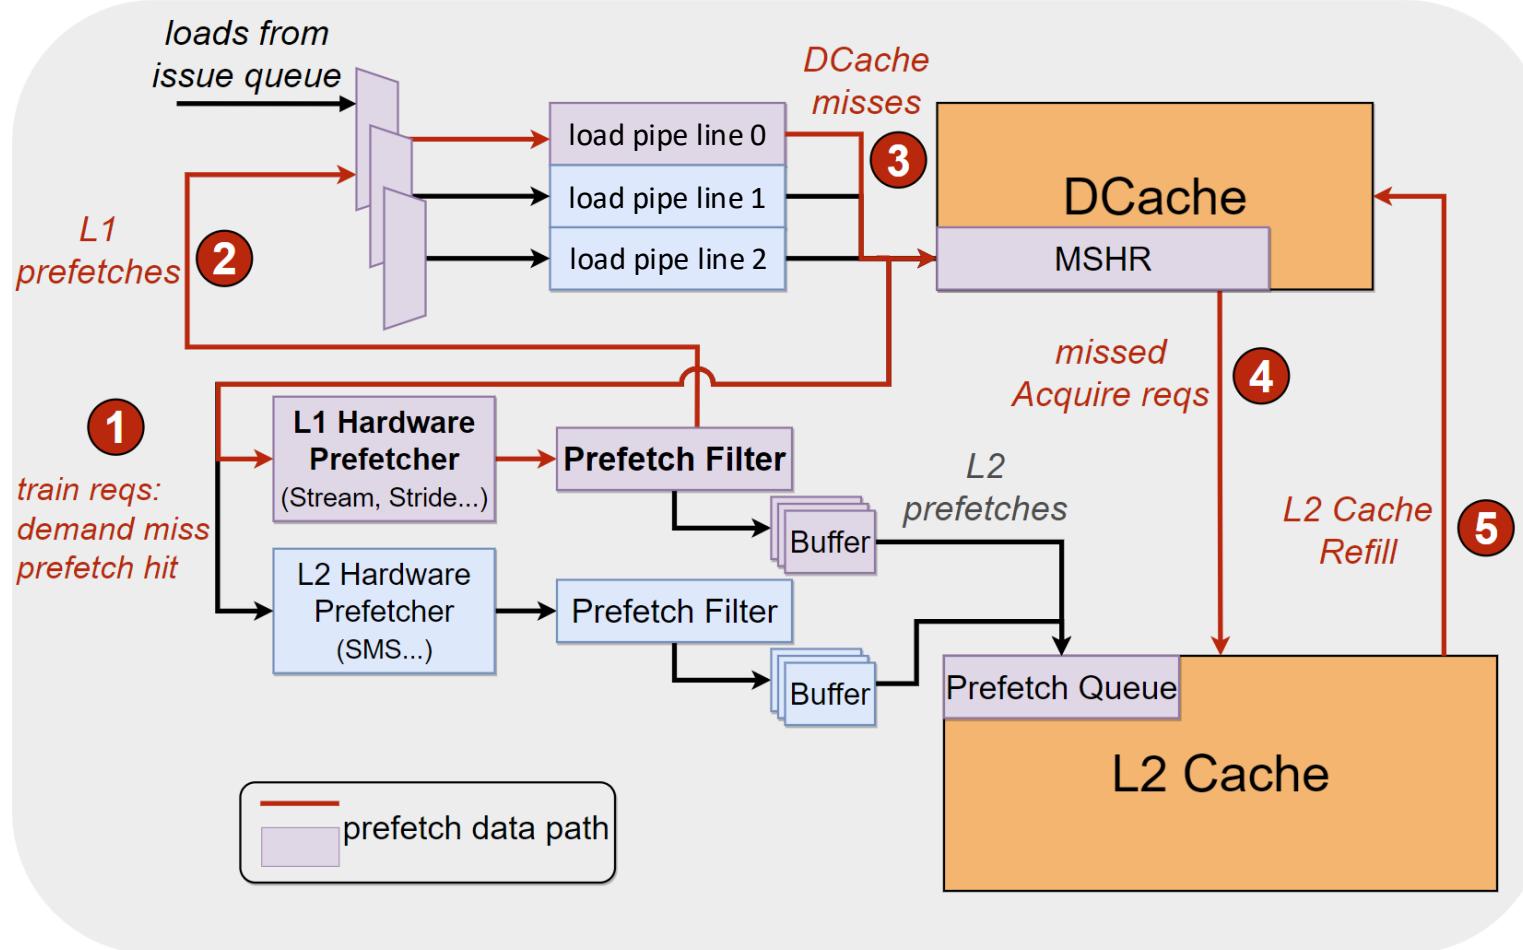
- Split load queue **based on different responsibilities**
 - Logic simpler, timing friendly and area acceptable
 - **Early commit:** load instruction can be retired sequentially from lq after write-back





MemBlock Feature 2: L1D Prefetch

- L1D prefetch workflow
 - Training pattern detection
 - **Reuse load pipelines**
 - Fetch data from D\$ MSHR
- Prefetch algorithm
 - **Stream^[1]**
 - $(x, x+1, x+2, x+3 \dots)$
 - **Stride^[2]**
 - $(x, x+n, x+2n, x+3n \dots)$



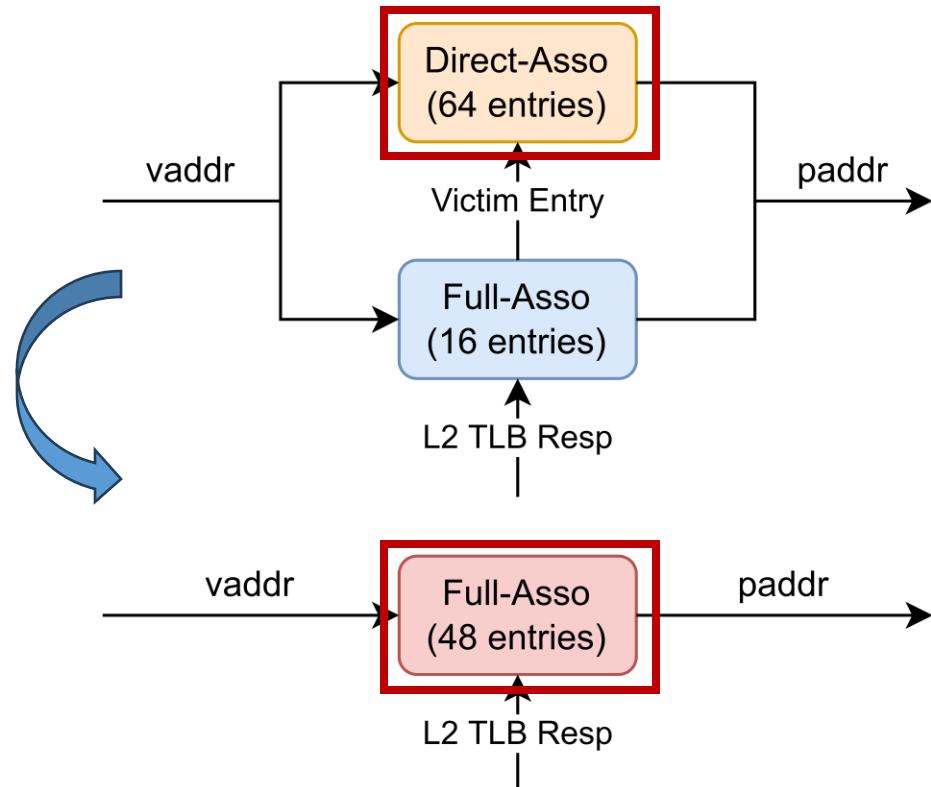
[1] S. Srinath, O. Mutlu, H. Kim and Y. N. Patt, "Feedback Directed Prefetching: Improving the Performance and Bandwidth-Efficiency of Hardware Prefetchers," 2007 IEEE 13th International Symposium on High Performance Computer Architecture, Scottsdale, AZ, USA, 2007

[2] J.-L. Baer and T.-F. Chen, "An effective on-chip preloading scheme to reduce data access penalty," Supercomputing '91:Proceedings of the 1991 ACM/IEEE Conference on Supercomputing, Albuquerque, NM, USA, 1991

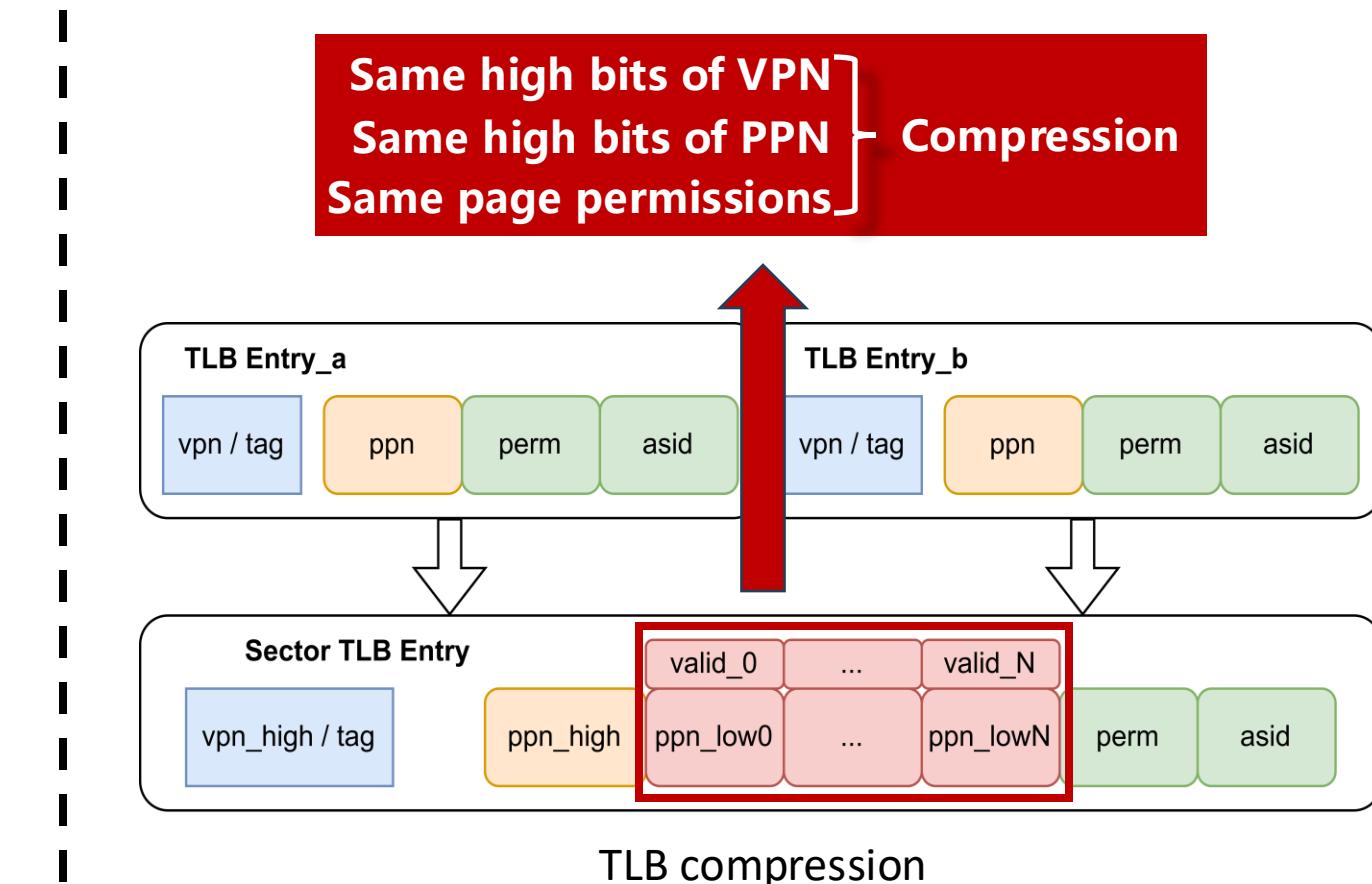


MemBlock Feature 3: Increase effective TLB capacity

- 16 fully associative entry + 64 direct mapping entry → **48 fully associative entry**
- Support **TLB compression**, merging contiguous page table entries

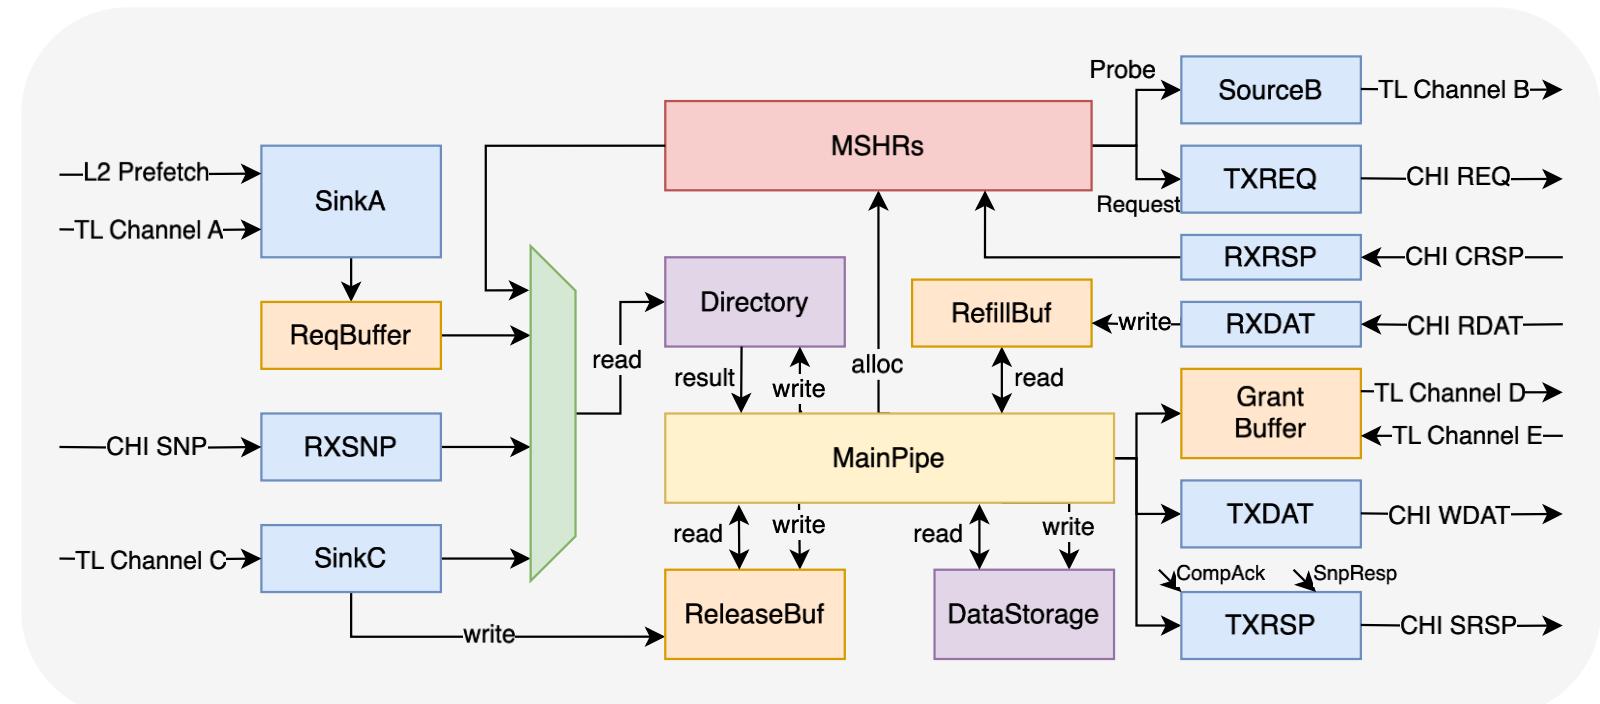


DTLB organization in Nanhu / Kunminghu



L2 Cache Upgrade Overview

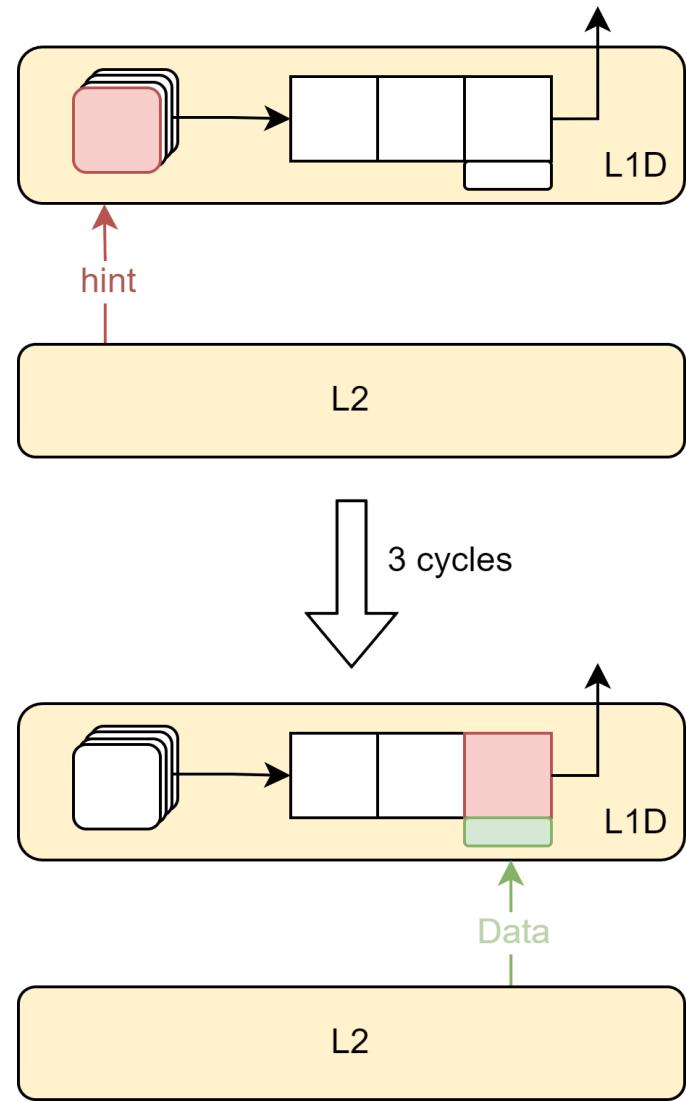
- Multiple concurrent pipelines → Non-blocking main pipeline
- L1-L2 Refill wake-up collaboratively
- Eliminate transaction serialization in the same set & Request merge
- Supported protocols
 - CHI Issue E.b
 - CHI Issue B
 - CHI Issue C
 - TileLink





L2 Cache Feature 1: L1-L2 Collaborative Optimization

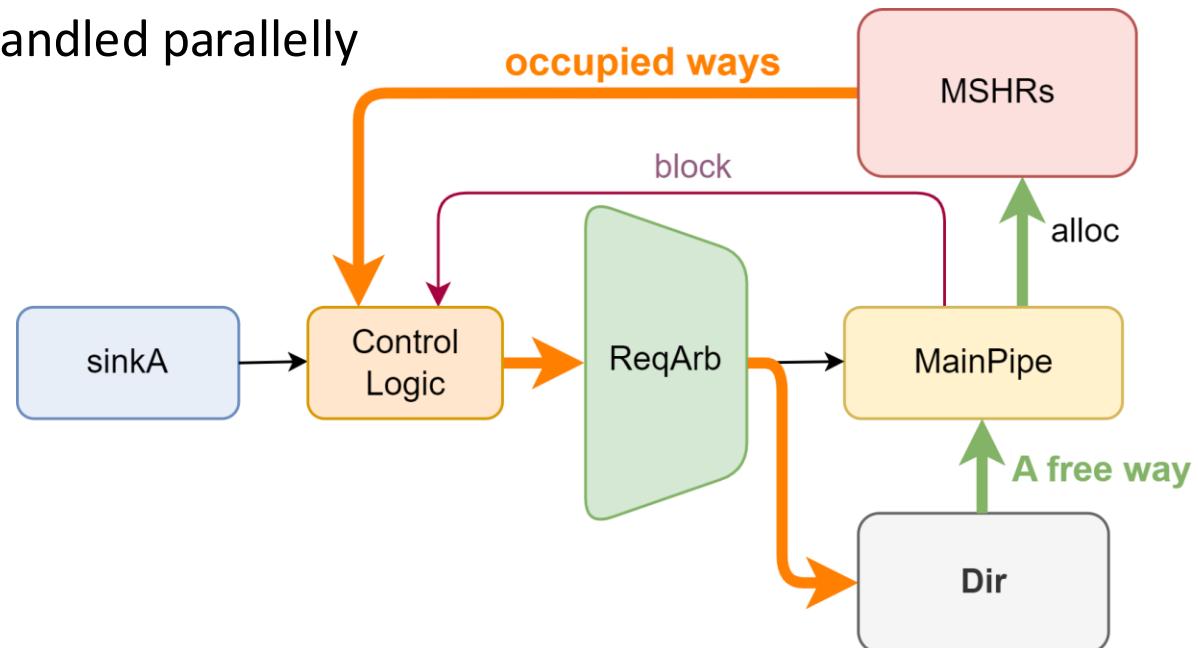
- Observation:
 - Interval between **load wake-up** and **use-refill-data** is 3 cycles
 - There's chance to hint LSU to replay before refill
- New Design:
 - Let L2 give **Hint signal in advance**
 - Speedup the wakeup and replay of load replay queue
- Implementation:
 - Set up monitor to track info from main pipeline
 - Need to calculate the exact Hint timing
 - Send hint signal to L1 3 cycles before refill





L2 Cache Feature 2: Eliminate Txns Serialization in same set

- Background: Txns to **same sets** were handled serially, reduces parallelism
- Difficulty: Same-set txns can affect each other, due to replacement
- New Design
 - Txns of **same-set but different-addr** can be handled parallelly
 - Txns of same-addr are handled serially
- Implementation
 - Record occupied ways in active MSHRs
 - Assign a free way** to the new request



L2 Cache Feature 3: Evict on Refill

- **Observation:**

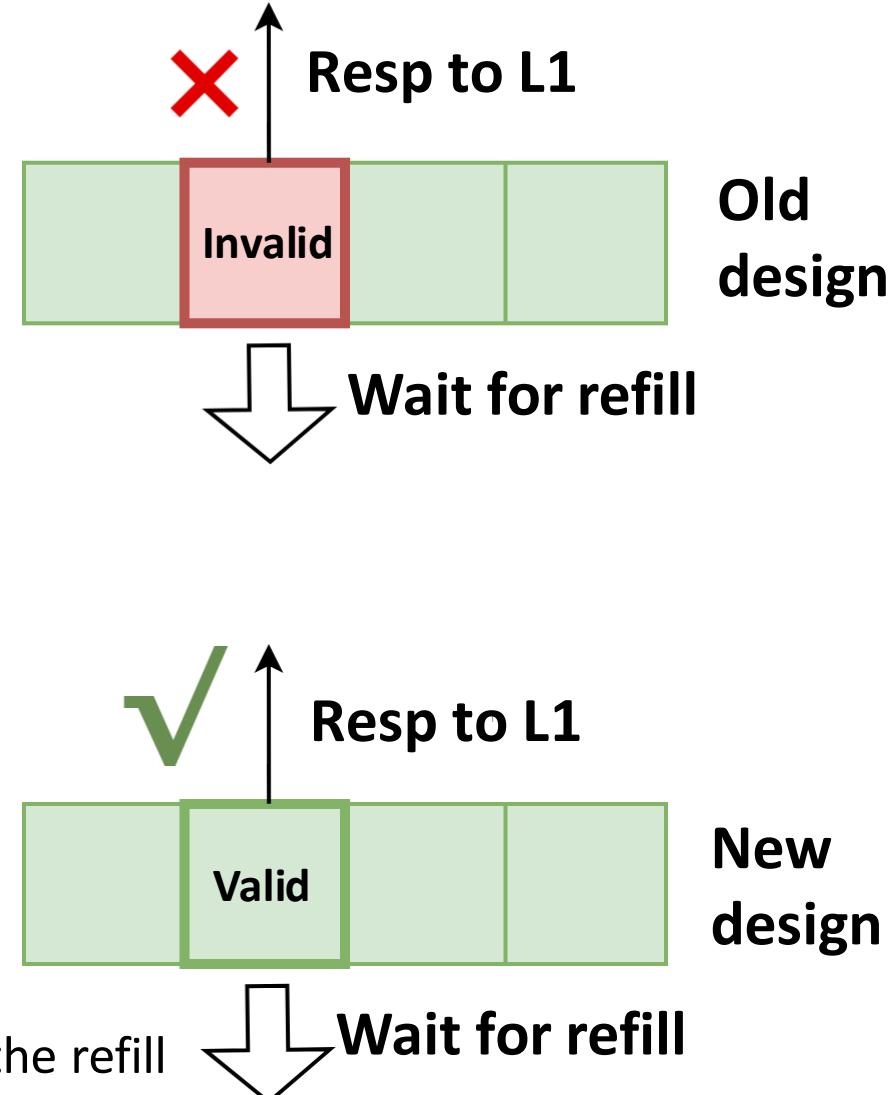
- On L1 acquire, L2 chooses a way for replacement
- But this way is **occupied until refill**, unable to serve L1

- **Improvement:**

- New design chooses replaced-way **after** data is refilled from L3
- So, it can still serve L1 when waiting for refill

- **Implementation:**

- Transfer it to L3 on acquire miss, make replacer untouched
- Wait for L3 to refill
- **Read directory again** and assign one way when MSHR handles the refill





Beyond RTL Design: Agile development & infrastructure

- Implement by Chisel HDL

- High readability
- High modifiability



Learn more from <https://www.chisel-lang.org/>

- we designed some open-source high-performance utilities / components

BinaryArbiterNode.scala

BitUtils.scala

ChiselDB.scala

CircularQueuePtr.scala

Constantin.scala

DataModuleTemplate.scala

ECC.scala

ExcitingUtils.scala

ExtractVerilogModules.scala

FastArbiter.scala

FileRegisters.scala

GTimer.scala

Hold.scala

IntBuffer.scala

LFSR64.scala

LatencyPipe.scala

LookupTree.scala

MIMOQueue.scala

Misc.scala

ParallelMux.scala

Pipeline.scala

PipelineConnect.scala

PriorityMuxDefault.scala

PriorityMuxGen.scala

RegMap.scala

Replacement.scala

ResetGen.scala

SRAMTemplate.scala

StopWatch.scala

TLClientsMerger.scala

TLEdgeBuffer.scala

Check it out! <https://github.com/OpenXiangShan/Utility>



Beyond RTL Design: Agile development & infrastructure

- Highly configurable
 - setup parameters as you wish

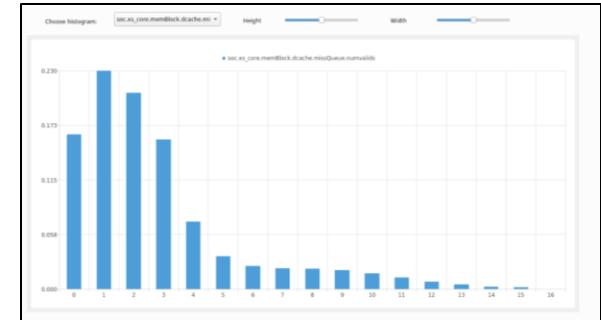
```
IBufSize: Int = 48,  
DecodeWidth: Int = 6,  
RenameWidth: Int = 6,  
CommitWidth: Int = 6,  
FtqSize: Int = 64,  
EnableLoadFastWakeUp: Boolean = true,  
IssQueSize: Int = 16,  
NRPhyRegs: Int = 192,  
LoadQueueSize: Int = 80,  
LoadQueueNWriteBanks: Int = 8,  
StoreQueueSize: Int = 64,  
StoreQueueNWriteBanks: Int = 8,  
RobSize: Int = 256,  
dpParams: DispatchParameters =  
  DispatchParameters(  
    IntDqSize = 16,  
    FpDqSize = 16,  
    LsDqSize = 16,  
    IntDqDeqWidth = 4,  
    FpDqDeqWidth = 4,  
    LsDqDeqWidth = 4  
)  
  
  exuParameters: ExuParameters =  
    ExuParameters(  
      JmpCnt = 1,  
      AluCnt = 4,  
      MulCnt = 0,  
      MduCnt = 2,  
      FmacCnt = 4,  
      FmiscCnt = 2,  
      FmiscDivSqrtCnt = 0,  
      LduCnt = 2,  
      StuCnt = 2  
)  
  
  prefetcher: Option[PrefetcherParams] =  
    Some(SMSParams()),  
  LoadPipelineWidth: Int = 2,  
  StorePipelineWidth: Int = 2,  
  StoreBufferSize: Int = 16,  
  StoreBufferThreshold: Int = 7,  
  EnableLoadToLoadForward: Boolean = true,  
  EnableFastForward: Boolean = false,  
  EnableLdVioCheckAfterReset: Boolean = true,  
  EnableSoftPrefetchAfterReset: Boolean = true,  
  EnableCacheErrorAfterReset: Boolean = true,  
  EnablePTWPreferCache: Boolean = true,  
  EnableAccurateLoadError: Boolean = true,
```

Some key parameters

- Easy-to-use performance counter

- Multiple collection types

- Cumulative counter
- Delay counter
- Transaction counter
- Queue counter



- Multiple analysis style

- Key-value list
- Histogram
- Database
- Time series visualization





XiangShan: Open-Source High-Performance Processor

XiangShan Project

- Embraces innovative agile development workflow
- Fills the gap in open-source high-performance processors
- Addresses the needs from both industry and academia

- 1st generation: YQH (Yanqihu) & 2nd generation: NH (Nanhu)
- 3rd generation: KMH (Kunminghu)
 - Support RVA23 profile and L2 Cache with CHI protocol
 - 3GHz @7nm, estimated SPEC CPU2006 45@3GHz
- Open-sourced at <https://github.com/OpenXiangShan/XiangShan>
- **NEXT PART: Hands-on Development**
 - Scan the QR to get the env introduction and slides



Thanks!