

XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

The XiangShan Team

APPT'25@Athens, Greece

July 14, 2025


What we will cover in this tutorial

- Introduction of XiangShan (9:45 – 10:20, 35 minutes)
 - Project Overview and Research Infrastructure
- Introduction of XS-Gem5 (10:20 – 11:00, 40 minutes)
 - Calibration and Development Tools
- **Development Tools Live Demonstration (11:15 – 12:30, 75 minutes)**
 - Introduction of their usages – How to develop on XiangShan with MinJie
 - Simulation, Debugging, and Performance Evaluation
 - Research Demo



APPT 2025


Processor Development Workflow



XiangShan provides a development platform to speed up the workflow.



MinJie: Infrastructures for developing chips



* Xu Y, Yu Z, Tang D, et al. Towards developing high performance RISC-V processors using agile methodology, 2022 55th MICRO.



In this tutorial, we will

- Go through the development workflows including:

Simulation

- RTL Generation
- RTL Simulation
-

Func. Verification

- Nexus-AM
- NEMU
- Dfftest
- TL-Test
-

Perf. Verification

- Checkpoint
- Gem5
- Top-Down
- Constantin
-

- ***Scan the QR code to download slides and instructions.***
 - or access <https://tutorial.xiangshan.cc/appt25/>
 - XiangShan tutorial <https://github.com/OpenXiangShan/xs-env>



Environment Preparation

- **OS:** Ubuntu 22.04 LTS
- **CPU:** Not limited but determines the compilation speed
- **Memory:** 32G or more
- **Disk:** 20G or more
- **Network:** Need a smooth network environment



APPT 2025

- **Setup Instructions:** <https://tutorial.xiangshan.cc/appt25/StartUp/>
- Instructions include:
 - (1) riscv-gnu-toolchain/Verilator/...
 - (2) download XiangShan project
 - (3) download tutorial materials



Chisel Compilation



- Compile RTL and build simulator with Verilator

```
$ make emu -j4 # After environment preparation
```

Options:

<i>CONFIG=MinimalConfig</i>	<i>Configuration of XiangShan</i>
<i>// EMU_THREADS=4</i>	<i>Simulation threads</i>
<i>// EMU_TRACE=1</i>	<i>Enable waveform dump</i>
<i>// WITH_DRAMSIM=1</i>	<i>Enable DRAMSim3 for DRAM simulation</i>
<i>// WITH_CHISELDB = 1</i>	<i>Enable ChiselDB feature</i>
<i>// WITH_CONSTANTIN = 1</i>	<i>Enable Constantin feature</i>

- Compilation might take ~**20** mins

Run RTL Simulation with Verilator

- We have prepared a pre-built XiangShan simulator

run pre-built simulator

```
$ cd $XS_PROJECT_ROOT/tutorial/p1-basic-func
```

```
$ ./emu -i hello.bin --no-diff 2>hello.err
```

Some key options:

<code>-i</code>	<i># Workload to run</i>
<code>-C / -I</code>	<i># Max cycles /Max Insts</i>
<code>--diff=PATH / --no-diff</code>	<i># Path of Reference Model / disable difftest</i>

Great!


We have learned the basic simulation process of Xiangshan.

Once we get the RTL simulator ready, what's next?

How can we

- **generate desired workloads**
- **find and fix functional bugs**
- **do performance analysis**
- **do research on micro-architecture**

MinJie Functional Verification



Nexus-AM: Bare Metal Runtime

- **Purpose**
 - Generate workloads **agilely** without an OS
 - Provide runtime framework for **bare metal machines** like XiangShan
- **We present a bare metal runtime called Nexus-AM**
 - Light-weight, easy to use
 - Implement basic **system call** interfaces and exception handlers
 - Support multiple **ISAs and configurations**




Nexus-AM

- **Hands-on:** build Coremark workload with Nexus-AM
- **Showcase**

```
$ bash build-am.sh
```

```
# cd $AM_HOME/apps/coremark
# make ARCH=riscv64-xs \
    ITERATIONS=1 LINUX_GNU_TOOLCHAIN=1 TOTAL_DATA_SIZE=400
# ls -l build
#   coremark-riscv64-xs.bin    Binary image of the program
#   coremark-riscv64-xs.elf    ELF of the program
#   coremark-riscv64-xs.txt    Disassembly of the program
```

MinJie Functional Verification





NEMU: ISA REF

- **Purpose**

- Golden model for verification
- Simple yet high performance

- **We present NEMU**

- An **instruction set simulator**, similar to Spike
- Through **optimization techniques**, achieve performance similar to QEMU.
- **A set of APIs** is provided to assist XiangShan in comparing and verifying the **architectural states**

- Showcase

```
$ bash build-nemu.sh
```

```
# cd $NEMU_HOME
```

```
# make clean
```

```
# make riscv64-xs_defconfig
```

```
# make -j build NEMU as the bare metal machine, can run the Coremark from the previous step
```

```
# make clean-softfloat
```

```
# make riscv64-xs-ref_defconfig
```

```
# make -j build NEMU as the reference model for XiangShan
```

- **Hands-on:** run Coremark workload on NEMU
- **Showcase**

```
$ bash run-nemu.sh
```

```
# cd $NEMU_HOME run in batch mode, faster  
# ./build/riscv64-nemu-interpretter -b \ set workspace to Coremark  
$AM_HOME/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin
```




DiffTest: ISA Co-simulation Framework

- **Basic flows**


- Processor instructions commit/other states update
- The simulator executes the same instructions
- Compare the architectural states between DUT and REF
- Abort or continue

- **Provided as verification infrastructure for processors**

- APIs for HDLs such as Chisel/Verilog
- RTL simulators such as Verilator, VCS
- RISC-V Instruction Set Simulators, such as **NEMU**, **Spike**

- **SMP-DiffTest: co-simulation on an SMP processor**

- SMP Linux kernel and multi-threading programs
- Online checking of cache coherency and memory consistency



Basic architecture of DiffTest

```
while (1) {  
    icnt = cpu_step();  
    nemu_step(icnt);  
    r1s = cpu_getregs();  
    r2s = nemu_getregs();  
    if (r1s != r2s) { abort(); }  
}
```

Online checking

Difftest

- **Hands-on:** run Coremark workload on XiangShan, difftest with NEMU
- **Showcase**

Running Coremark takes about 3 minutes

```
$ bash run-emu.sh
```

```
# ./emu -i \ use coremark.bin
```

```
$AM_HOME/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin
```

```
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \ difftest with NEMU
```

```
2>perf.out \ redirect stderr to file
```

```
./emu -i $AM_HOME/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin --diff $NEMU_HOME/build/riscv64-nemu-interpretor-so 2>perf.err
emu compiled at Jun 17 2025, 15:53:52
[WARNING] /home/guest/tanghaojin/XiangShan/build/constantin.txt does not exist, so all constants default to initialized values.
Using simulated 32768B flash
Core 0's Commit SHA is: c8bbd5c6ac, dirty: 0
Using simulated 8386560MB RAM
The image is /home/guest/tanghaojin/nexus-am/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin
The reference model is /home/guest/tanghaojin/NEMU/build/riscv64-nemu-interpretor-so
The first instruction of core 0 has committed. DiffTest enabled.
Running CoreMark for 1 iterations
2K performance run parameters for coremark.
CoreMark Size      : 666
Total time (ms)    : 1943
Iterations         : 1
Compiler version   : GCC11.4.0
seedcrc            : 0xe9f5
[0]crclist         : 0xe714
[0]crcmatrix       : 0x1fd7
[0]crcstate        : 0x8e3a
[0]crcfinal        : 0xe714
Finished in 1943 ms.
=====
CoreMark Iterations/Sec 514.67
Core 0: HIT GOOD TRAP at pc = 0x80001de2
Core-0 instrCnt = 379,235, cycleCnt = 253,872, IPC = 1.493804
Seed=0 Guest cycle spent: 253,877 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 32,729ms
```

Difftest

- **Hands-on:**
 - We injected a bug in a pre-built XiangShan processor
 - Now, let's trigger it by Difftest
- **Showcase**

```
$ bash run-emu-diff.sh
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-1-teration-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so    # difftest with NEMU
```

- Difftest will report the error once failed
- Dump info like PC, GPRs, etc.

```
===== Commit Group Trace (Core 0) =====
commit group [00]: pc 00800015f4 cmtcnt 2
commit group [01]: pc 00800015f8 cmtcnt 1
commit group [02]: pc 00800015fa cmtcnt 2
commit group [03]: pc 0080001600 cmtcnt 1
commit group [04]: pc 008000160e cmtcnt 3
commit group [05]: pc 0080001614 cmtcnt 1
commit group [06]: pc 008000162a cmtcnt 1
commit group [07]: pc 008000162e cmtcnt 4
commit group [08]: pc 008000163e cmtcnt 10
commit group [09]: pc 0080001658 cmtcnt 7
commit group [10]: pc 008000166a cmtcnt 6
commit group [11]: pc 0080001a6e cmtcnt 4
commit group [12]: pc 0080001a7a cmtcnt 1
commit group [13]: pc 0080001a7c cmtcnt 2
commit group [14]: pc 008000167a cmtcnt 2
commit group [15]: pc 008000166c cmtcnt 7 <--
```

```
===== REF Regs =====
----- Integer Registers -----
$0: 0x0000000000000000 ra: 0x00000000800015f0 sp: 0x000000008000cef0 gp: 0x0000000000000000
tp: 0x0000000000000000 t0: 0x0000000000000000 t1: 0x0000000000000001 t2: 0x0000000000000009
s0: 0x0000000000000000 s1: 0x0000000000000000 a0: 0x000000000000029a a1: 0x000000008000cf08
a2: 0x0000000000000002 a3: 0x000000000000029a a4: 0x0000000000000002 a5: 0x0000000000000007
a6: 0x0000000000000002 a7: 0x0000000000000003 s2: 0x0000000000000000 s3: 0x0000000000000000
s4: 0x0000000000000000 s5: 0x0000000000000000 s6: 0x0000000000000000 s7: 0x0000000000000000
s8: 0x0000000000000000 s9: 0x0000000000000000 s10: 0x0000000000000000 s11: 0x0000000000000000
t3: 0x00000000800039b0 t4: 0x0000000000000001 t5: 0x000000008000cd91 t6: 0x0000000000000031
----- Float Registers -----
ft0: 0xffffffff00000000 ft1: 0xffffffff00000000 ft2: 0xffffffff00000000 ft3: 0xffffffff00000000
ft4: 0xffffffff00000000 ft5: 0xffffffff00000000 ft6: 0xffffffff00000000 ft7: 0xffffffff00000000
fs0: 0xffffffff00000000 fs1: 0xffffffff00000000 fa0: 0xffffffff00000000 fa1: 0xffffffff00000000
fa2: 0xffffffff00000000 fa3: 0xffffffff00000000 fa4: 0xffffffff00000000 fa5: 0xffffffff00000000
fa6: 0xffffffff00000000 fa7: 0xffffffff00000000 fs2: 0xffffffff00000000 fs3: 0xffffffff00000000
fs4: 0xffffffff00000000 fs5: 0xffffffff00000000 fs6: 0xffffffff00000000 fs7: 0xffffffff00000000
fs8: 0xffffffff00000000 fs9: 0xffffffff00000000 fs10: 0xffffffff00000000 fs11: 0xffffffff00000000
ft8: 0xffffffff00000000 ft9: 0xffffffff00000000 ft10: 0xffffffff00000000 ft11: 0xffffffff00000000
fcsr: 0x0000000000000000 fflags: 0x0000000000000000 frm: 0x0000000000000000

privilegeMode: 3
a1 different at pc = 0x008000166c, right= 0x000000008000cf08, wrong = 0x000000008000cf10
a3 different at pc = 0x008000166c, right= 0x000000000000029a, wrong = 0x0000000000000000
a6 different at pc = 0x008000166c, right= 0x0000000000000002, wrong = 0x0000000000000001
Core 0: ABORT at pc = 0x80001694
Core-0 instrCnt = 1,264, cycleCnt = 12,760, IPC = 0.099060
Seed=0 Guest cycle spent: 12,764 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 1,677ms
```

Difftest


- **Hands-on:** Dump waveform after Difftest detects the bug
- **Showcase**

```
$ bash run-emu-dumpwave.sh
```

```
$ ls $NOOP_HOME/build/*.vcd -lh # There should be a .vcd waveform
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \  
-b 9000 \      the begin cycle of the wave, you may find XiangShan trap at 10000 cycles in previous step  
-e 11000 \    the end cycle of the wave  
--dump-wave \  set this option to dump wave, you will find *.vcd on $NOOP_HOME/build
```


MinJie Functional Verification





LightSSS: Light-weight Simulation SnapShot

- **Re-run simulation is time-consuming!**
 - Especially for long workload like SPEC CPU
 - Snapshot is the way out
- **LightSSS: snapshots of the process with fork()**
 - Record snapshots of the process with function fork()
 - When bug is detected, it will be waked up and generate waveform of several cycles from the latest snapshot
- **Advantage 1: Good portability and scalability**
 - Snapshots for any external models (such as C++)
 - No need to understand details of external models
- **Advantage 2: Low overhead of snapshots**
 - ~500us for taking a snapshot
 - Far less overhead than RTL snapshots by Verilator



- **Hands-on:** use lightSSS to dump waveform
- **Showcase**

```
$ bash run-emu-lightsss.sh
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-1-iteration-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \  
--enable-fork \      No longer need to use complex parameters  
2 > lightsss.err
```

- LightSSS is working if you see:


```
privilegeMode: 3
  a1 different at pc = 0x008000166c, right= 0x000000008000cf08, wrong = 0x000000008000cf10
  a3 different at pc = 0x008000166c, right= 0x000000000000029a, wrong = 0x0000000000000000
  a6 different at pc = 0x008000166c, right= 0x0000000000000002, wrong = 0x0000000000000001
[FORK_INFO pid(116199)] the oldest checkpoint start to dump wave and dump nemu log...
[FORK_INFO pid(116199)] dump wave to /home/guest/tanghaojin/XiangShan/build/2025-06-19@17:31:04_1.vcd...
%Warning: previous dump at t=0, requesting t=0, dump call ignored
%Warning: previous dump at t=0, requesting t=0, dump call ignored
The first instruction of core 0 has committed. Dfftest enabled.
Running CoreMark for 1 iterations

===== In the last commit group =====
the first commit instr pc of DUT is 0x000000008000166c
the first commit instr pc of REF is 0x000000008000166c

===== Commit Group Trace (Core 0) =====
commit group [00]: pc 00800015f4 cmtcnt 2
commit group [01]: pc 00800015f8 cmtcnt 1
commit group [02]: pc 00800015fa cmtcnt 2
```

- After that, simulation will start again from the existing latest snapshot

MinJie Functional Verification





ChiselDB: Debug-friendly Structured Database

- **Motivation:**

- Waveforms are **large** in size and **hard to apply** further analysis
- Need to analyze structured data like memory transaction trace

- **We present ChiselDB**

- **Highlights:**

- Inserting probes between module interfaces in hardware
- DPI-C: Using C++ function in Chisel code to transfer data
- Persist in database, SQL queries supported

- **Design source code:** *XiangShan/utility/src/main/scala/utility/ChiselDB.scala*
- **Usage:** Create table

```
// API: def createTable[T <: Record](tableName: String, hw: T): Table[T]
import huancun.utils.ChiselDB

class MyBundle extends Bundle {
    val fieldA = UInt(10.W)
    val fieldB = UInt(20.W)
}

val table = ChiselDB.createTable("MyTableName", new MyBundle)
```

- **Usage:** Register probes

```
/* APIs
def log(data: T, en: Bool, site: String = "", clock: Clock, reset: Reset)
def log(data: Valid[T], site: String, clock: Clock, reset: Reset): Unit
def log(data: DecoupledIO[T], site: String, clock: Clock, reset: Reset): Unit
*/
table.log(
  data = my_data /* hardware of type T */,
  en = my_cond,   site = "MyCallSite",
  clock = clock,   reset = reset
)
```

- **Example:** *XiangShan/src/main/scala/xiangshan/cache/mmu/L2TLB.scala*

```
class L2TlbMissQueueDB(implicit p: Parameters) extends TlbBundle {
  val vpn = UInt(vpnLen.W)
}

val L2TlbMissQueueInDB, L2TlbMissQueueOutDB = Wire(new L2TlbMissQueueDB)
L2TlbMissQueueInDB.vpn := missQueue.io.in.bits.vpn
L2TlbMissQueueOutDB.vpn := missQueue.io.out.bits.vpn

val L2TlbMissQueueTable = ChiselDB.createTable(
  "L2TlbMissQueue_hart" + p(XSCoreParamsKey).HartId.toString, new L2TlbMissQueueDB)

L2TlbMissQueueTable.log(L2TlbMissQueueInDB, missQueue.io.in.fire, "L2TlbMissQueueIn", clock, reset)
L2TlbMissQueueTable.log(L2TlbMissQueueOutDB, missQueue.io.out.fire, "L2TlbMissQueueOut", clock, reset)
```

- **TL-Log: Persistence Transactions/Database**
 - Record bus transactions to SQLite3 database using ChiselDB
 - Support offline bug detect and performance analysis

Time	Name	Channel	Opcode	Permission	State	SourceID	SinkID	Address	Data
116854134	L2_L1I_0	A	AcquireBlock	Grow	NtoB	1	0	803d3680	0000000000000000 0000000000000000 0000000000000000 0000000000000000
116854146	L2_L1I_0	D	GrantData	Cap	toT	1	16	803d3680	fa8437839043903 06090063843c23 84382300093783 378300093023c3bd
116854147	L2_L1I_0	D	GrantData	Cap	toT	1	16	803d3680	03e32e07bb030089 f0ef8526c881f49b b60300893783bbdf 855a010925832e07
123191840	L2_L1I_0	C	ReleaseData	Shrink	TtON	0	0	803d3680	fa8437839043903 06090063843c23 84382300093783 378300093023c3bd
123191841	L2_L1I_0	B	Probe	Cap	toN	0	0	803d3680	0000000000000000 0000000000000000 0000000000000000 0000000000000000
123191841	L2_L1I_0	C	ReleaseData	Shrink	TtoN	0	0	803d3680	03e32e07bb030089 f0ef8526c881f49b b60300893783bbdf 855a010925832e07
123191848	L2_L1I_0	D	ReleaseAck	Cap	toT	0	31	803d3680	0000000020ab05b 0000000020fab45b 0000000020fab85b 0000000020fab5b
123191851	L3_L2_0	C	ReleaseData	Shrink	TtoN	31	0	803d3680	fa8437839043903 06090063843c23 84382300093783 378300093023c3bd
123191852	L3_L2_0	C	ReleaseData	Shrink	TtoN	31	0	803d3680	03e32e07bb030089 f0ef8526c881f49b b60300893783bbdf 855a010925832e07
123191862	L3_L2_0	D	ReleaseAck	Cap	toT	31	16	803d3680	86a6f7043583dcd5 f0ef856a865a8762 3583b7654481e17 865a86a68762f704
123191863	L2_L1I_0	C	ProbeAck	Shrink	TtoN	0	0	803d3680	0000000000000000 0000000000000000 0000000000000000 0000000000000000
123191878	L3_L2_0	C	Release	Report	NtoN	30	0	803d3680	5597bf494981aa09 8522cc2585930000 f84ef15dda0f00ef 001947036775e84a

Example: requests sequence in database that violate cache coherence

- Hands-on: Analysis of Cache Coherence Violation

Inject bug – We set all Release Data (L2->L3) as a constant value

```
$ cd ../p2-chiseldb && cat cdb_err.patch
```

```
--- a/src/main/scala/huancun/noninclusive/SinkC.scala
+++ b/src/main/scala/huancun/noninclusive/SinkC.scala

@@ -99,7 +99,7 @@ class SinkC(implicit p: Parameters) extends
BaseSinkC {
-    buffer(insertIdx)(count) := c.bits.data
+    buffer(insertIdx)(count) := 0xABCDEF.U
```

- Hands-on: Analysis of Cache Coherence Violation

```
# simulate using pre-built emu (ChiselDB enabled)
```

```
$ bash chiseldb_step1_run.sh
```

```
# ./emu-cdb-err -i $NOOP_HOME/ready-to-run/Linux.bin \  
--diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpretter-so \  
--dump-db # set this argument to dump data base  
2>Linux.err
```

- **Hands-on: Analysis of Cache Coherence Violation**

- **Difftest Info**

- Error Address: 0x80041940

[illegible]

- Hands-on: Analysis of Cache Coherence Violation

```
# Analyze
```

```
$ bash chiseldb_step2_analyze.sh
```

```
# 1. sqlite query for all transactions on Eaddr
```

```
# 2. use script to parse TLLog
```

```
# sqlite3 $NOOP_HOME/build/*.db \
```

```
"select * from TLLOG where ADDRESS=0x80041940" | \
```

```
sh $NOOP_HOME/scripts/utis/parseTLLog.sh
```



ChiselDB:

Result: [Time | To_From | Channel | Opcode | Permission | Address | Data]

Data successfully transferred from L1D to L2

8922 L2_L1D_0 C ReleaseData Shrink TtoN 80041940

0000000000000000 0000000000000000 0000000080016094 0000000000000000

Data successfully transferred from L2 to L3

13820 L3_L2_0 C ProbeAckData Shrink TtoN 80041940

0000000000000000 0000000000000000 0000000080016094 0000000000000000

But when L1D acquires Eaddr again, data loaded from L3 is wrong

13906 L2_L1D_0 A AcquireBlock Grow NtoB 80041940

13914 L3_L2_0 A AcquireBlock Grow NtoB 80041940


13960 L3_L2_0 D GrantData Cap toT 80041940

0000000000abcdef0000000000000000 0000000000000000 0000000000000000


So there must be something wrong when L3 records Release Data

TL-Test: Cache Verification Framework

- As an infrastructure for cache verification
 - Support **Tilelink protocol** and **cache coherence** check
 - **Fast test cases generation** and **constrained random test**



- **TL-Trace: Trace-based testbench generation**
 - Convert cache access trace into testbench
 - Fast reproduction of functional bugs and analysis of performance problems



TL-Test

Project Structure

```
$ cd $TLT_HOME && tree -d -L 1
```

```
.
├── assets
├── configs
├── dut
├── main
├── run
└── scripts
```

Enter TL-Test tutorial

```
$ cd ../tutorial/p3-tl-test
```


- Hands-on: Analysis of Cache Coherence Violation

Inject bug – We wrongly shift the Grant Data (L2->L1)

```
$ cat tlt_err.patch
```

```
--- a/tl-test-new/dut/CoupledL2/src/main/scala/coupledL2/GrantBuffer.scala
+++ b/tl-test-new/dut/CoupledL2/src/main/scala/coupledL2/GrantBuffer.scala
@@ -94,7 +94,7 @@ class GrantBuffer(implicit p: Parameters) extends
LazyModule {
-    d.data := data
+    d.data := data << 8
```

- Hands-on: Analysis of Cache Coherence Violation

```
# We provide a pre-compiled TL-Test
```

```
# Run the TL-Test
```

```
$ bash tltest_step1_run.sh
```

```
# cd $TLT_HOME/run && ./tltest_v3lt 2>&1 | tee tltest_v3lt.log
```

- Hands-on: Analysis of Cache Coherence Violation
- TL-Test Info
 - Error Address: 0x14000

```
[2236] [tl-test-new-INFO] #0 L2[0].C[0] [data complete D] [GrantData toT] source: 0x5, addr: 0x14000, alias: 0x2, data: [ 00 f3 ... 00 33 ... ]
[2236] [tl-test-new-INFO] #0 L2[0].C[0] [data complete D] [GrantData toT] A -> D latency: 518

dut: [ 00 f3 ... 00 33 ... ]
ref: [ f3 75 ... 33 3b ... ]
[2236] [tl-test-new-ERROR] [tlc_assert failure at int GlobalBoard<unsigned long>::data_check(TLLocalContext *, const uint8_t *, const uint8_t *, std::string) [T = unsigned long]:277]
[2236] [tl-test-new-ERROR] [tlc_assert failure from system #0]
[2236] [tl-test-new-ERROR] info: Data mismatch from status SB_VALID!
[2236] [tl-test-new-ERROR] stack backtrace:
[2236] [tl-test-new-ERROR] #0 /tltest_v3lt(_ZN11GlobalBoardImE10data_checkEP14TLLocalContextPKhS4_NSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEEE
```

- Hands-on: Analysis of Cache Coherence Violation

```
# Analyze
```

```
$ bash tltest_step2_analyze.sh
```

```
# 1. grep all transactions on Eaddr(0x4000)
```

```
# 2. use TL-Log to help debugging
```

```
# grep "addr: 0x14000" $TLT_HOME/run/tltest_v3lt.Log
```

Result: [Time | Core | Channel | Opcode | Source | Address | Data]

L1D acquires Eaddr

[234] L2[0].C[0] [fire A] [AcquireBlock NtoB] source: 0x4, addr: 0x14000

L1D release Eaddr, and data successfully transferred from L1D to L2

[1064] L2[0].C[0] [fire C] [ReleaseData TtoN] source: 0x4, addr: 0x14000, data: [f3 75 ...]

[1066] L2[0].C[0] [fire C] [ReleaseData TtoN] source: 0x4, addr: 0x14000, data: [33 3b ...]


But when L2 grants data of Eaddr, data loaded from L2 is wrong

[2234] L2[0].C[0] [fire D] [GrantData toT] source: 0x5, addr: 0x14000, data: [00 f3 ...]


[2236] L2[0].C[0] [fire D] [GrantData toT] source: 0x5, addr: 0x14000, data: [00 33 ...]

So there must be something wrong when L2 Grant Data

MinJie Performance Verification




MinJie Performance Verification





Checkpoint: Agile Performance Evaluation




Time for performance modelling of single-core XiangShan on SPEC CPU2006

	RTL-simulation	FPGA	RTL-Simulation w/ Checkpoint
Compile	20 minutes	5 hours	20 minutes
Simulation	958 years@2KHz (2K CPS)	7 days@100MHz	5.5 hours with enough x86 servers
Our Approach			



Checkpoint: Agile Performance Evaluation

- Uniform Checkpoint
 - Divide a program to equal segments for parallel simulation with uniform sampling



- SimPoint Checkpoint
 - Select representative segments with cluster analysis and weighted sampling



Checkpoint: SimPoint Checkpoint

- Step 0: Prepare NEMU environment for SimPoint checkpoint

```
$ cd ../p4-checkpoint && bash simpoint_step0_prepare.sh
```

```
# simpoint_step0_prepare.sh (~20s)
```

```
# cd $NEMU_HOME  
# git submodule update --init  
# cd $NEMU_HOME/resource/simpoint/simpoint_repo  
# make clean && make                # generate simpoint generator binary
```

```
# cd $NEMU_HOME  
# make clean  
# make riscv64-xs-cpt_defconfig && make -j 8    # compile NEMU
```

```
# cd $NEMU_HOME/resource/gcpt_restore  
# rm -rf $XS_PROJECT_ROOT/tutorial/part5-checkpoint/gcpt  
# make -C $NEMU_HOME/resource/gcpt_restore/ \    # generate gcpt restorer binary  
O=$XS_PROJECT_ROOT/tutorial/part5-checkpoint/gcpt \ # directory of results  
GCPT_PAYLOAD_PATH=$XS_PROJECT_ROOT/tutorial/part5-  
checkpoint/bin/stream_100000.bin
```

Checkpoint: SimPoint Checkpoint

- Step 1: Execute the workload and collect program behavior

```
$ bash simpoint_step1_profiling.sh
```

```
# simpoint_step1_profiling.sh (~20s)
# source simpoint_env.sh                # configure environment variables

# rm -rf $RESULT

# $NEMU ${BBL_PATH}/${workload}.bin \    # specify workload
#   -b \                                # run with batch mode
#   -D $RESULT \                         # directory where the checkpoint is generated
#   -C $profiling_result_name \          # name of this task
#   -w stream \                          # name of workload
#   --simpoint-profile \                 # is simpoint profiling
#   --cpt-interval ${interval} \         # simpoint interval instructions: 50,000,000
#   > >(tee $log/stream-out.txt) 2> >(tee ${log}/stream-err.txt) # redirect stdout and stderr
```

Checkpoint: SimPoint Checkpoint

```
-----
Your clock granularity/precision appears to be 511 microseconds.
Each test below will take on the order of 351243 microseconds.
(= 687 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
```

```
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
```

Function	Best Rate MB/s	Avg time	Min time	Max time
Copy:	4595.3	0.017409	0.017409	0.017410
Scale:	455.5	0.175621	0.175621	0.175622
Add:	559.3	0.214918	0.214535	0.215047
Triad:	559.4	0.214662	0.214534	0.215046


```
-----
```

```
Solution Validates: avg error less than 1.000000e-06 on all three arrays
-----
```

```
[src/monitor/monitor.c:206,parse_args] Doing Simpoint Profiling
[src/checkpoint/path_manager.cpp:55,init] Cpt id: -1
[src/checkpoint/path_manager.cpp:73,setSimpointProfilingOutputDir] Created /nfs/home/zhangziyue/tutorial_env/tutorial/p4-checkpoint/simpoint_result/simpoint-profiling/stream/
```

```
[src/checkpoint/simpoint.cpp:84,init] Doing simpoint profiling with interval 50000000
[src/checkpoint/serializer.cpp:485,next_index] set next index 0
[src/checkpoint/path_manager.cpp:92,setCheckpointingOutputDir] donot set checkpoint path without Checkpoint mode
[src/isa/riscv64/init.c:220,init_isa] NEMU will start from pc 0x80000000
[src/device/io/port-io.c:35,add_pio_map_with_diff] Add port-io map 'uartlite' at [0x000000000000003f8, 0x00000000000000404]
[src/device/io/port-io.c:35,add_pio_map_with_diff] Add port-io map 'screen' at [0x00000000000000100, 0x00000000000000107]
[src/device/io/port-io.c:35,add_pio_map_with_diff] Add port-io map 'keyboard' at [0x00000000000000060, 0x00000000000000063]
[src/device/sdcard.c:137,init_sdcard] Can not find sdcard image:
[src/monitor/image_loader.c:204,load_img] Loading image (checkpoint/bare metal app/bbl) form cmdline: /nfs/home/zhangziyue/tutorial_env/tutorial/p4-checkpoint/gcpt/build/gcpt.bin
```

*NEMU will do profiling while
simulating workload STREAM at
intervals of 50,000,000 instructions.*



Checkpoint: SimPoint Checkpoint

- Step 2: Cluster and obtain multiple representative slices

```
$ bash simpoint_step2_cluster.sh
```

```
# simpoint_step2_cluster.sh
```

```
# export CLUSTER=$RESULT/cluster/${workload} && mkdir -p $CLUSTER
```

```
# mkdir -p $LOG_PATH/cluster_logs/cluster
```

```
# random1=`head -20 /dev/urandom | cksum | cut -c 1-6`
```

```
# random2=`head -20 /dev/urandom | cksum | cut -c 1-6`
```

```
# $NEMU_HOME/resource/simpoint/simpoint_repo/bin/simpoint \
```

```
#   -loadFVFile $PROFILING_RES/${workload}/simpoint_bbv.gz \    # load a frequency vector file
```

```
#   -saveSimpoints $CLUSTER/simpoints0 \                        # file to save simpoints
```

```
#   -saveSimpointWeights $CLUSTER/weights0 \                   # file to save simpoints weights
```

```
#   -inputVectorsGzipped \                                     # input vectors have been compressed with gzip
```

```
#   -maxK 3 \          # maximum number of clusters to use
```

```
#   -numInitSeeds 2 \   # times of different random initialization for each run, taking only the best clustering
```

```
#   -iters 1000 \       # maximum number of iterations that should perform
```

```
#   -seedkm ${random1} \ # random seed for choosing initial k-means centers
```

```
#   -seedproj ${random2} \ # random seed for random linear projection
```

```
#   > >(tee $Log/${workload}-out.txt) 2> >(tee $Log/${workload}-err.txt)
```

```
    # redirect stdout and stderr
```

Checkpoint: SimPoint Checkpoint

```
-----  
Run number 3 of at most 4, k = 2  
-----
```

```
-----  
Initialization seed trial #1 of 2; initialization seed = 344399  
-----
```

```
-----  
Initialized k-means centers using random sampling: 2 centers
```

```
Number of k-means iterations performed: 1
```

```
BIC score: -39.5836
```

```
Distortion: 10.4664
```

```
Distortions/cluster: 8.08959 2.37678
```

```
Variance: 0.951488
```

```
Variances/cluster: 0.898843 1.18839  
-----
```

Compute the information of
K-means clustering

```
-----  
Initialization seed trial #2 of 2; initialization seed = 344400  
-----
```

```
-----  
Initialized k-means centers using random sampling: 2 centers
```

```
Number of k-means iterations performed: 4
```

```
BIC score: -25.3732
```

```
Distortion: 8.89606
```

```
Distortions/cluster: 4.93125 3.96481
```

```
Variance: 0.808733
```

```
Variances/cluster: 1.23281 0.566402
```

```
The best initialization seed trial was #2  
-----
```

Obtain multiple program
representative checkpoints

```
-----  
Post-processing runs  
-----
```

```
-----  
For the BIC threshold, the best clustering was run 2 (k = 3)  
-----
```

Checkpoint: SimPoint Checkpoint

- Step 3: Generate corresponding Checkpoints based on clustering results

```
$ bash simpoint_step3_genspt.sh # It may take about 2 minutes
```

```
# simpoint_step3_genspt.sh (~2mins)
```

```
# export CLUSTER=$RESULT/cluster  
# mkdir -p $LOG_PATH/checkpoint_logs
```

```
# $NEMU ${BBL_PATH}/${workLoad}.bin \  
# -b \ run with batch mode  
# -D $RESULT \ directory where the checkpoint is generated  
# -C checkpoint \ name of this task  
# -w stream \ name of workload  
# -S $CLUSTER \ simpoints results director  
# --cpt-interval $interval \ simpoint interval instructions: 50,000,000  
# > >(tee $log/stream-out.txt) 2> >(tee $log/stream-err.txt)  
\ redirect stdout and stderr
```


Checkpoint: SimPoint Checkpoint

tool-agnostic

```
[/nfs/home/zhangziyue/tutorial_env/NEMU/src/isa/riscv64/include/ /instr/special h:38 execute] nemu_trap case 101
[src/profiling/profiling_control.c:22,start_profiling] Start profiling. Setting inst count base to Current inst count 270365805.
[src/checkpoint/serializer.cpp:432,instrsCouldTakeCpt] Should take cpt now: 8 next point 0
[src/checkpoint/serializer.cpp:265,serializeRegs] Writing int registers to checkpoint memory @[0xedde0, 0xedee0) [0xedde0, 0xedee0)
[src/checkpoint/serializer.cpp:273,serializeRegs] Writing float registers to checkpoint memory @[0xedee8, 0xedfe8) [0xedee8, 0xedfe8)
[src/checkpoint/serializer.cpp:284,serializeRegs] Writing vector registers to checkpoint memory @[0xedee8, 0xedfe8) [0xfdf8, 0xfe1f8)
[src/checkpoint/serializer.cpp:293,serializeRegs] Writing PC: 0x3f80588552 at addr 0xecdb8
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x105: 0xffffffff801f7318
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x106: 0x7
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x10a: 0xf0
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x10c: 0x1
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x140: 0xffffffffd880179480
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x141: 0x3f80599524
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x142: 0xf
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x143: 0x3f806b31d0
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x14d: 0x203bd8
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x180: 0x8000700000100c54
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x200: 0x200000000
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x24d: 0xffffffffffffffff
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x300: 0xa000066a2
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x301: 0x80000000003411af
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x302: 0xf0b509
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x303: 0x3666
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x304: 0x20a
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x305: 0x80100500
[src/checkpoint/serializer.cpp:316,serializeRegs] CSR 0x306: 0xffffffff

[src/checkpoint/serializer.cpp:335,serializeRegs] Writing CSR to checkpoint memory @[0xedff0, 0xf5ff0) [0xedff0, 0xf5ff0)
[src/checkpoint/serializer.cpp:343,serializeRegs] Touching Flag: 0xbeef at addr 0xecdb0
[src/checkpoint/serializer.cpp:347,serializeRegs] Record mode flag: 0x0 at addr 0xecdc0
[src/checkpoint/serializer.cpp:352,serializeRegs] Record time: 0x0 at addr 0xecdc8
[src/checkpoint/serializer.cpp:356,serializeRegs] Record time: 0x0 at addr 0xecdd0
[src/checkpoint/serializer.cpp:84,serializePMem] Host physical address: 0x100000000 size: 200000000
[src/checkpoint/serializer.cpp:104,serializePMem] Using GZ format generate checkpoint

[src/checkpoint/serializer.cpp:151,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp:151,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp:151,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp:151,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp:151,serializePMem] Written 0x4 bytes
[src/checkpoint/serializer.cpp:250,serializePMem] Checkpoint done!
```


Start making checkpoint

Save int, float & vector registers

Save CSR registers

Save other information

write checkpoints into .gz



Basic format of
RISC-V Checkpoint

Checkpoint: SimPoint Checkpoint

- Step 4: Run the SimPoint checkpoint on NEMU or XiangShan
- Here we take XiangShan as an example

```
$ bash simpoint_step4_run_xs.sh
```

```
# simpoint_step4_run_xs.sh (~ 50s)
```

```
# ./emu \  
#   -i `find $RESULT/checkpoint/stream -type f -name "*_gz" | tail -1` \  
#                                     get the path of workload  
#   --diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpretor-so \  
#                                     Enable the reference standard design path for diff test  
#   --max-cycles=50000 \  
#   2>simpoint.err Maximum execution instruction count
```

Checkpoint: SimPoint Checkpoint

- Step 5: Dump configuration files for batch running on Gem5/XiangShan

```
$ python3 simpoint_step5_dumpresult.py  
$ ls -l simpoint_result/checkpoint
```

- Collect profiling/checkpoint information to generate a list file for Gem5 execution

```
# checkpoint.lst  
stream_5 stream/5 0 0 20 20  
stream_11 stream/11 0 0 20 20  
stream_1 stream/1 0 0 20 20  
Workload Name | Checkpoint Path | Instructions to Skip (default 0) | Functional Warmup Instructions (default 0) | Detailed Warmup Instructions (default 20M) | Sample Instructions (default 20M)
```

- Collect cluster information to generate a JSON file for XiangShan execution

```
# cluster.json  
{"stream": {"insts": "654228721",  
"points": {"7": "0.615385", "0": "0.153846", "11": "0.230769"}}}  
Workload Name | Instruction Count | Checkpoint Index | Checkpoint Weight
```

Checkpoints: Uniform Checkpoint

- **Step 0: Prepare NEMU environment for checkpoints**

```
$ cd $XS_PROJECT_ROOT/tutorial/p4-checkpoint  
$ bash simpoint_step0_prepare.sh
```


- **Step 1: Generate uniform checkpoints using NEMU**

```
$ bash uniform_cpt.sh
```

- **Step 2: Run uniform checkpoint on NEMU or XiangShan**
- Taking XiangShan as an example:

```
$ bash uniform_run_xs.sh
```

MinJie Performance Verification





XSPerf: Simulation Performance Collection

- **Purpose**

- Collect simulation performance data under different requirements

- **XSPerf**

- Multiple types for various fine-grained performance analysis
 - **Accumulation**: basic accumulator counter (log with stderr)
 - **Histogram**: count the distribution (log with stderr)
 - **Rolling**: rolling curve collection and visualization (ChiselDB)

XSPerf: Accumulate & Histogram

- Example

Accumulation:

```
[PERF][time=10000] ctrlBlock.rob: commitInstr,      1500  
[PERF][time=10000] ctrlBlock.rob: waitLoadCycle,    800
```

Histogram:

```
[PERF][time=10000] l2cache: acquire_period_10_20,    15  
[PERF][time=10000] l2cache: acquire_period_20_30,   200  
[PERF][time=10000] l2cache: acquire_period_30_40,    60
```

XSPerf: Accumulate & Histogram

- Usage

- Add `XSPerfAccumulate('name', signal)` or `XSPerfHistogram('name', signal)`
- **By default**, it is printed after simulation finishes
- Set `DebugOptions(EnablePerfDebug = false)` to turn off

- Showcase

```
$ cd ../p5-xs-perf
$ bash xs-perf-log.sh | head -n 30
```

```
# print the log result of p1-basic-func
# cat ${XS_PROJECT_ROOT}/tutorial/p1-basic-func/perf.err
```

XSPerf: Accumulate & Histogram

```
bash xs-perf-log.sh | head -n 30
```

```
cat ${XS_PROJECT_ROOT}/tutorial/p1-basic-func/perf.err
```


```
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_fire,          247
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_PutFullData_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_PutFullData_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_PutPartialData_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_PutPartialData_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_ArithmeticData_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_ArithmeticData_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_LogicalData_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_LogicalData_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_Get_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_Get_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_Hint_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_Hint_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_AcquireBlock_fire,          247
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_AcquireBlock_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_AcquirePerm_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_AcquirePerm_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOClean_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOClean_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOFlush_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOFlush_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOInval_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_A_channel_CBOInval_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_fire,          494
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_AccessAck_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_AccessAck_stall,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_AccessAckData_fire,          0
[PERF ][time=      330342] SimTop.l_soc.socMisc.busPMU: L3_Mem_L3_bank_0_D_channel_AccessAckData_stall,          0
```

Find more parsing scripts at <https://github.com/OpenXiangShan/env-scripts/blob/main/perf/perf.py>


XSPerf: Rolling

- **Advantages**

- Visualization of performance analysis
- Detailed performance differences among segments
- Help identify parameter-sensitive critical regions



compare a performance matrix between two versions code
under the same workload



visualization of performance trends



XSPerf: Rolling

sampling precision(interval)

- Usage

- Add `XSPerfRolling('name', perfCnt, granularity, clock, reset)` wherever you want
- Compile with `WITH_CHISELDB=1 WITH_ROLLINGDB=1`, run with `--dump-db`

- Showcase

```
# Build emu with rolling (time consuming, use pre-built emu instead)
```

```
# bash xs-perf-prepare.sh
```

```
# cd ${NOOP_HOME}
# make clean
# make emu EMU_THREADS=4 WITH_CHISELDB=1 WITH_ROLLINGDB=1 -j8 \
#     PGO_WORKLOAD=${NOOP_HOME}/ready-to-run/coremark-2-iteration.bin \
#     PGO_MAX_CYCLE=10000 PGO_EMU_ARGS=--no-diff LLVM_PROFDATA=llvm-profdata
# ./build/emu -i ./ready-to-run/coremark-2-iteration.bin \
#     --diff ./ready-to-run/riscv64-nemu-interpreter-so --dump-db
# cp `find $NOOP_HOME/build/ -type f -name "*.db" | tail -1` \
#     ${XS_PROJECT_ROOT}/tutorial/p5-xs-perf/xs-perf-rolling.db
```




XSPerf: Rolling

- **Hands-on:** check the results and plot the curve of performance segments

```
# tutorial/p5-xs-perf/  
$ bash xs-perf-rolling.sh
```

```
# cd ${NOOP_HOME}/scripts/rolling  
# python3 rollingplot.py  
${XS_PROJECT_ROOT}/tutorial/p6-xs-  
perf/XsPerfRolling-test.db ipc  
# ls ${NOOP_HOME}/scripts/rolling/results  
  
perf.png
```



rolling curve example



Top-Down: Method for Performance Analysis

- **Purpose**



- Organize scattered performance events in a **hierarchical manner**
- **Accurately** calculate one event's impact on processor performance

- We apply the **Top-Down** approach to XS-Gem5 and XiangShan

- Complete **targeted optimization adaptation** for RISC-V instruction set
- Optimize **performance counters** according to XiangShan microarchitecture
- Further refine the **hierarchical design** of Top-Down model
- Without missing or duplicating any events
- Without assuming the blocking cycles of performance events in advance

Top-Down

- Use Top and Backend as examples





Top-Down in RTL

- Setting Performance Counters in RTL Code

XiangShan/src/main/scala/xiangshan/backend/dispatch/Dispatch.scala

```
val stallReason = Wire(chiselTypeOf(io.stallReason.reason))  
// ...  
TopDownCounters.values.foreach(ctr =>  
  XSPerfAccumulate(ctr.toString(), PopCount(stallReason.map(_ === ctr.id.U)))  
)
```

- Do simulation to collect performance counter data



Top-Down in RTL

- Analyze the performance counter data through Top-Down method

XiangShan/scripts/top-down/configs.py

```
xs_coarse_rename_map = {  
    'OverrideBubble': 'MergeFrontend',  
    'FtqFullStall': 'MergeFrontend',  
    'IntDqStall': 'MergeCoreDQStall', # attribute perf-counters to Top-Down hierarchy  
    'FpDqStall': 'MergeCoreDQStall'  
}
```

- Obtain analysis results and make targeted optimizations



Top-Down in RTL

- **Hands-on:** analyze the congestion causes in RTL using the Top-Down tool

```
#run the Top-Down analysis scripts (it takes ~1min)
```


```
$ bash rtl-top-down.sh
```

```
# cd ${NOOP_HOME}/scripts/top-down && \  
python3 top_down.py \  
-s /opt/SPEC06_EmuTasks_topdown \      #path of performance counter results  
-j /opt/SPEC06_EmuTasks_topdown.json #json file of checkpoints configuration  
  
# ls ${NOOP_HOME}/scripts/top-down/results  
result.png results.csv results-weighted.csv #output of visual analysis results
```




Top-Down in RTL

- Result Example: Top-Down visual analysis results



ConstantIn

- **Motivation**

- Want to test the performance under different parameters
- But re-compilation is time consuming
- Can we change parameters (constants) at runtime?


- **ConstantIn**

- DPI-C: Using C++ function in Chisel code (through Verilog using BlackBox)
- Signal values configured at runtime rather than compile time



ConstantIn

- Constants can be changed when initializing
- And remain unchanged during runtime





ConstantIn

- Usage: Create signals

```
# XiangShan/coupledL2/src/main/scala/coupledL2/prefetch/TemporalPrefetch.scala
```

```
139 require(cacheParams.hartIds.size == 1)
140 val hartid = cacheParams.hartIds.head
141 // 0 / 1: whether to enable temporal prefetcher
142 private val enableTP = WireInit(Constantin.createRecord("enableTP"+hartid.toString, initValue = 1.U))
143 // 0 ~ N: throttle cycles for each prefetch request
144 private val tpThrottleCycles = WireInit(Constantin.createRecord("tp_throttleCycles"+hartid.toString, initValue = 4.U(3.W)))
145 // 0 / 1: whether request to set as trigger on meta hit
146 private val hitAsTrigger = WireInit(Constantin.createRecord("tp_hitAsTrigger"+hartid.toString, initValue = 1.U))
147 // 1 ~ triggerQueueDepth: enqueue threshold for triggerQueue
148 private val triggerThres = WireInit(Constantin.createRecord("tp_triggerThres"+hartid.toString, initValue = 1.U(3.W)))
149 // 1 ~ tpEntryMaxLen: record threshold for recorder and sender (storage size will not be affected)
150 private val recordThres = WireInit(Constantin.createRecord("tp_recordThres"+hartid.toString, initValue = tpEntryMaxLen.U))
151 // 0 / 1: whether to train on vaddr
152 private val trainOnVaddr = WireInit(Constantin.createRecord("tp_trainOnVaddr"+hartid.toString, initValue = 0.U))
153 // 0 / 1: whether to eliminate L1 prefetch request training
154 private val trainOnL1PF = WireInit(Constantin.createRecord("tp_trainOnL1PF"+hartid.toString, initValue = 0.U))
```



ConstantIn

- **Hands-on:** Pass constant via standard input stream
 - **By default**, read constant via `${NOOP_HOME}/build/constantin.txt`
 - **For better demonstration**, it is changed to stdin (corresponding modifications can refer to `p6-constantin/utility.patch`)

```
$ cd ../p6-constantin
```

```
# bash step0-build.sh      # Build emu with constantin (time consuming, use pre-built emu instead)
```

```
$ bash step1-basic.sh      # Run emu and pass constant
```

```
please input total constant number
```

```
2
```

```
please input each constant ([constant name] [value])
```

```
pbop_delayQueueLatency 150
```

```
vbop_delayQueueLatency 150
```



ConstantIn

- **Feature:** *AutoSolving* for automatically finding better constant values
- **Process**
 - Enable *AutoSolving* & prepare signals you need
 - Compile the emu with `WITH_CONSTANTIN=1`
 - Set parameters for *AutoSolving* in a configuration file
 - Run the script to find automatically



ConstantIn

- **Feature:** *AutoSolving* for automatically finding better constant values

```
# Automatical parameter solver configuration file
```

```
$ cat my_constantin.json
```

- **Parameters defined**
 - properties of signals
 - target of *AutoSolving*
 - parameters of genetic algorithm
 - parameters of XS emulator

```
cat my_constantin.json
{
  "constants": [
    {
      "name": "vbop_delayQueueLatency",
      "width": 8,
      "guide": 256,
      "init": 177
    },
    {
      "name": "pbop_delayQueueLatency",
      "width": 8,
      "guide": 256,
      "init": 242
    }
  ],
  "opt_target": [
    {"l2cache.topDown: L2MissMatch,": {"policy": "min", "baseline": 409} },
    {"backend.ctrlBlock.rob: clock_cycle,": {"policy": "min", "baseline": 217759} }
  ],
  "population_num": 4,
  "iteration_num": 2,
  "crossover_rate": 50,
  "mutation_rate": 50,

  "emu_threads": 8,
  "concurrent_emu": 4,
  "max_instr": 50000,
  "seed": 3888,
  "work_load": "${XS_PROJECT_ROOT}/tutorial/p6-constantin/maprobe-riscv64-xs.bin"
}
```



ConstantIn

- Hands-on: *AutoSolving* example

```
# Run auto solver, output will be the optimal constant currently found
```

```
$ bash step2-solve.sh
```


```
# ~2min
```

```
# The solver generates optimal parameters
```

```
opt constant in this round is [['vbop_delayQueueLatency', 242], ['pbop_delayQueueLatency', 135]] fitness is -316
```

```
opt constant for gene algrithom is [['vbop_delayQueueLatency', 243], ['pbop_delayQueueLatency', 135]] fitness -316
```


MinJie Performance Verification






Chisel: Object-Oriented Hardware Design

- Code Efficiency
 - 5-10x fewer LOCs vs. Verilog
- Reusable IP blocks
 - Reduce code duplication with *Inheritance*

• Task #1: Design an L2 Cache for RISC-V Rocket-chip core • Who: A 5-year engineer vs. a senior student		
	A 5-year Engineer	An Undergraduate
Experience	Familiar w/ OpenSparc T1; Modified Xilinx Cache	Course projects on CPU design; Using Chisel for 9 months
Language	Verilog	Chisel
Time	6 weeks	3 days
LOCs	~1700	~350
Results	Unable to boot Linux	Boot Linux on multicore RISC-V ; support NIC w/ DMA mode
• 1st Round results: Chisel is more productive than Verilog by 14X with only 1/5 LOC		




Pixel
(gate-level)

Circuit-oriented
Processor Design



Processor




Object
(Reusable Modules)

Object-oriented
Processor Design



Processor





Configurable



Example: How to implement a prefetcher


- How to abstract prefetching algorithms?
 - Common parameters
 - Common module IO





Example: How to implement a prefetcher

- Use an *out-of-the-box* L2 Cache prefetch algorithm template





Example: How to implement a prefetcher

- Use an *out-of-the-box* L2 Cache prefetch algorithm template

```
# apply patch for the out-of-the-box prefetch template
```

```
$ git -C $NOOP_HOME apply prefetch_template.patch
```

```
# apply patch for a next-line prefetcher
```

```
$ git -C $NOOP_HOME apply next_line.patch
```

1. Define corresponding parameters
2. Implement prefetcher algorithm (e.g., a next-liner prefetcher)

```
--- a/src/main/scala/coupledL2/prefetch/PrefetchTemplate.scala
+++ b/src/main/scala/coupledL2/prefetch/PrefetchTemplate.scala
@@ -43,6 +43,11 @@ abstract class MyPrefetchModule(implicit val p: Parameters) extends Module with
   class MyPrefetch(implicit p: Parameters) extends MyPrefetchModule {
     val io = IO(new L2PrefetchIO())

-    // TODO
+    val addr = io.train.bits.addr + (1 << offsetBits).U
+
+    io <> DontCare
+    io.req.valid := io.train.valid
+    io.req.bits.tag := parseFullAddress(addr)._1
+    io.req.bits.set := parseFullAddress(addr)._2
+    io.train.ready := true.B
   }
```



Example: How to implement a prefetcher

- Use an *out-of-the-box* L2 Cache prefetch algorithm template
3. Instantiate an object and wire its connection

```

--- a/src/main/scala/coupledL2/prefetch/Prefetcher.scala
+++ b/src/main/scala/coupledL2/prefetch/Prefetcher.scala
@@ -291,6 +291,8 @@ class Prefetcher(implicit p: Parameters) extends PrefetchModule {
    val tp = if (hasTPPrefetcher) Some(Module(new TemporalPrefetch())) else None
    // prefetch from upper level
    val pfRcv = if (hasReceiver) Some(Module(new PrefetchReceiver())) else None
+   val hasMyPrefetch = prefetchers.exists(_.isInstanceOf[MyPrefetchParameters])
+   val myPrefetch = if (hasMyPrefetch) Some(Module(new MyPrefetch())) else None

    // ===== Connection for each Prefetcher =====
    // Rcv > VBOP > PBOP > TP
@@ -341,6 +343,10 @@ class Prefetcher(implicit p: Parameters) extends PrefetchModule {

    tp.get.io.tpmeta_port <> tpio.tpmeta_port.get
  }
+   if (hasMyPrefetch) {
+     myPrefetch.get.io <> DontCare
+     myPrefetch.get.io.train <> io.train
+   }

  private val mbistPl = MbistPipeline.PlaceMbistPipeline(2, "MbistPipeL2Prefetcher", cach

    // ===== Connection of all Prefetchers =====
@@ -356,6 +362,7 @@ class Prefetcher(implicit p: Parameters) extends PrefetchModule {
    case _: PrefetchReceiverParams => pfRcv.get.io.req
    case _: BOPParameters          => bopReq
    case _: TPParameters           => tp.get.io.req
+   case _: MyPrefetchParameters   => myPrefetch.get.io.req
  },
  out = pftQueue.io.enq,

```

Instantiate myPrefetch object

Connect training interface

Connect request interface



Example: How to implement a prefetcher

- Use an *out-of-the-box* L2 Cache prefetch algorithm template
 4. Open the prefetch in Configs

```
--- a/src/main/scala/top/Configs.scala
+++ b/src/main/scala/top/Configs.scala
@@ -319,7 +319,8 @@ case class L2CacheConfig
    // NOTICE: the prefetchers are sorted by priority
    prefetch = (if (p.prefetcher.nonEmpty) Seq(PrefetchReceiverParams()) else Nil) ++
      Seq(BOPParameters()) ++
-      (if (tp) Seq(TPParameters()) else Nil),
+      (if (tp) Seq(TPParameters()) else Nil) ++
+      Seq(MyPrefetchParameters()),
    enableL2Flush = enableFlush,
    enablePerf = !site(DebugOptionsKey).FPGAPlatform && site(DebugOptionsKey).EnablePerfDebug,
    enableRollingDB = site(DebugOptionsKey).EnableRollingDB,
```




Example: How to implement a prefetcher

- Use an *out-of-the-box* L2 Cache prefetch algorithm template

```
[814] ===== Inclusive TL-CHI L2 (256.0KB * 4-bank) =====  
[814] prefetch (sort by priority):  
[814] 1. Receiver for prefetches from L1  
[814] 2. Best-offset prefetch (hybrid vbop and pbop)  
[814] 3. My prefetch :)
```


MinJie Performance Verification





XS-Gem5: Architecture Performance Evaluation Tool

- **Purpose**

- Explore microarchitecture designs efficiently
- Conduct fast end-to-end performance evaluations

- **XS-Gem5**

- Align with XiangShan RTL based on Gem5 O3CPU with modifications
- Support Dfftest and Checkpoint
- Scripts for XiangShan performance analysis

XS-Gem5

- **Hands-on:** compile and build XS-Gem5

```
$ cd ../p7-xs-gem5 # Enter XS-Gem5 tutorial directory
```

```
$ bash 0-gem5_prepare.sh # We have built this for you, skip it on our server to save cpu time and memory
```

```
prepare_gem5() {  
    pushd $gem5_home && \  
    cd ext/dramsim3 && \# build DRAMSim3  
    git clone git@github.com:umd-memsys/DRAMSim3.git DRAMsim3 && \  
    cd DRAMsim3 && mkdir -p build && cd build && cmake .. && make -j `nproc` && \  
    popd  
}  
build_gem5() {  
    pushd $gem5_home && \  
    scons build/RISCV/gem5.opt --gold-linker -j `nproc` && \  
    popd  
}
```

XS-Gem5

- **Hands-on:** run bare-metal workload on XS-Gem5

```
$ bash 1-gem5_run_coremark.sh
```

```
#~9s
pushd $gem5_home && \
export GCBV_REF_SO=$gem5_home/ext/NEMU/build/riscv64-nemu-interpreter-so && \
mkdir -p util/xs_scripts/coremark && \ # Setup CoreMark working directory
cd util/xs_scripts/coremark && \
$gem5_home/build/RISCV/gem5.opt $gem5_home/configs/example/xiangshan.py \
--raw-cpt \
--generic-rv-cpt=$NOOP_HOME/ready-to-run/coremark-2-iteration.bin && \
popd
```

```
build/RISCV/mem/physical.cc:608: info: First 4 bytes are 0x93 0x0 0x0 0x0
build/RISCV/sim/system.cc:561: info: Restored from Xiangshan RISC-V Checkpoint
**** REAL SIMULATION ****
build/RISCV/sim/simulate.cc:194: info: Entering event queue @ 0. Starting simulation...
build/RISCV/cpu/base.cc:1266: warn: Start memcpy to NEMU from 0x7934dca00000, size=8589934592
build/RISCV/cpu/base.cc:1269: warn: Start regcpy to NEMU
build/RISCV/dev/serial/uartlite.cc:35: warn: Write to other uartlite addr 12 is not implemented
Running CoreMark for 2 iterations
2K performance run parameters for coremark.
CoreMark Size      : 666
Total time (ms)    : 88
Iterations         : 2
Compiler version   : GCC13.2.0
seedcrc            : 0xe9f5
[0]crclist         : 0xe714
[0]crcmatrix       : 0x1fd7
[0]crcstate        : 0x8e3a
[0]crcfinal        : 0x72be
Finished in 88 ms.
=====
CoreMark Iterations/Sec 22727
Exiting @ tick 291149559 because m5_exit instruction encountered when simulating XS
~/tutorial_env/tutorial/p7-xs-gem5
```

- **Hands-on:** analyze the performance counters

```
$ bash 2-gem5_counter.sh | head -n 20
```

```
cat $gem5_home/util/xs_scripts/coremark/m5out/stats.txt
```

```
----- Begin Simulation Statistics -----
simSeconds          0.000291          # Number of seconds simulated (Second)
simTicks            291149559         # Number of ticks simulated (Tick)
finalTick           291149559         # Number of ticks from beginning of simulation (restored from checkpoints and never reset)
(Tick)
simFreq             1000000000000      # The number of ticks per simulated second ((Tick/Second))
hostSeconds         11.23             # Real time elapsed on the host (Second)
hostTickRate        25929758         # The number of ticks simulated per host second (ticks/s) ((Tick/Second))
hostMemory          17142900         # Number of bytes of host memory used (Byte)
simInsts            663614           # Number of instructions simulated (Count)
simOps              663614           # Number of ops (including micro ops) simulated (Count)
hostInstRate        59101            # Simulator instruction rate (inst/s) ((Count/Second))
hostOpRate          59101            # Simulator op (including micro ops) rate (op/s) ((Count/Second))
system.bridge.power_state.pwrStateResidencyTicks::UNDEFINED 291149559 # Cumulative time (in ticks) in various power states (Tick)
system.clk_domain.clock 333         # Clock period in ticks (Tick)
system.cpu.numCycles 874328         # Number of cpu cycles simulated (Cycle)
system.cpu.numWorkItemsStarted 0      # Number of work items this cpu started (Count)
system.cpu.numWorkItemsCompleted 0    # Number of work items this cpu completed (Count)
system.cpu.instsAdded 676980        # Number of instructions added to the IQ (excludes non-spec) (Count)
system.cpu.nonSpecInstsAdded 18      # Number of non-speculative instructions added to the IQ (Count)
```



XS-Gem5: Performance Analysis Scripts - TopDown

- **Hands-on:** run the Top-Down analysis script

```
$ bash 3-gem5_topdown.sh
```

```
pushd gem5_data_proc && \  
python3 batch.py \  
  -s $gem5_home/util/xs_scripts/coremark \  
  -t --topdown-raw && \  
popd
```

```
Process finished job: m5out_0  
/nfs/home/zhangziyue/tutorial_env/gem5/util/xs_scripts/coremark/m5out/stats.txt  
{'m5out_0': {'OtherFetchStall': 0, 'OtherStall': 0, 'CommitSquash': 163020, 'ResumeUnblock': 0, 'OtherMemSta  
ll': 60, 'Atomic': 0, 'MemCommitRateLimit': 2382, 'MemNotReady': 1718802, 'MemSquashed': 0, 'StoreMemBound':  
  0, 'StoreL3Bound': 0, 'StoreL2Bound': 0, 'StoreL1Bound': 2232, 'LoadMemBound': 4470, 'LoadL3Bound': 0, 'Loa  
dL2Bound': 0, 'LoadL1Bound': 672, 'InstNotReady': 204, 'VectorLongExecute': 0, 'ScalarLongExecute': 64212, '  
SerializeStall': 63792, 'InstSquashed': 44832, 'InstMisPred': 0, 'FetchBufferInvalid': 0, 'SquashStall': 729  
0, 'TrapStall': 0, 'IntStall': 0, 'BpStall': 116700, 'DTlbStall': 192, 'ITlbStall': 0, 'IcacheStall': 224154  
, 'NoStall': 778901, 'ipc': 0.758999, 'Insts': 663614, 'Cycles': 874328, 'point': '0', 'workload': 'm5out',  
'bmk': 'm5out'}}
```



XS-Gem5: Performance Analysis Scripts – Cache MPKI

- **Hands-on:** run the Cache MPKI analysis script

```
$ bash 4-gem5_cache.sh
```

```
pushd gem5_data_proc && \  
mkdir -p results && \  
export PYTHONPATH=`pwd` && \  
python3 batch.py \  
    -s /home/share/xs-model-11bank \  
    -o gem5-cache-example.csv \  
    --cache && \  
python3 simpoint_cpt/compute_weighted.py \  
    -r gem5-cache-example.csv \  
    -j /home/share/xs-model-11bank/cluster-0-0.json \  
    -o weighted.csv
```

Note: Here we take the results from existing output from running SPEC CPU 2006 checkpoints.



XS-Gem5: Performance Analysis Scripts – Cache MPKI

	Cycles	Insts	dcache_miss	l1d.data.MPKI	l1d.data.miss	...	l3_acc_l1_pref	l3_miss	l3_miss_l1d_pref	cpi	coverage
mcf	3.027e+07	2.000e+07	5.413e+06	270.673	5.413e+06	...	1.575e+06	698680.709	575668.623	1.513	1.000
omnetpp	1.774e+07	2.000e+07	1.154e+06	57.713	1.154e+06	...	7.398e+05	492397.438	413052.400	0.887	1.000
astar	1.618e+07	2.000e+07	5.450e+05	27.252	5.450e+05	...	2.588e+05	48744.997	41617.766	0.809	1.000
gobmk	1.155e+07	2.000e+07	1.097e+05	5.485	1.097e+05	...	3.230e+04	14055.195	10964.674	0.577	1.000
soplex	1.119e+07	2.000e+07	6.678e+05	33.391	6.678e+05	...	7.884e+05	175472.264	155213.001	0.559	1.000
milc	1.071e+07	2.000e+07	1.375e+06	68.772	1.375e+06	...	9.581e+05	702851.431	692453.110	0.535	1.000
bzip2	9.541e+06	2.000e+07	2.442e+05	12.212	2.442e+05	...	5.423e+04	4635.961	3321.498	0.477	0.988
gromacs	9.386e+06	2.000e+07	6.421e+05	32.104	6.421e+05	...	1.865e+04	7220.444	6523.288	0.469	1.000
gcc	9.261e+06	2.000e+07	2.785e+05	13.926	2.785e+05	...	3.880e+05	114394.979	104616.710	0.463	1.000
sjeng	9.021e+06	2.000e+07	2.879e+04	1.439	2.879e+04	...	1.448e+04	20958.788	13269.176	0.451	1.000
perlbench	8.118e+06	2.000e+07	1.068e+05	5.338	1.068e+05	...	3.076e+04	23625.296	18072.642	0.406	0.820
bwaves	8.071e+06	2.000e+07	1.296e+05	6.480	1.296e+05	...	5.677e+05	89426.320	87864.638	0.404	1.000
lbm	7.920e+06	2.000e+07	1.606e+05	8.029	1.606e+05	...	6.896e+05	122392.254	119831.519	0.396	1.000
zeusmp	7.450e+06	2.000e+07	2.486e+05	12.428	2.486e+05	...	1.805e+05	85596.788	82176.235	0.373	1.000
tonto	6.759e+06	2.000e+07	1.067e+05	5.334	1.067e+05	...	7.943e+04	5793.963	5075.233	0.338	1.000
calculix	6.720e+06	2.000e+07	1.863e+04	0.932	1.863e+04	...	8.803e+03	3756.312	3651.163	0.336	1.000
GemsFDTD	6.559e+06	2.000e+07	2.564e+05	12.818	2.564e+05	...	5.049e+05	374770.295	359822.201	0.328	1.000
leslie3d	6.552e+06	2.000e+07	4.255e+05	21.274	4.255e+05	...	3.793e+05	144669.463	142838.003	0.328	1.000
xalancbmk	6.515e+06	2.000e+07	3.998e+05	19.989	3.998e+05	...	2.924e+05	38865.836	30536.527	0.326	1.000
namd	6.494e+06	2.000e+07	2.828e+05	14.142	2.828e+05	...	5.306e+03	4312.841	4150.029	0.325	1.000
povray	6.162e+06	2.000e+07	3.497e+05	17.487	3.497e+05	...	3.765e+03	3190.261	3044.166	0.308	1.000
dealII	5.946e+06	2.000e+07	1.630e+05	8.148	1.630e+05	...	8.439e+04	10109.797	7289.549	0.297	1.000
sphinx3	5.736e+06	2.000e+07	2.923e+05	14.614	2.923e+05	...	3.447e+05	20540.266	19216.285	0.287	1.000
h264ref	4.941e+06	2.000e+07	7.420e+04	3.710	7.420e+04	...	2.751e+04	10893.001	9186.843	0.247	1.000
wrf	4.908e+06	2.000e+07	5.679e+04	2.839	5.679e+04	...	1.133e+05	12708.207	12142.350	0.245	1.000
hmmer	4.489e+06	2.000e+07	2.084e+04	1.042	2.084e+04	...	2.627e+04	494.690	418.687	0.224	1.000
cactusADM	4.474e+06	2.000e+07	4.308e+04	2.154	4.308e+04	...	9.758e+04	34987.191	34295.763	0.224	1.000
libquantum	4.390e+06	2.000e+07	2.371e+04	1.185	2.371e+04	...	5.110e+05	69486.561	68474.121	0.220	1.000
gamess	4.261e+06	2.000e+07	4.574e+04	2.287	4.574e+04	...	3.286e+03	2324.679	2228.684	0.213	1.000

[29 rows x 17 columns]



XS-Gem5: Performance Analysis Scripts

- **Hands-on:** run the SPEC CPU score calculation script

```
$ bash 5-gem5_spec06_score.sh
```

```
pushd gem5_data_proc && \  
mkdir -p results && \  
export PYTHONPATH=`pwd` && \  
python3 batch.py -s ../data/xs-model-11bank \  
-o gem5-score-example.csv && \  
python3 simpoint_cpt/compute_weighted.py \  
-r gem5-score-example.csv \  
-j ../data/xs-model-11bank/cluster-0-0.json \  
--score score.csv
```

Note: Here we take the results from existing output from running SPEC CPU 2006 checkpoints.



SPEC CPU score

```
===== SPEC06 =====
===== Int =====
      time  ref_time  score  coverage
perlbench  289.646   9770.0  11.244    1.0
bzip2      410.138   9650.0   7.843    1.0
gcc        167.812   8050.0  15.990    1.0
mcf        141.107   9120.0  21.544    1.0
gobmk      320.031  10490.0  10.926    1.0
hmmer      260.274   9330.0  11.949    1.0
sjeng      366.934  12100.0  10.992    1.0
libquantum 148.529  20720.0  46.500    1.0
h264ref    428.160  22130.0  17.229    1.0
omnetpp    139.743   6250.0  14.908    1.0
astar      229.743   7020.0  10.185    1.0
xalanmk    85.528    6900.0   26.892    1.0
```

```
Estimated Int score per GHz: 15.118648674723131
Estimated Int score @ 3.0GHz: 45.35594602416939
```


```
===== FP =====
      time  ref_time  score  coverage
bwaves     175.907  13590.0  25.752    1.0
games      451.597  19580.0  14.452    1.0
milc       142.358   9180.0  21.495    1.0
zeusmp     196.183   9100.0  15.462    1.0
gromacs    201.133   7140.0  11.833    1.0
cactusADM  252.748  11950.0  15.760    1.0
leslie3d   214.233   9400.0  14.626    1.0
namd       210.250   8020.0  12.715    1.0
dealII     144.375  11440.0  26.413    1.0
soplex     136.694   8340.0  20.337    1.0
povray      93.311   5320.0  19.005    1.0
calculix   490.428   8250.0   5.607    1.0
GemsFDTD   176.662  10610.0  20.019    1.0
tonto      260.012   9840.0  12.615    1.0
lbm        149.280  13740.0  30.681    1.0
wrf        250.853  11170.0  14.843    1.0
snhinx3    350.350  19490.0  18.543    1.0
```

```
Estimated FP score per GHz: 16.548846674376477
Estimated FP score @ 3.0GHz: 49.64654002312943
```

```
===== Overall =====
Estimated overall score per GHz: 15.941323843383495
Estimated overall score @ 3.0GHz: 47.82397153015049
```



Summary: MinJie Development Flows and Tools



Thanks!